

Selected Queries Resource Utilization

Todor Ivanov, Matteo Pergolesi

September 11, 2018

Contents

1 Q08	3
1.1 Type	3
1.2 Description	3
1.3 Performance	3
1.4 Phase 1 - Data preparation	3
1.5 Phase 2 - Query	3
1.6 Source	3
1.7 Notes	3
1.8 No compression - ORC vs Parquet	4
1.9 Snappy - ORC vs. Parquet	10
1.10 ORC - NoCompression vs. Snappy	16
1.11 Parquet - NoCompression vs Snappy	22
2 Q10	28
2.1 Type	28
2.2 Description	28
2.3 Performance	28
2.4 Phase 1	28
2.5 Source	28
2.6 Notes	28
2.7 No compression - ORC vs Parquet	29
2.8 Snappy - ORC vs. Parquet	35
2.9 ORC - NoCompression vs. Snappy	41
2.10 Parquet - NoCompression vs Snappy	47
3 Q12	53
3.1 Type	53
3.2 Description	53
3.3 Performance	53
3.4 Phase 1	53
3.5 Source	53
3.6 Notes	53
3.7 No compression - ORC vs Parquet	54
3.8 Snappy - ORC vs. Parquet	60
3.9 ORC - NoCompression vs. Snappy	66
3.10 Parquet - NoCompression vs Snappy	72
4 Q25	78
4.1 Type	78
4.2 Description	78
4.3 Performance	78
4.4 Phase 1	78
4.5 Phase 2	78
4.6 Phase 3	78
4.7 Source	78

CONTENTS

4.8 Notes	78
4.9 No compression - ORC vs Parquet	79
4.10 Snappy - ORC vs. Parquet	85
4.11 ORC - NoCompression vs. Snappy	91
4.12 Parquet - NoCompression vs Snappy	97

1 Q08

1.1 Type

MapReduce/Python

1.2 Description

For online sales, compare the total sales monetary amount in which customers checked online reviews 3 days before making the purchase and that of sales in which customers did not read reviews. Consider only online sales for a specific category in a given year.

1.3 Performance

	Default (sec.)		No compr. (sec.)		Snappy (sec.)	
Query	ORC	Parquet	ORC	Parquet	ORC	Parquet
Q08	913	675	911	712	900	639

1.4 Phase 1 - Data preparation

- Temp table 1: dates
- Temp table 2: users哪读评论 (get webclicks and use python sessionize)
- Temp table 3: sales on desired date range

1.5 Phase 2 - Query

- Sum on sales_reviews and (sales_all - sales_review).
- Result is a single line for money amount comparison

1.6 Source

<https://git.io/vpFen>

1.7 Notes

Phase 1 should be the interesting one (it's where you use file formats).

1.8 No compression - ORC vs Parquet

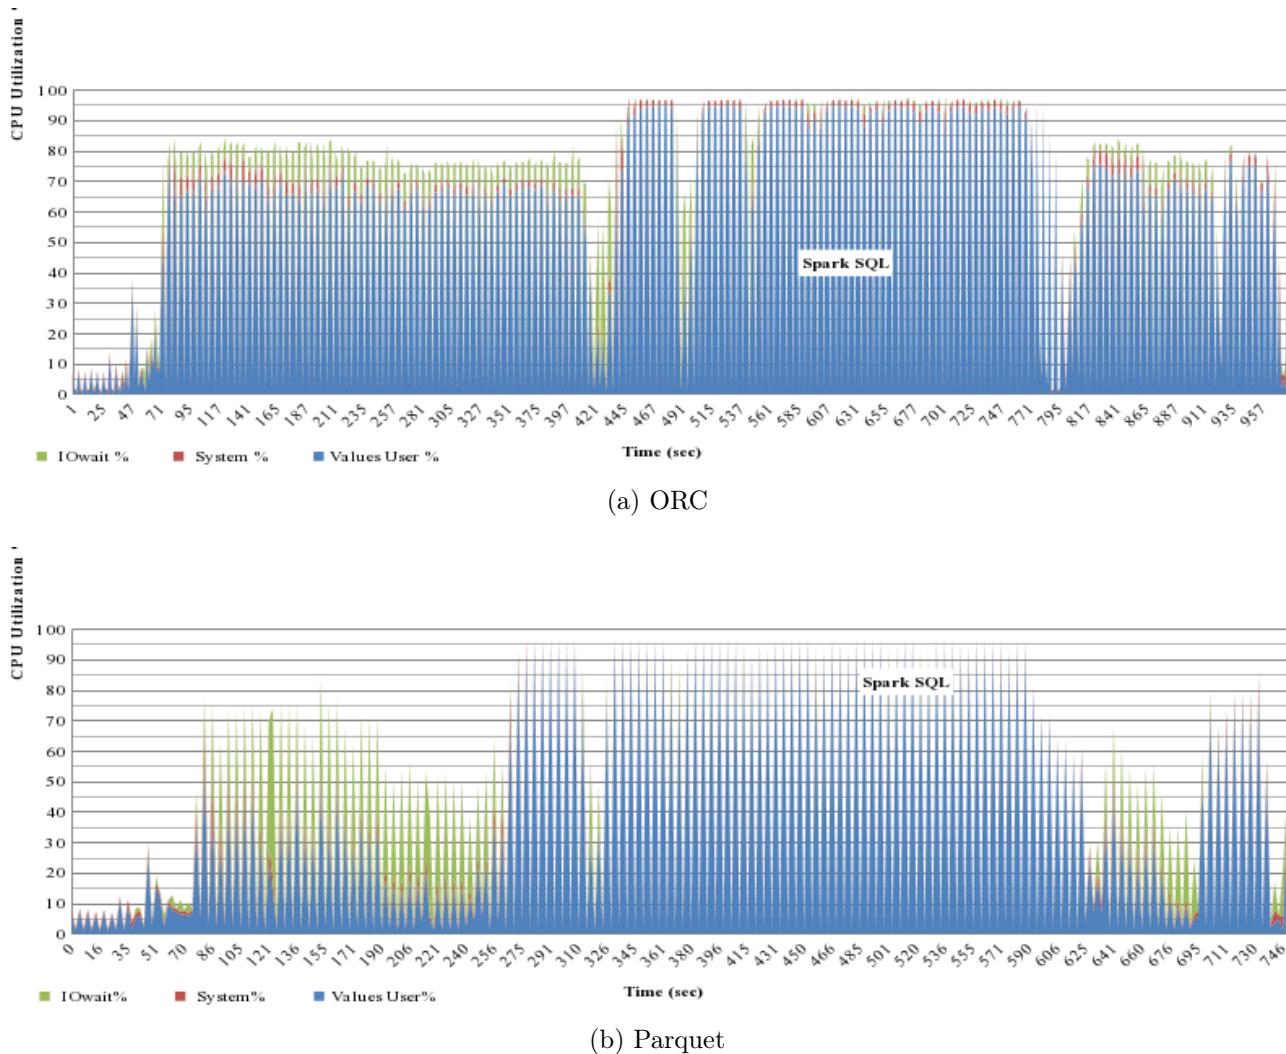


Figure 1: CPU: Lower utilization on Parquet (max is 80%)

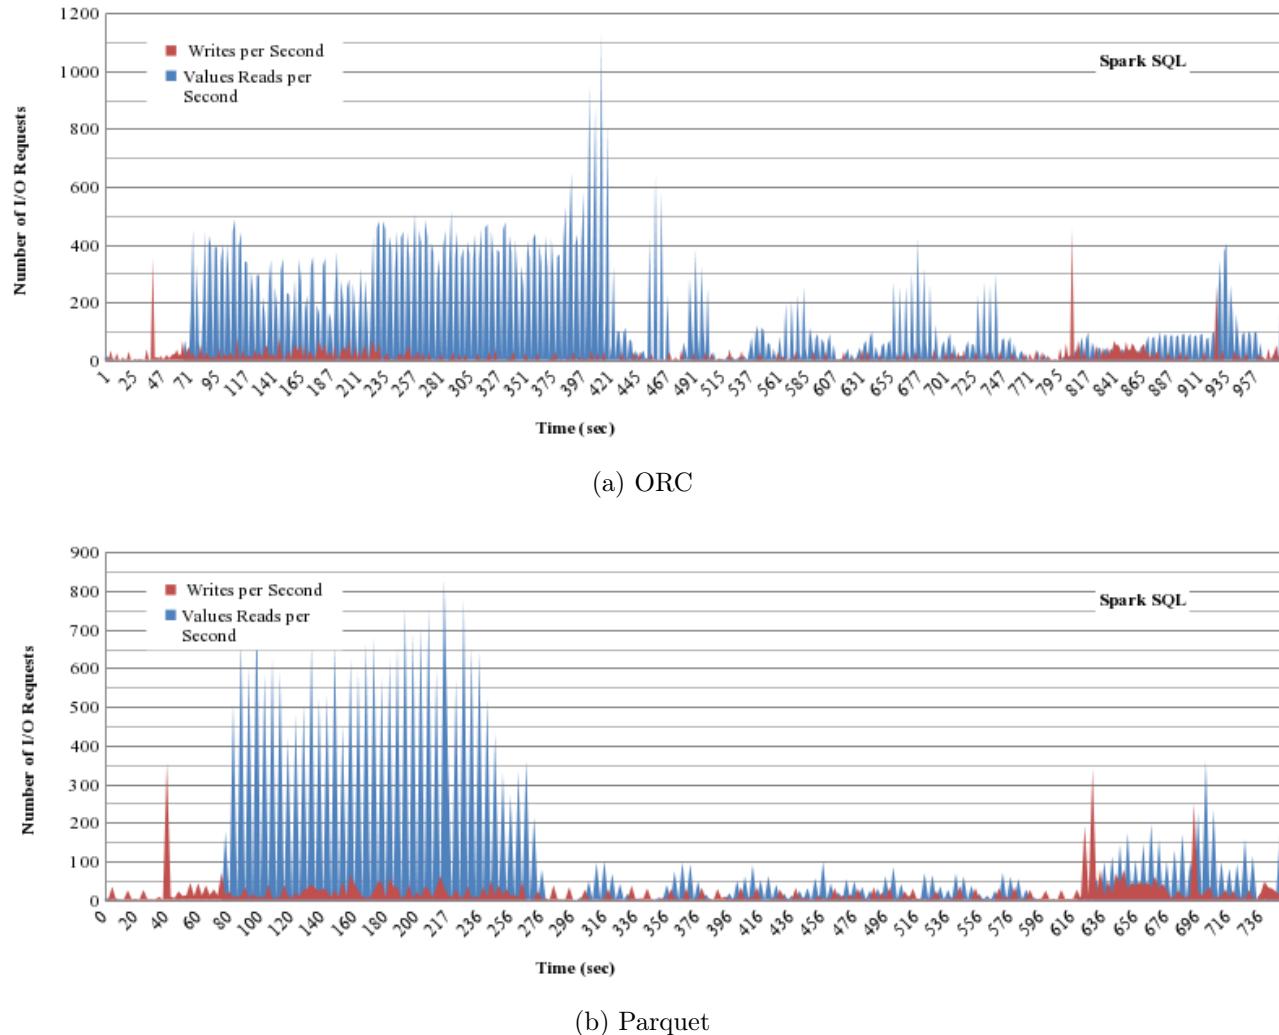


Figure 2: Disk Requests: More requests on Parquet.

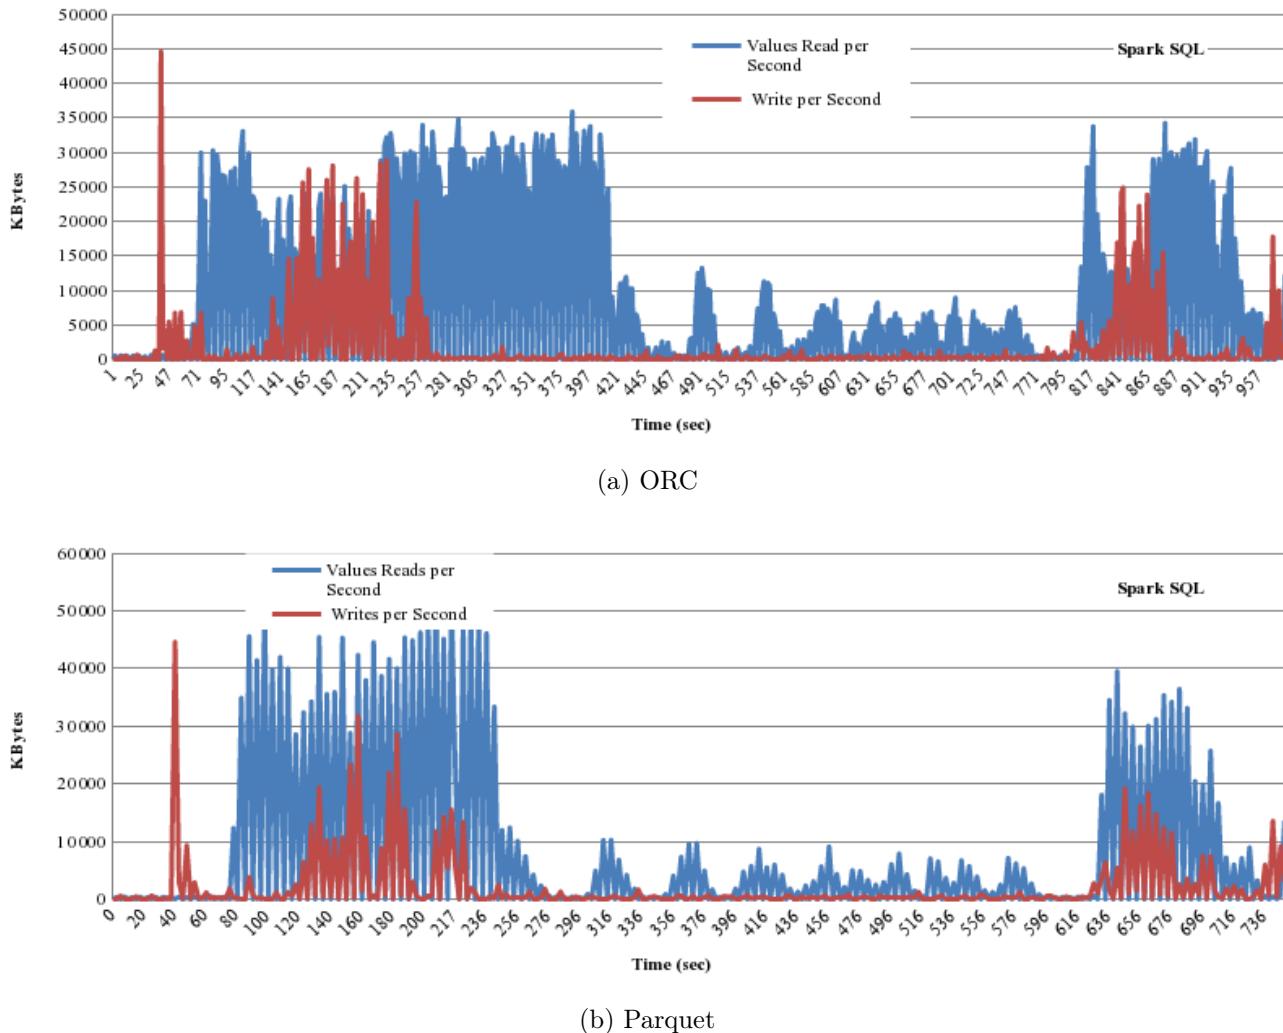


Figure 3: Disk Bandwidth: bandwidth utilization on Parquet.

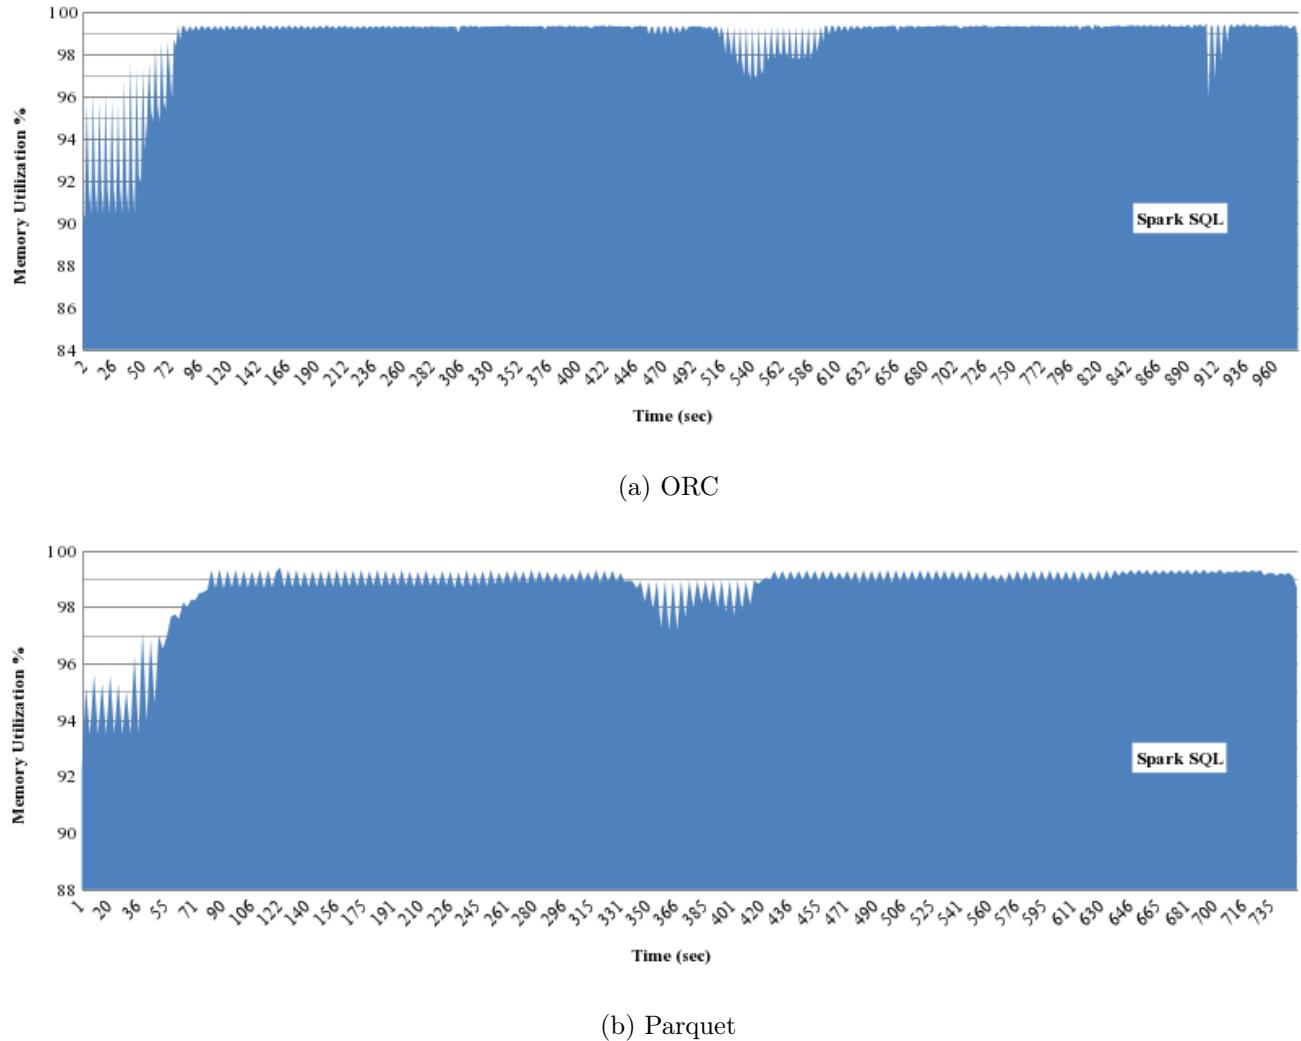


Figure 4: Free Memory.

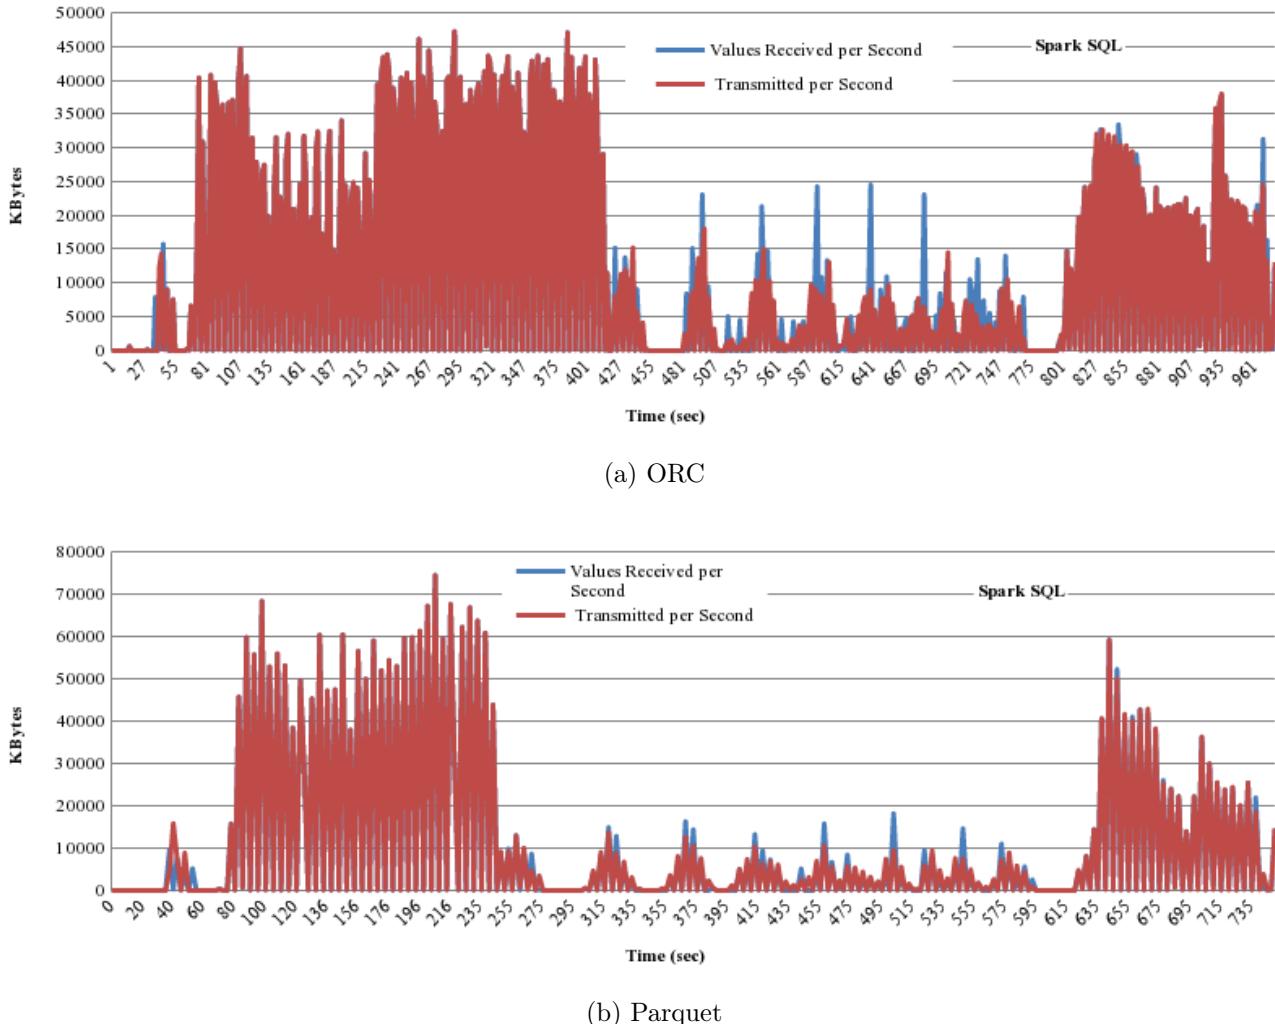


Figure 5: Network Utilization: More or less corresponding to DiskIO.

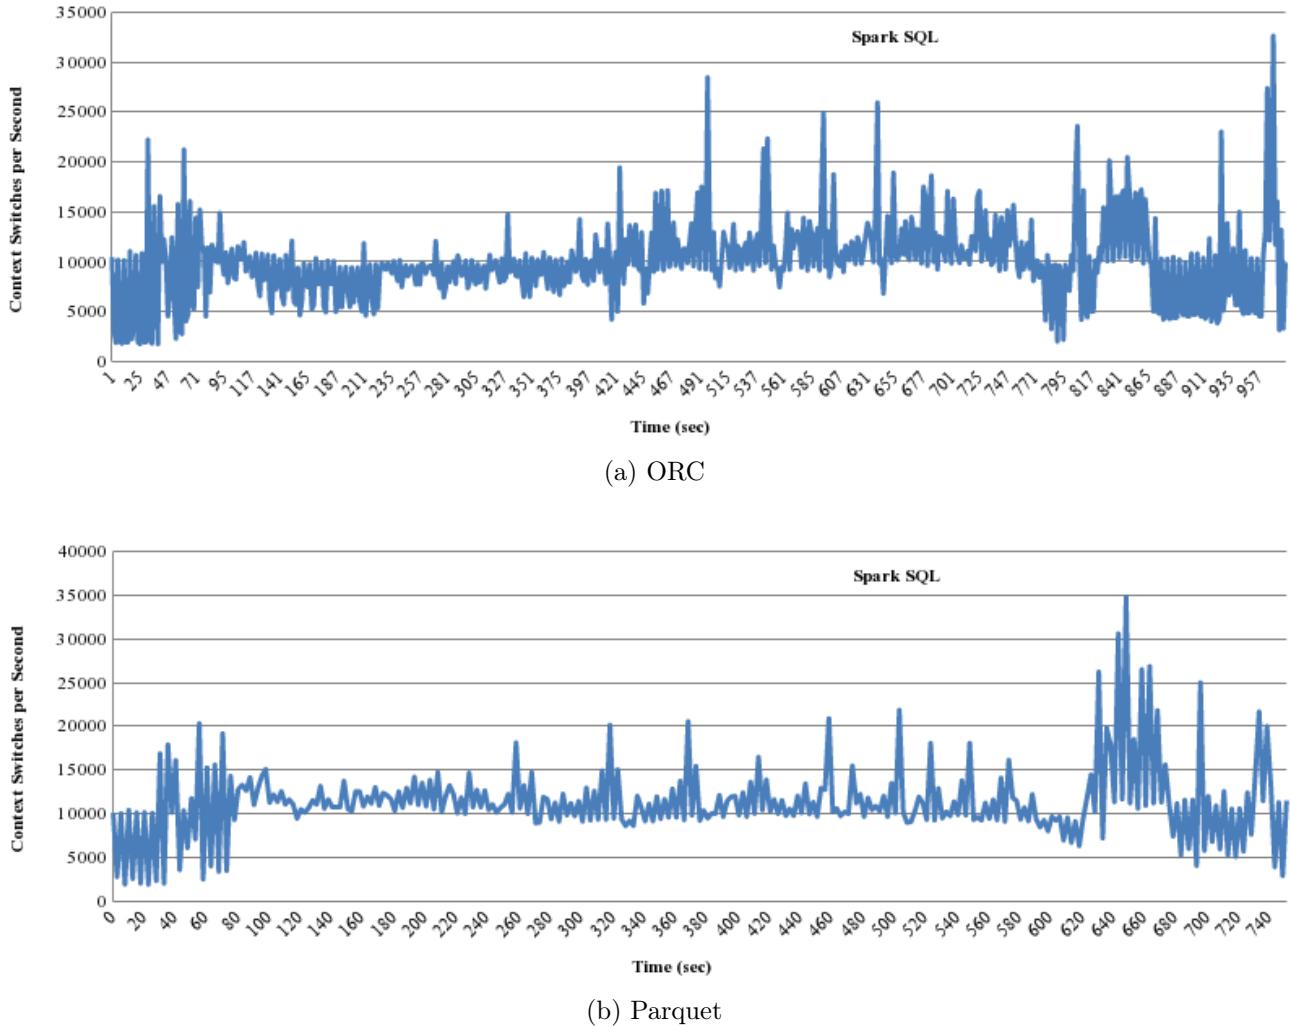


Figure 6: VMSTAT Context Switches

1.9 Snappy - ORC vs. Parquet

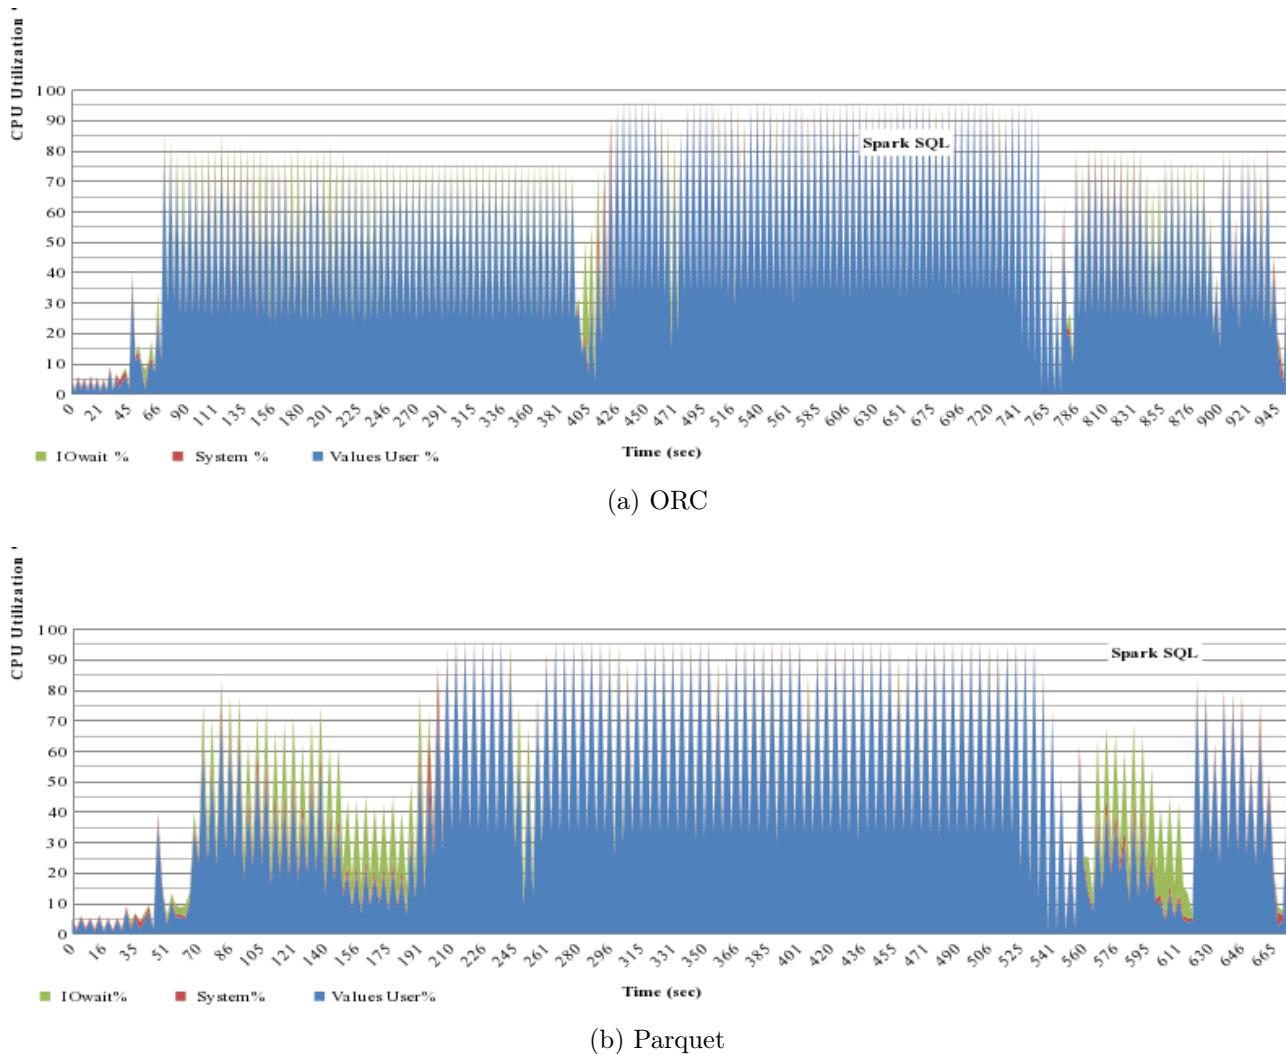


Figure 7: CPU: Low util on parquet (60%), short time for phase 1.

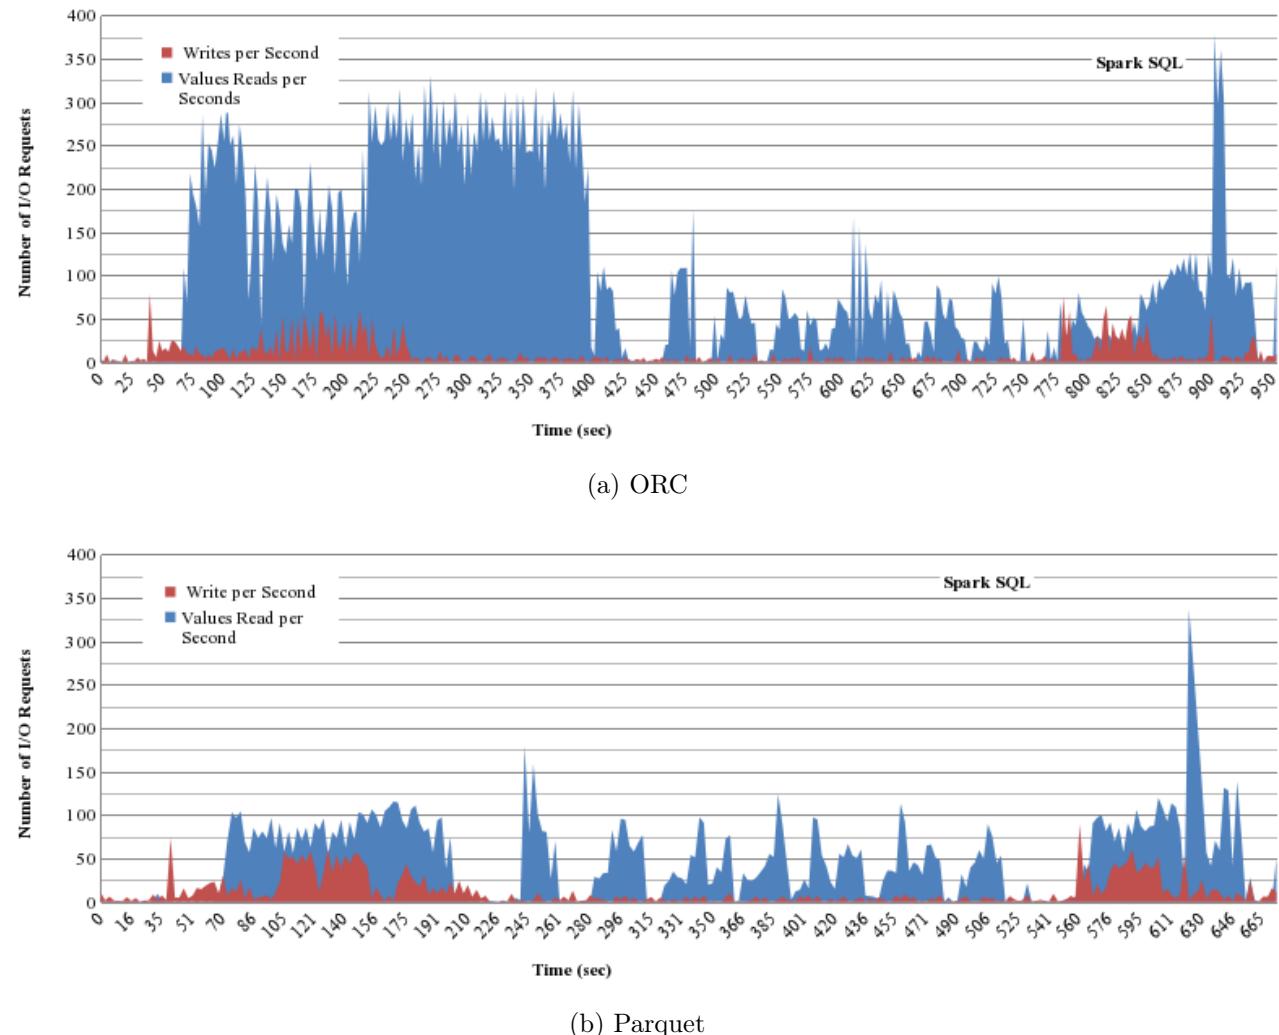


Figure 8: Disk Requests: Lower number of requests on Parquet.

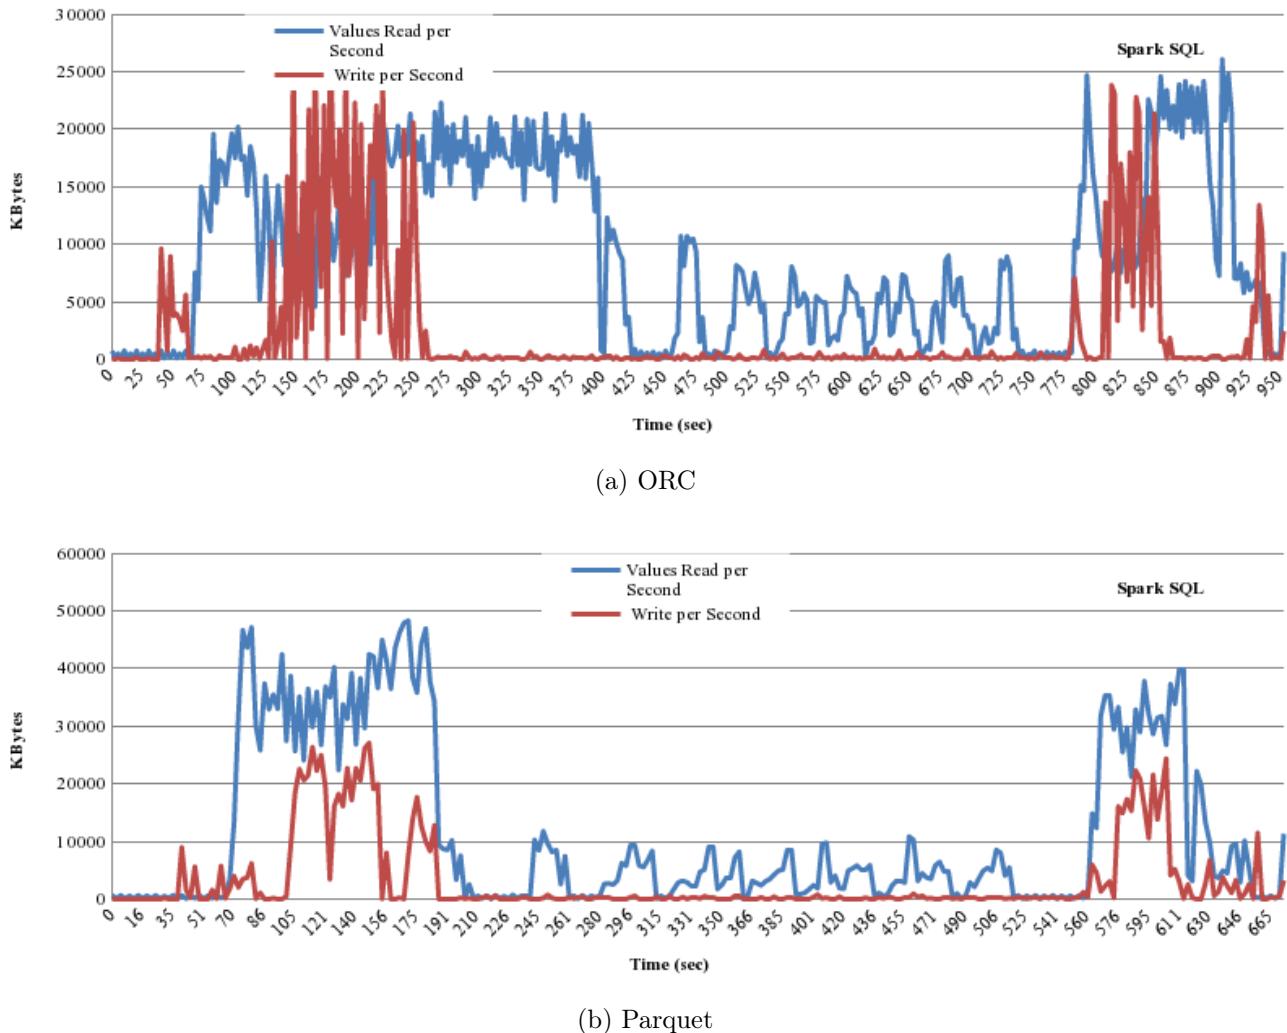


Figure 9: Disk Bandwidth: higher BW utilization on Parquet.

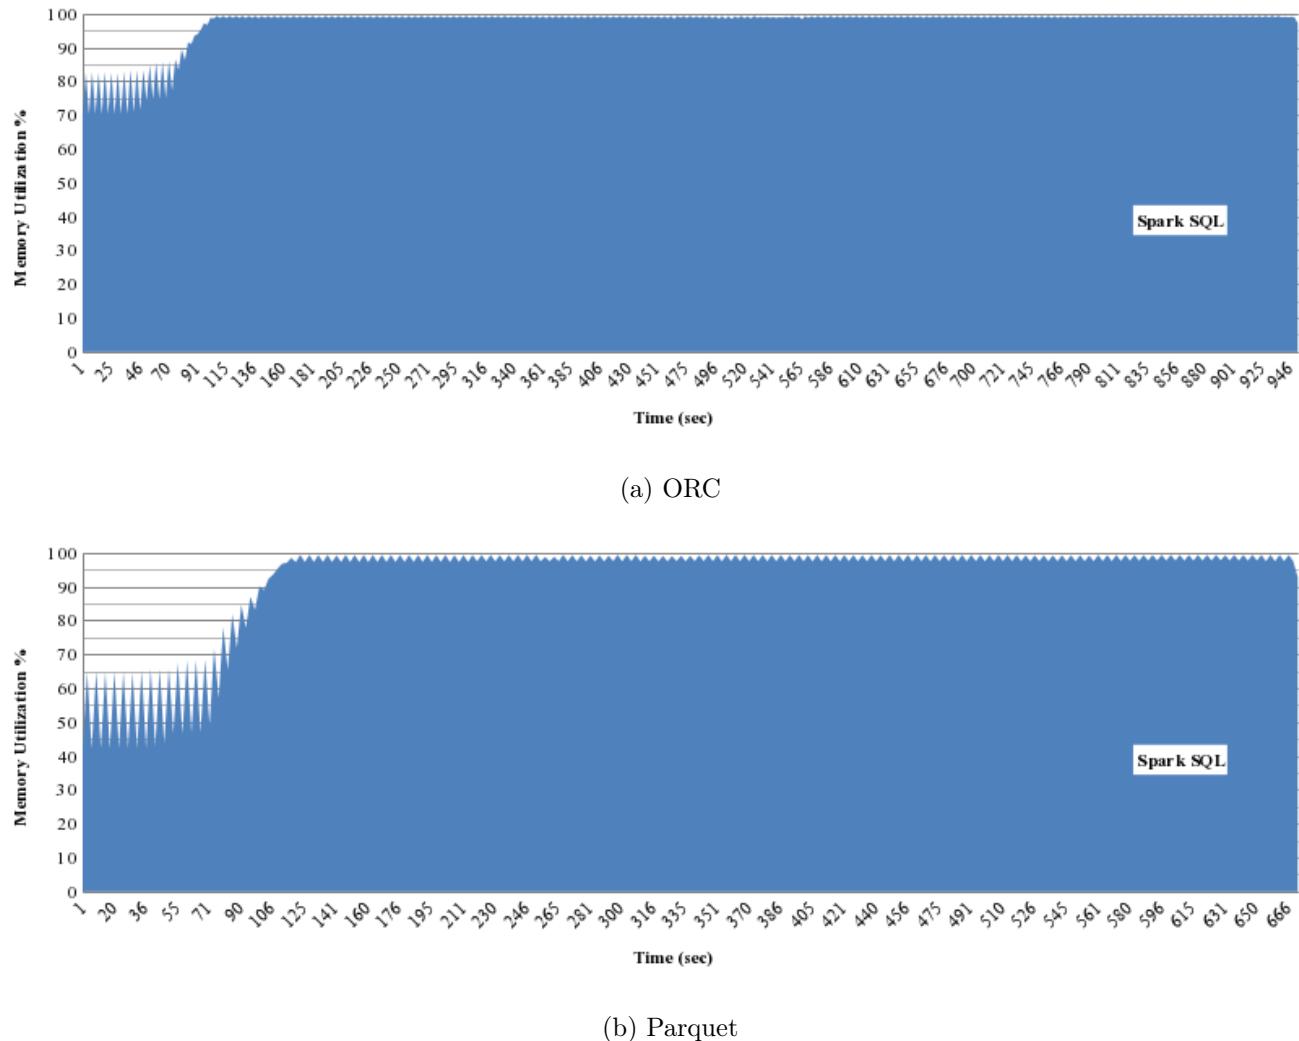


Figure 10: Free Memory.

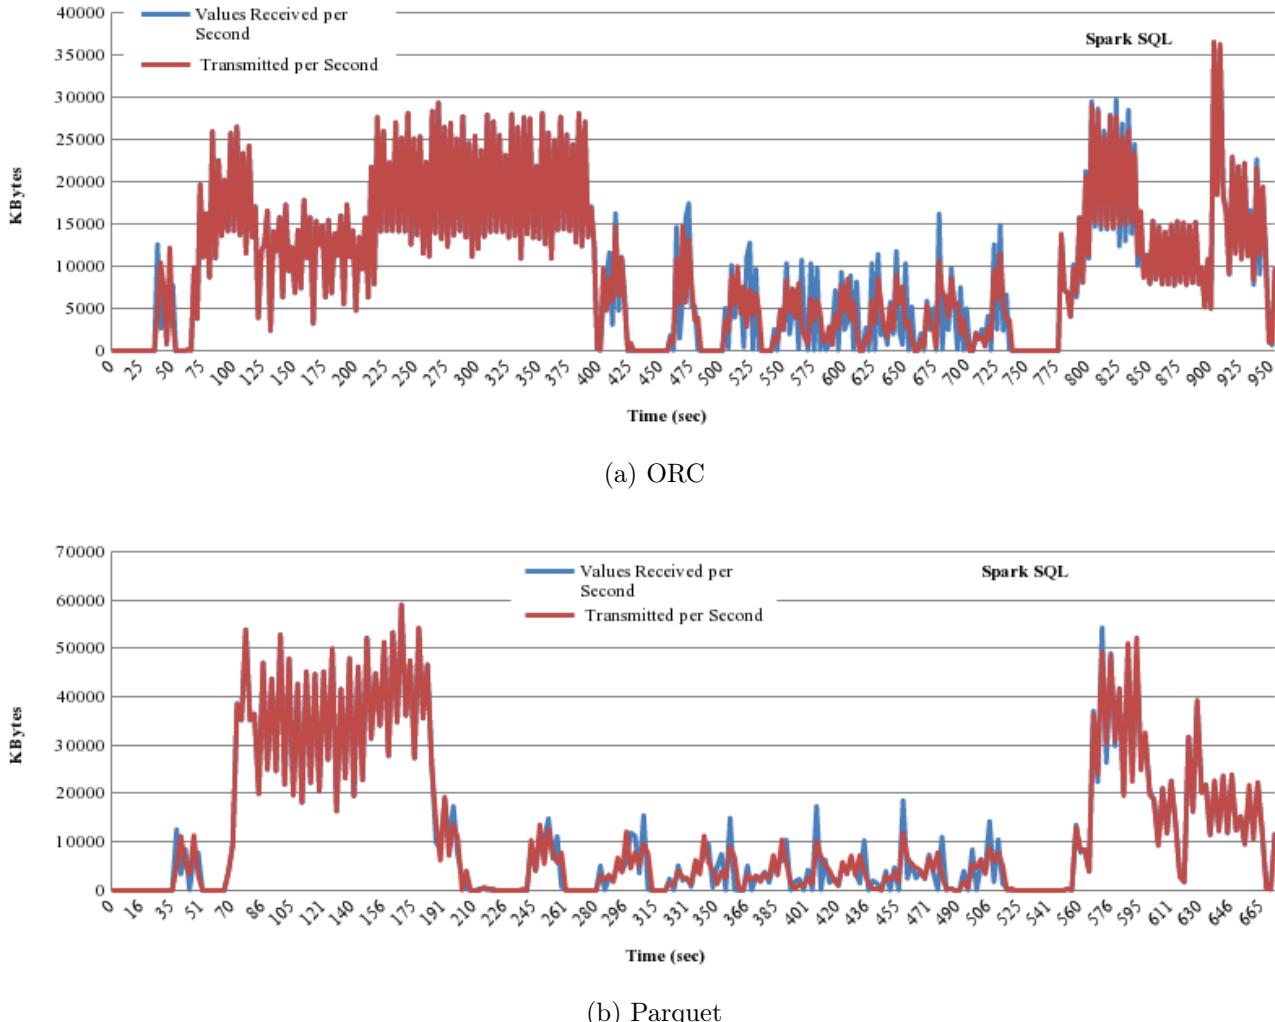


Figure 11: Network Utilization: more or less corresponding to DiskIO BW.

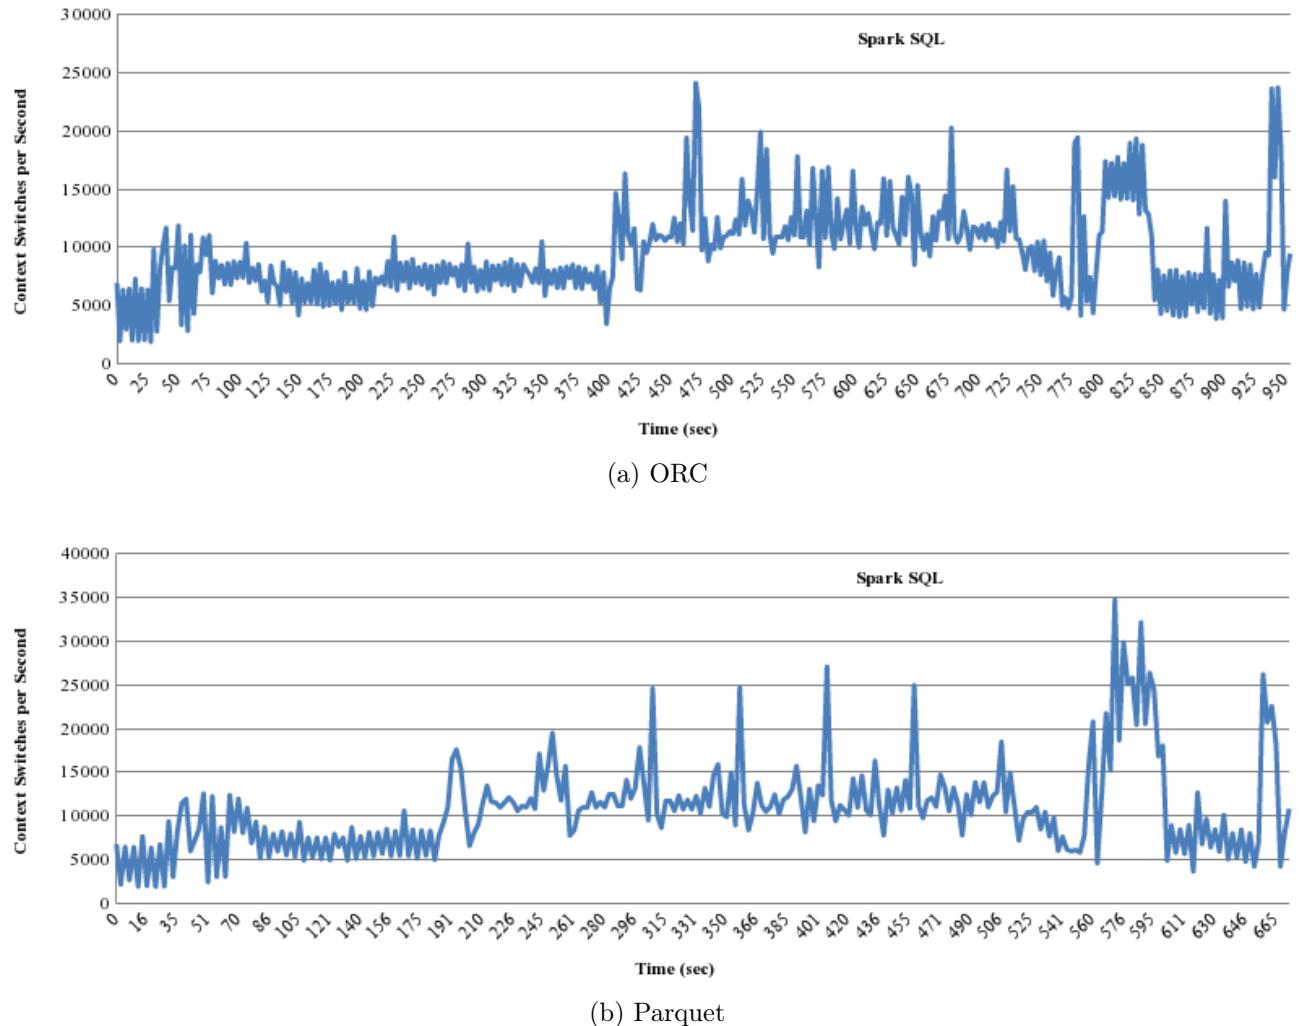


Figure 12: VMSTAT Context Switches

1.10 ORC - NoCompression vs. Snappy

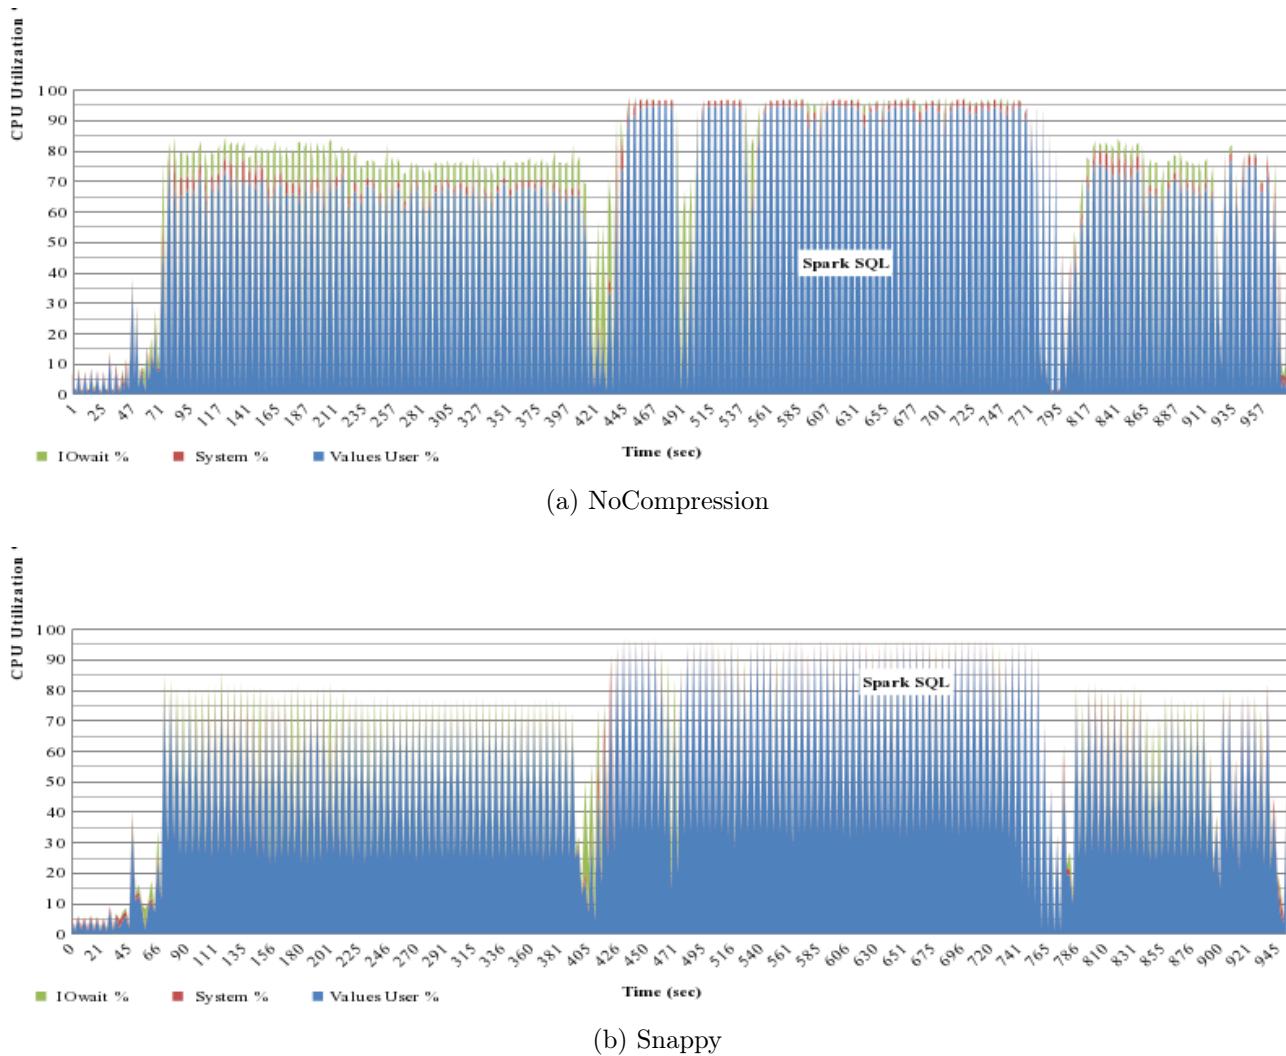


Figure 13: CPU: Same behavior.

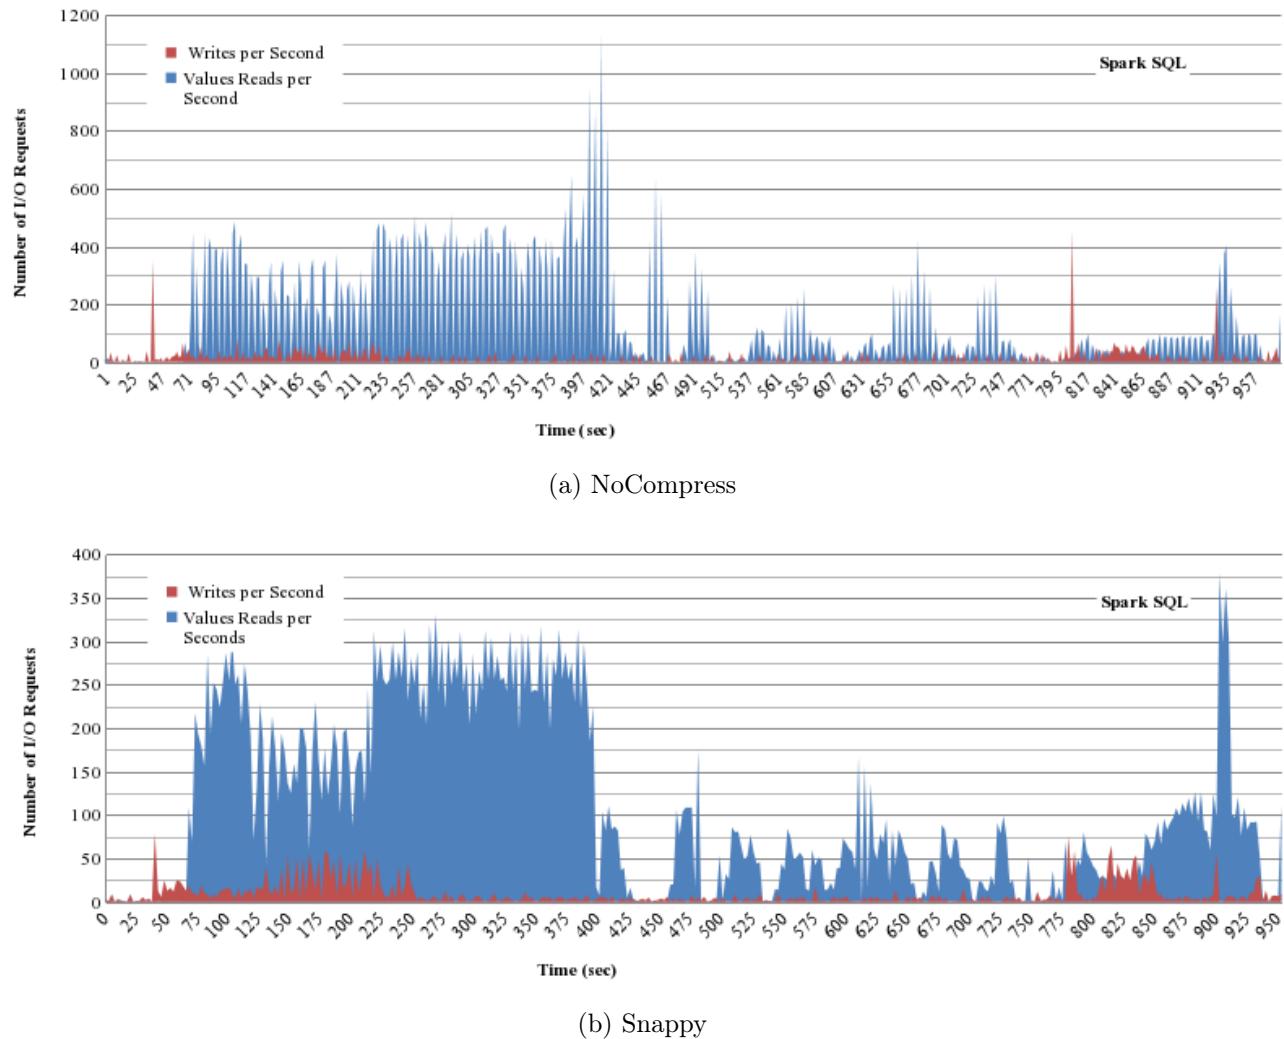


Figure 14: Disk Requests: Less requests with Snappy (good).

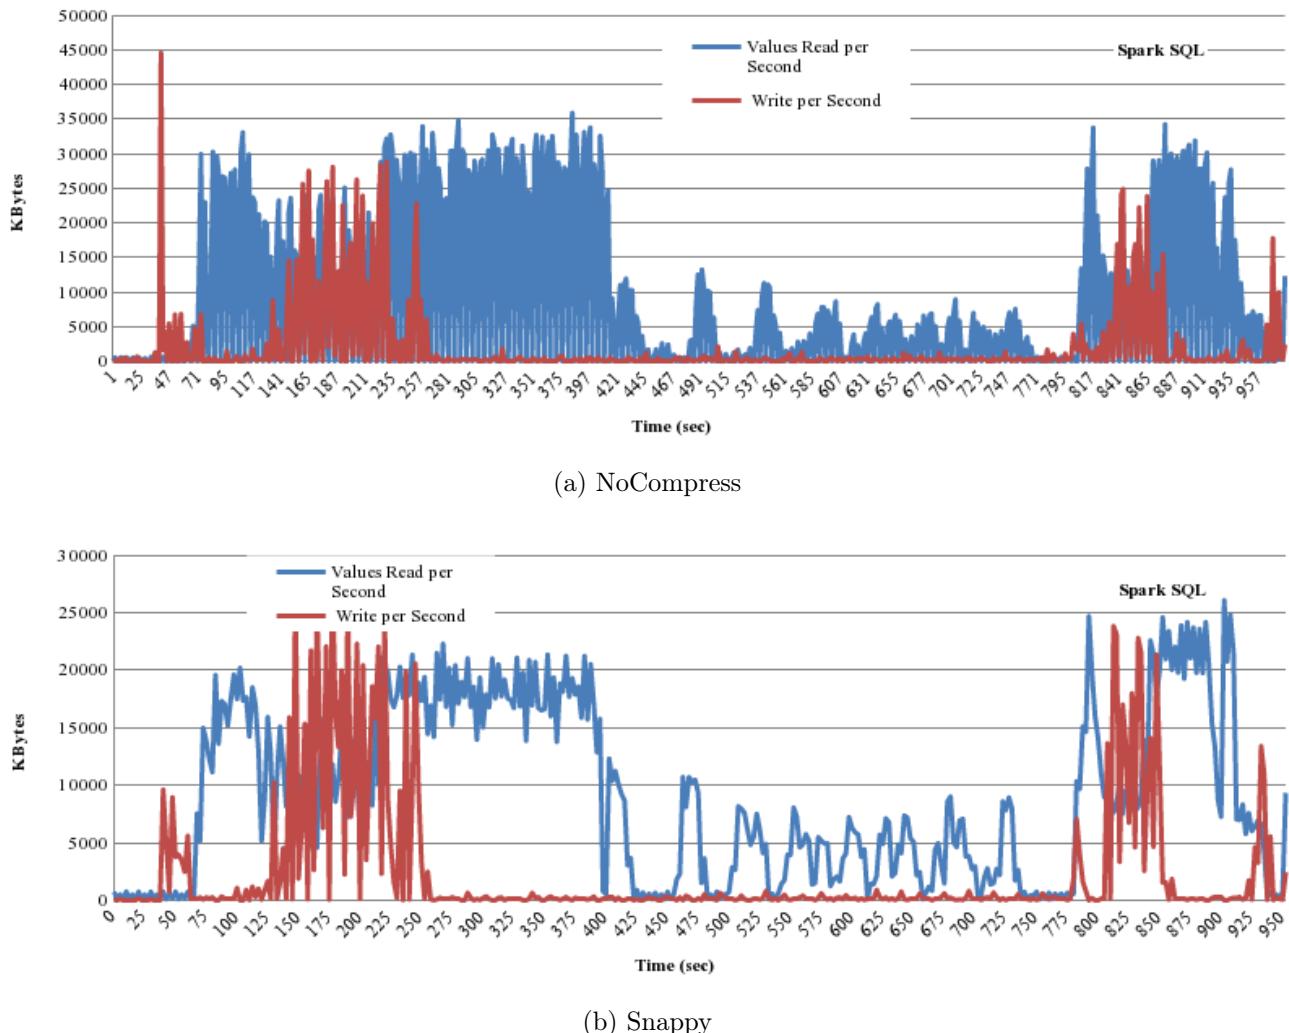


Figure 15: Disk Bandwidth: less bw utilization on snappy; goes up and falls to zero on no compression.

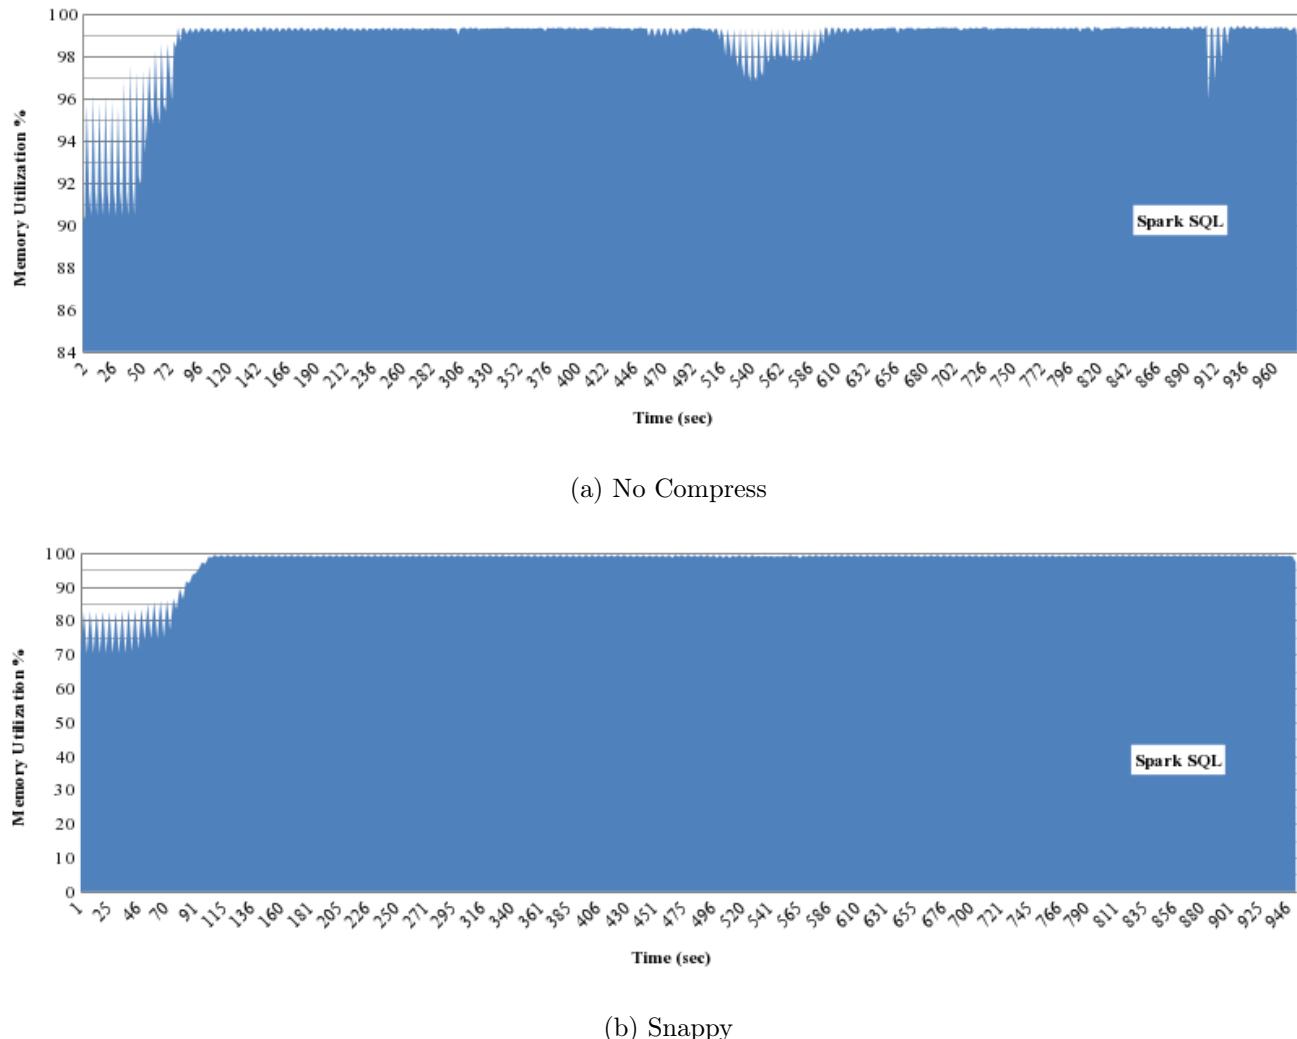


Figure 16: Free Memory.

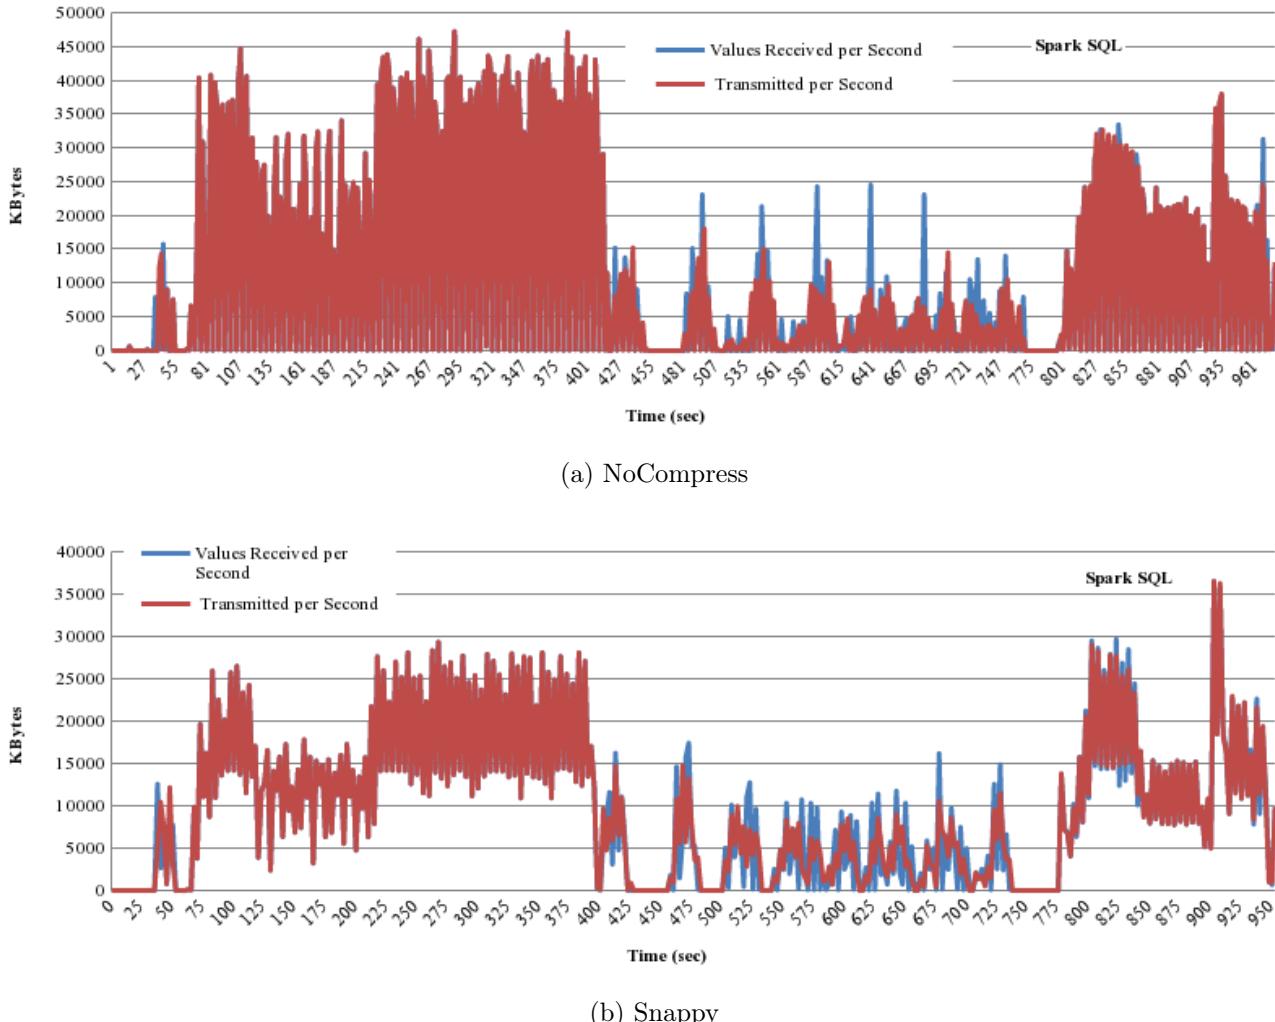


Figure 17: Network Utilization: Follows DiskIO BW.

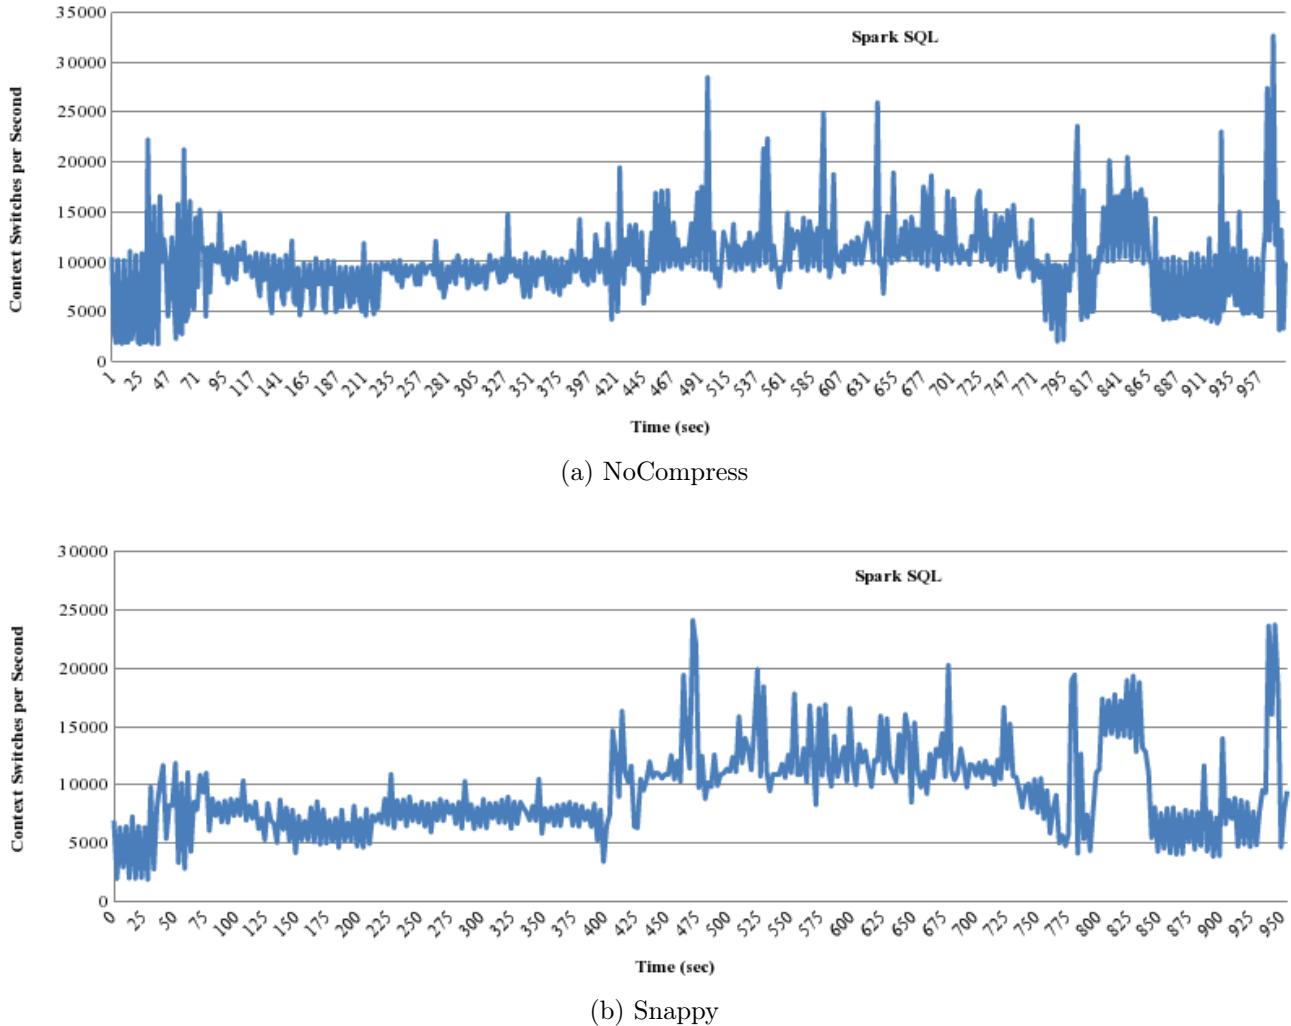
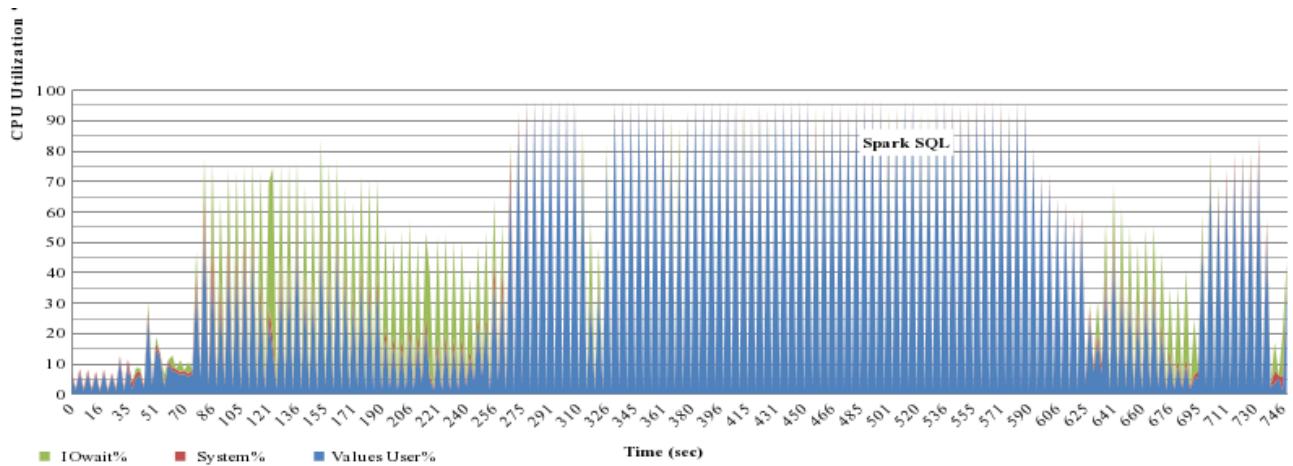
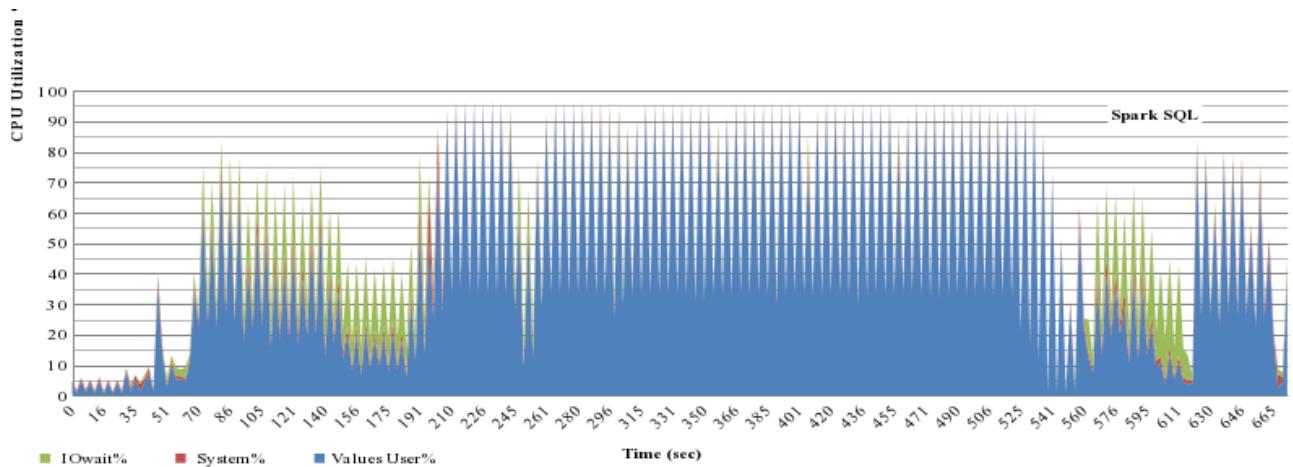


Figure 18: VMSTAT Context Switches

1.11 Parquet - NoCompression vs Snappy



(a) NoCompression



(b) Snappy

Figure 19: CPU: Same behavior (phase 1 on snappy takes less time).

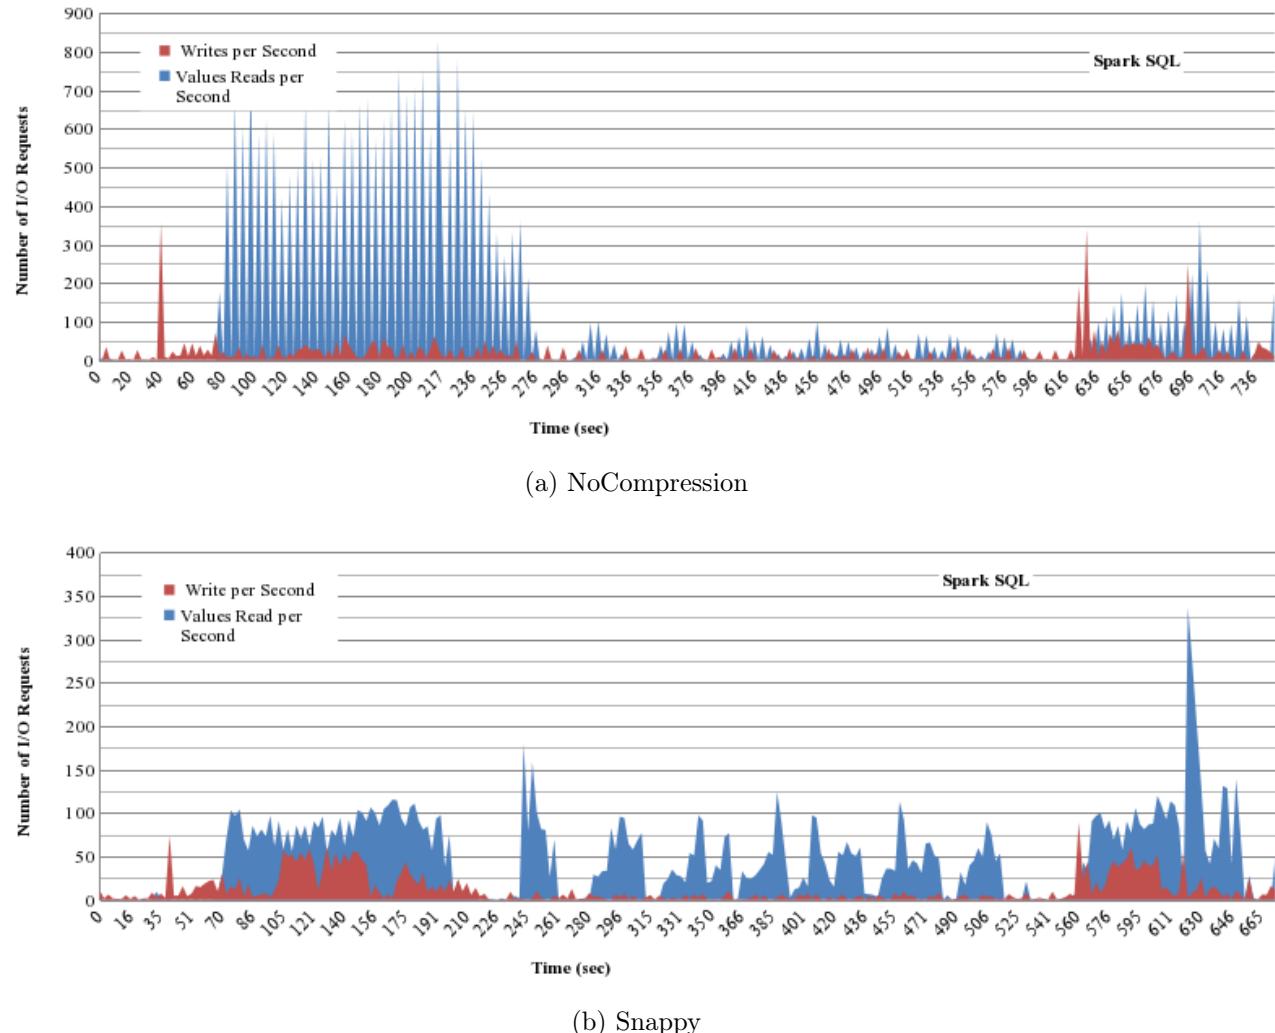


Figure 20: Disk Requests: So much less requests with Snappy!.

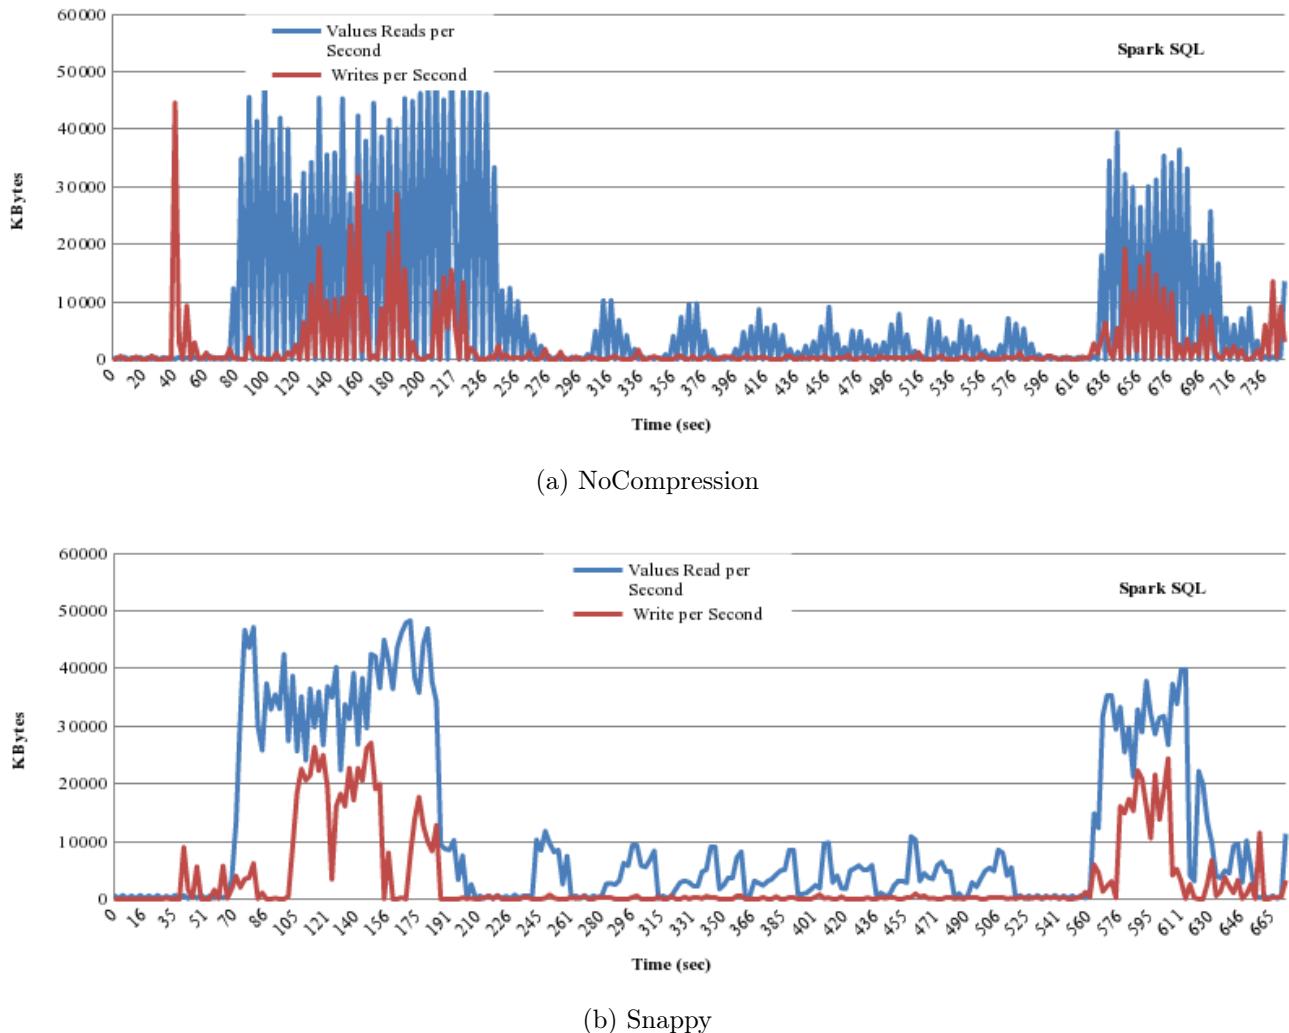


Figure 21: Disk Bandwidth: Less BW util with Snappy.

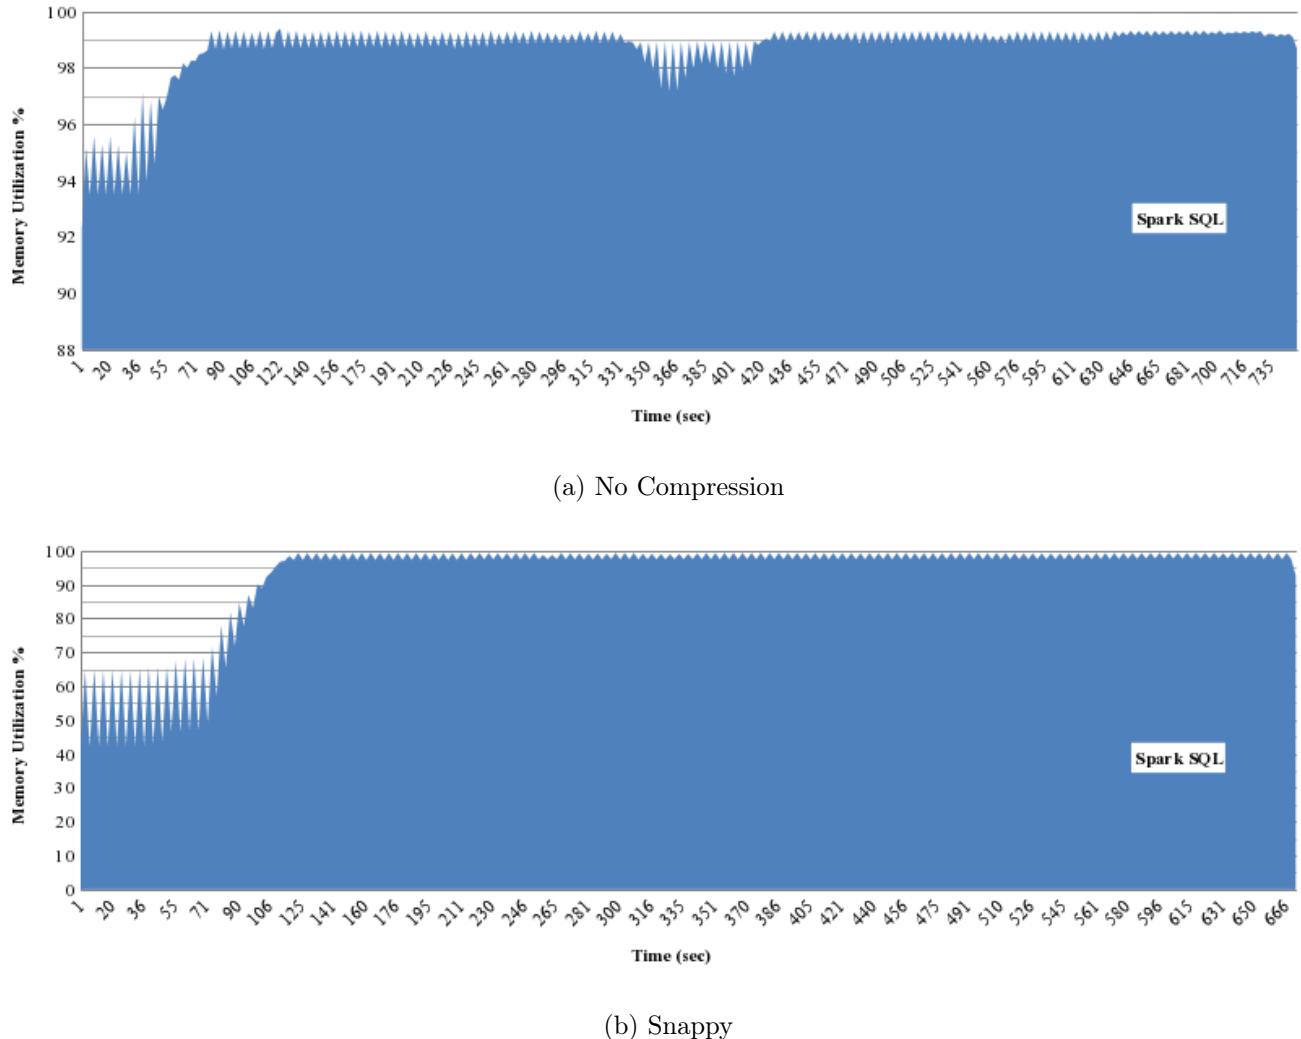


Figure 22: Free Memory.

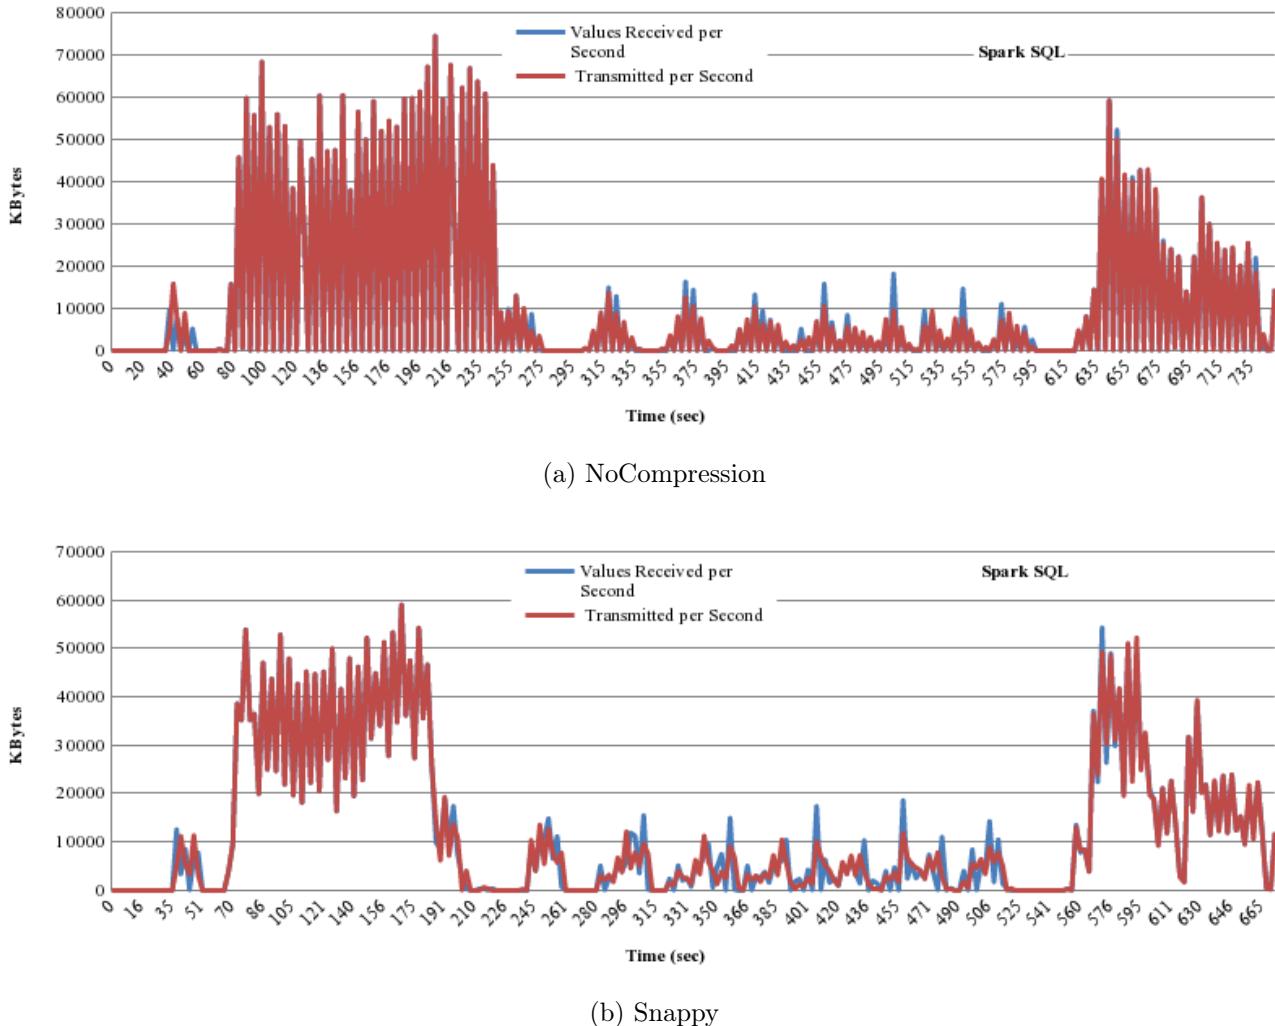


Figure 23: Network Utilization: Follows diskIO BW.

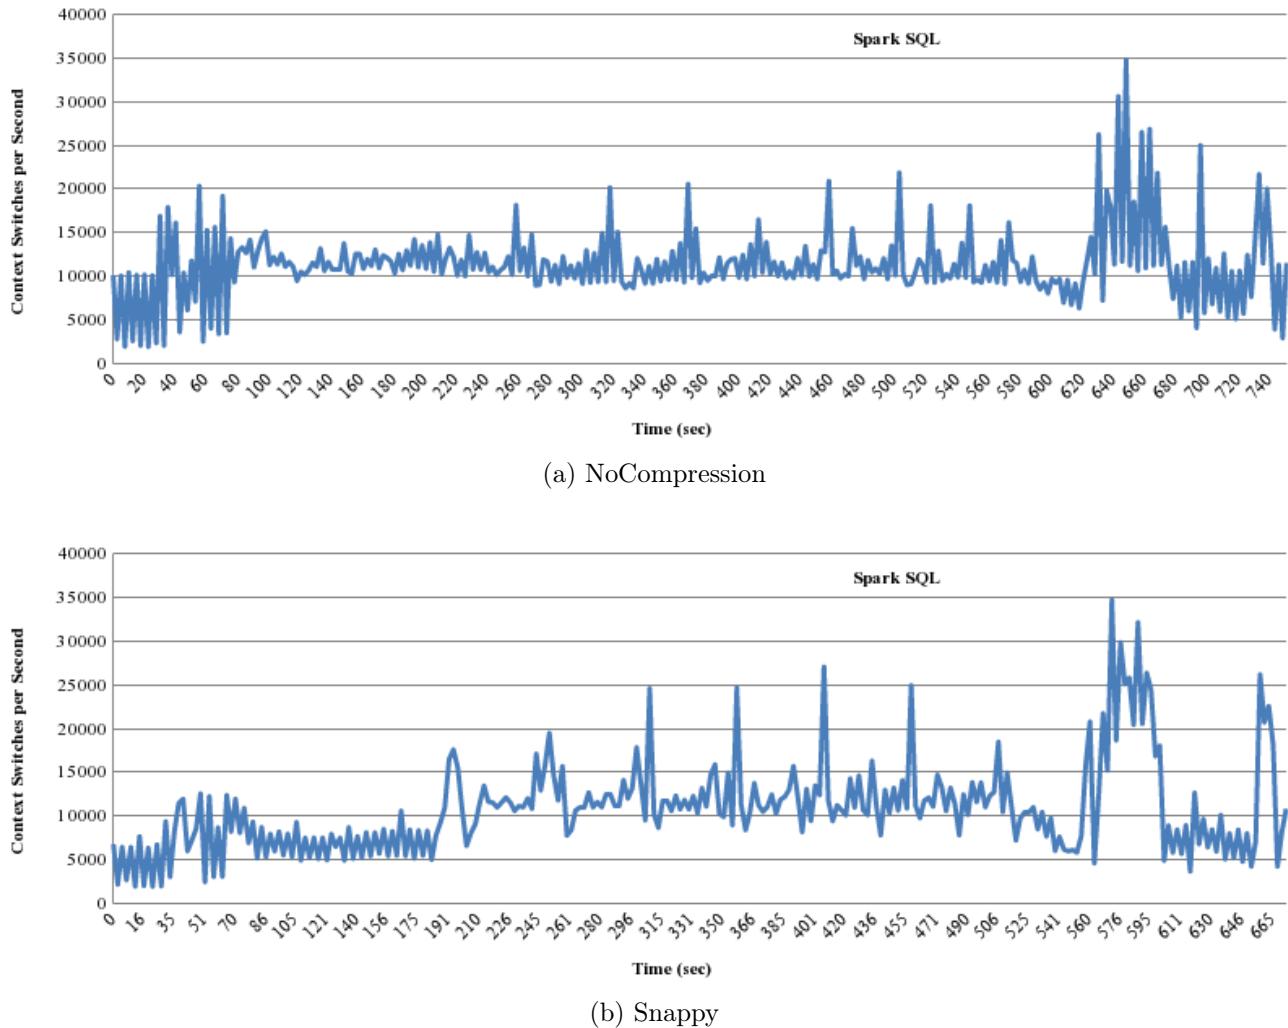


Figure 24: VMSTAT Context Switches

2 Q10

2.1 Type

HiveQL/OpenNLP

2.2 Description

For all products, extract sentences from its product reviews that contain positive or negative sentiment and display for each item the sentiment polarity of the extracted sentences (POS OR NEG) and the sentence and word in sentence leading to this classification.

2.3 Performance

	Default (sec.)		No compr. (sec.)		Snappy (sec.)	
Query	ORC	Parquet	ORC	Parquet	ORC	Parquet
Q10	1994	976	2026	1633	2070	3075

2.4 Phase 1

Selects all product reviews from “product_reviews” table and uses a UDF to extract sentiment. No temporary tables, works directly and only on file format data. Result of course is not compressed.

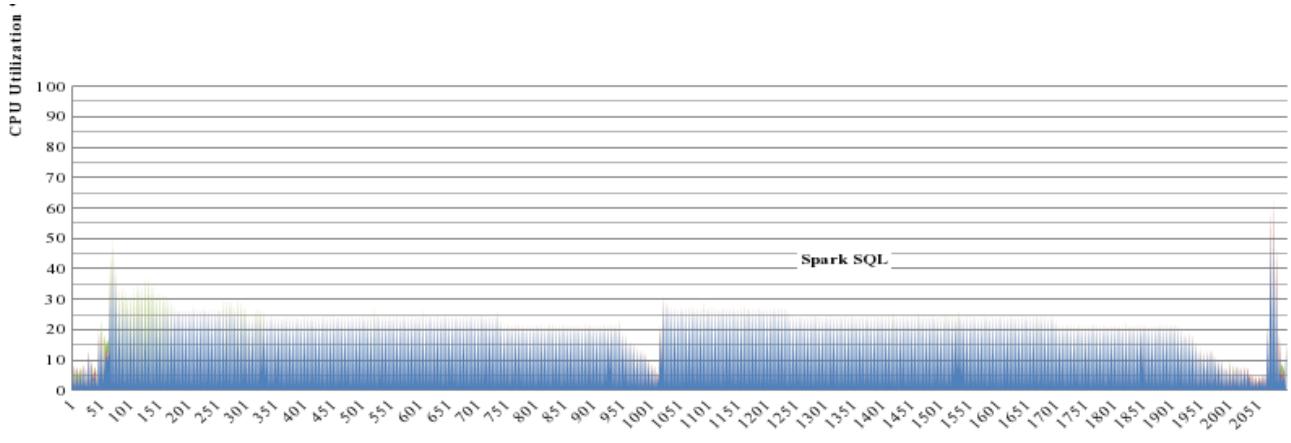
2.5 Source

GitHub: <https://git.io/vpFeK>

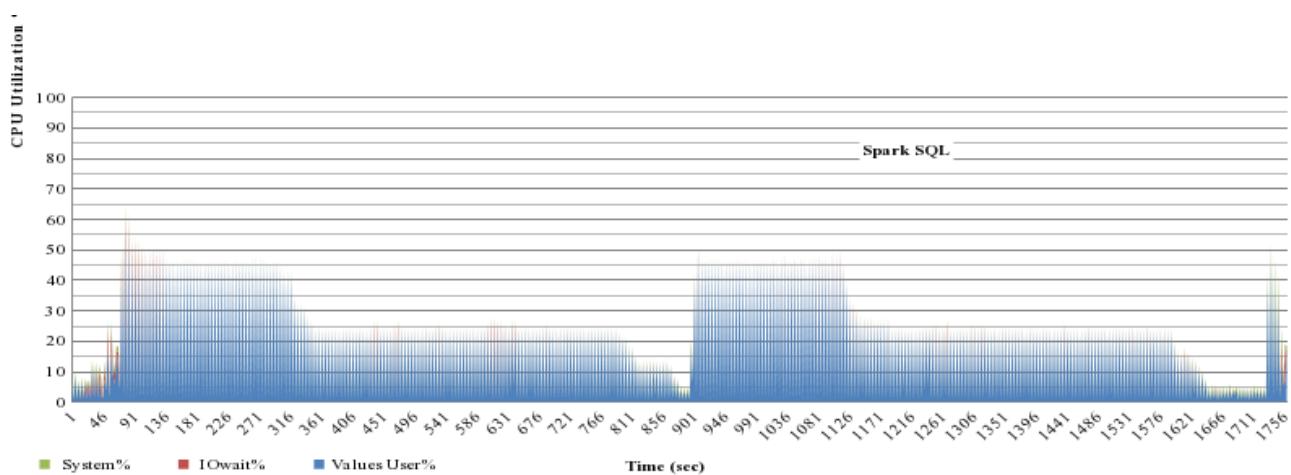
2.6 Notes

This type of queries perform better on uncompressed data. It is interesting to observe that with Snappy, ORC performs better for Q10. Check CPU and disk.

2.7 No compression - ORC vs Parquet



(a) ORC



(b) Parquet

Figure 25: CPU: Lower utilization on Parquet (max is 80%)

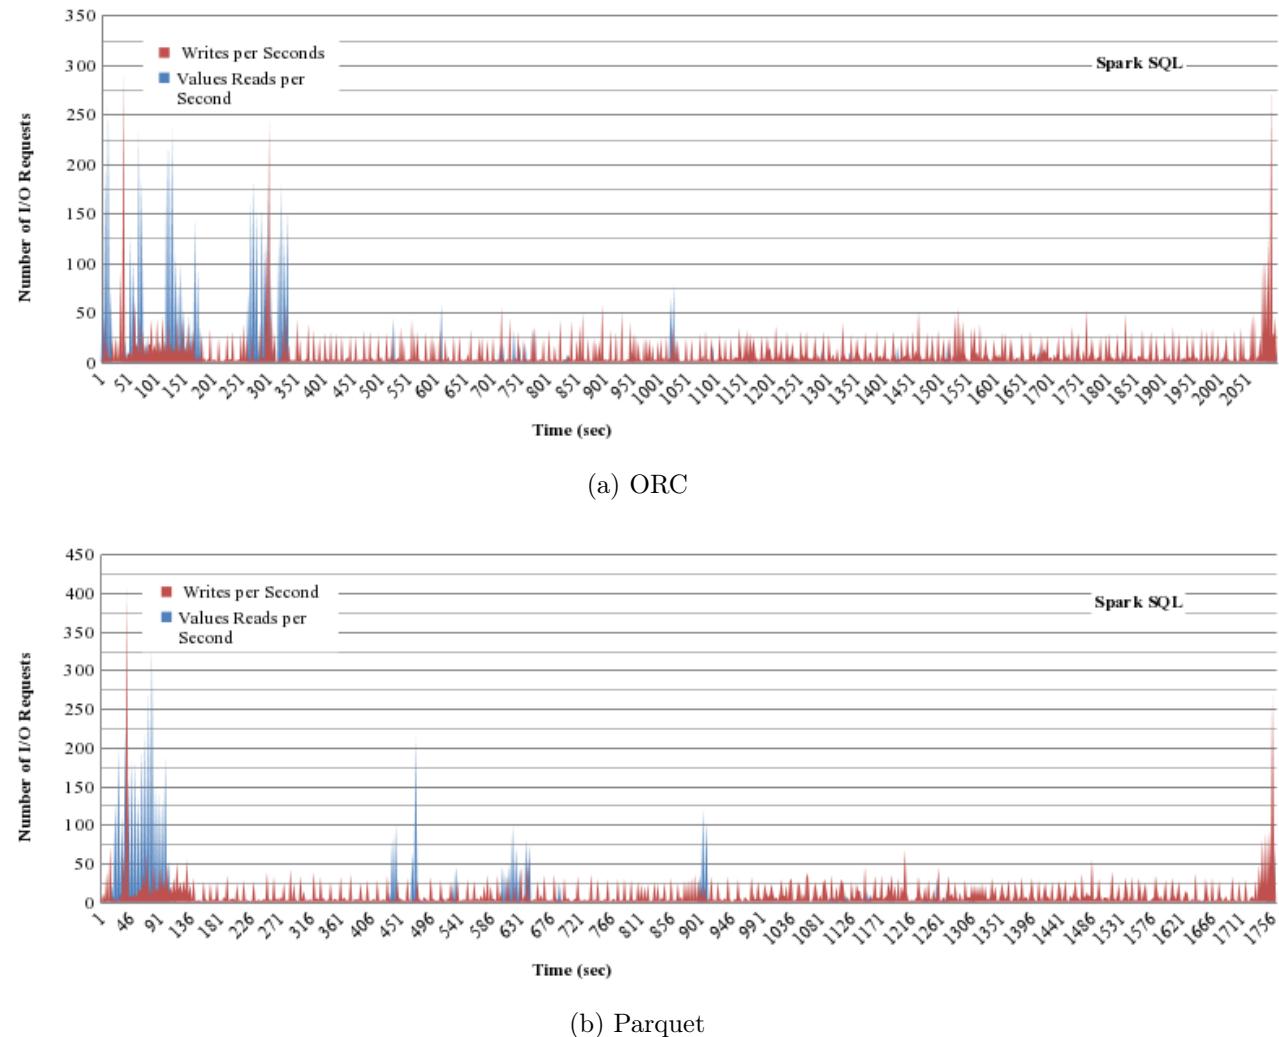


Figure 26: Disk Requests: More requests on Parquet.

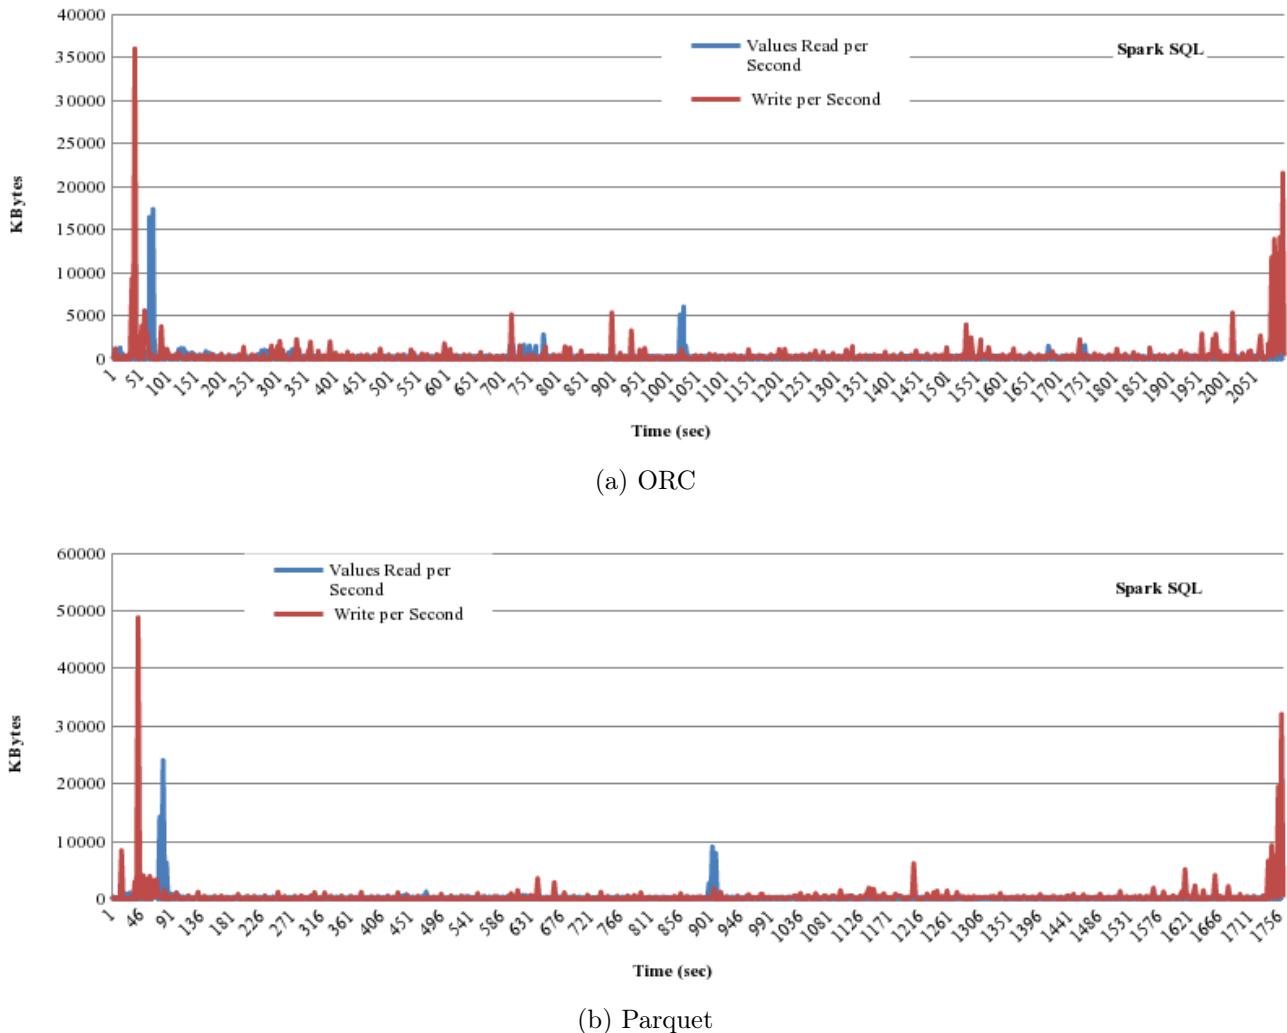


Figure 27: Disk Bandwidth: bandwidth utilization on Parquet.

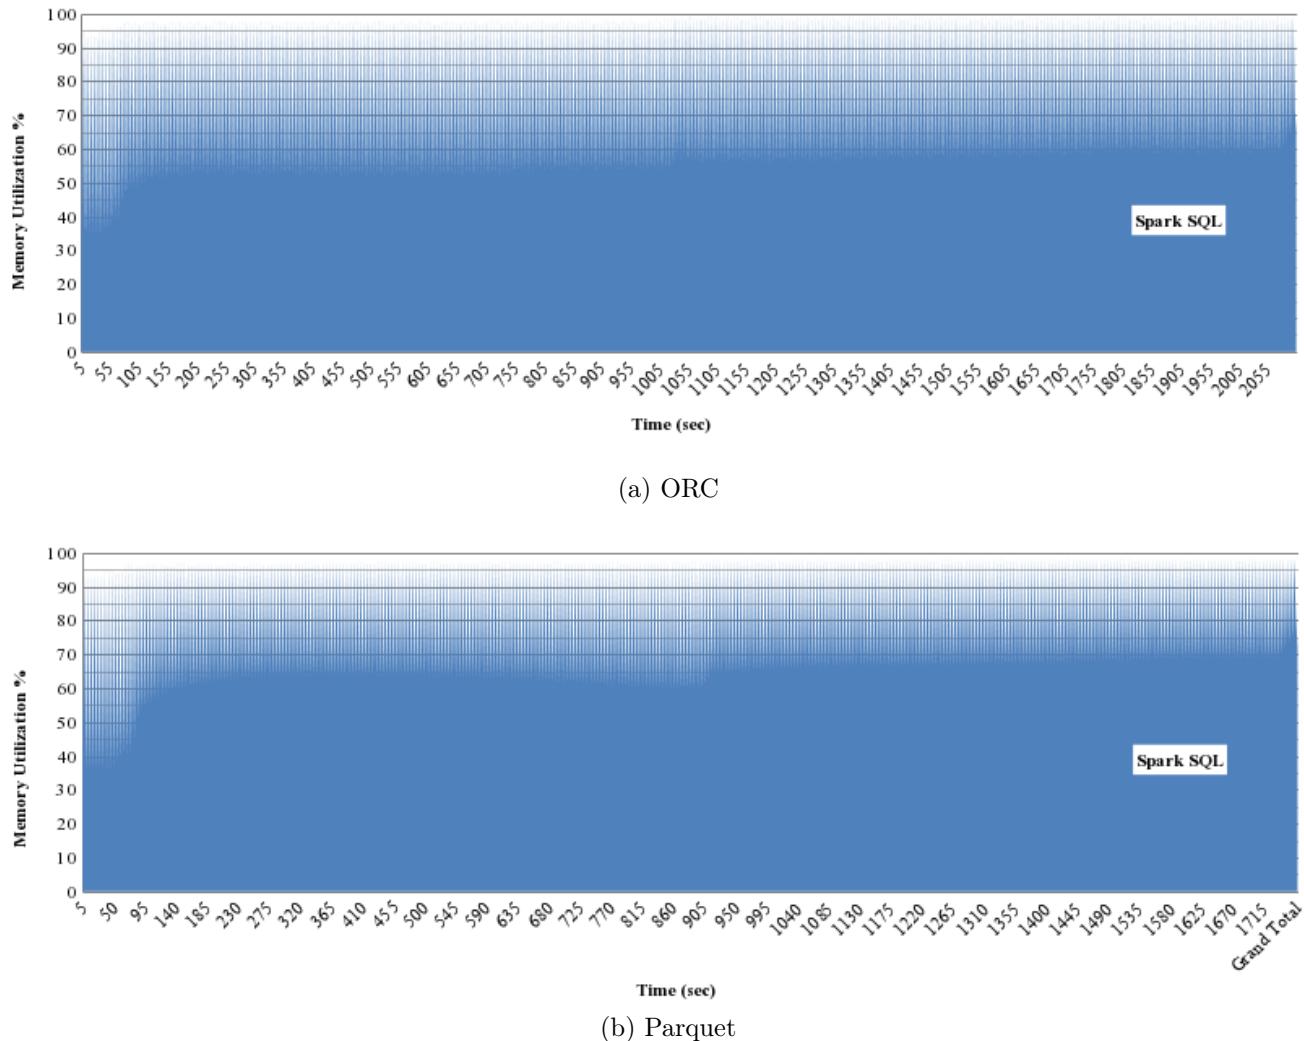


Figure 28: Free Memory.

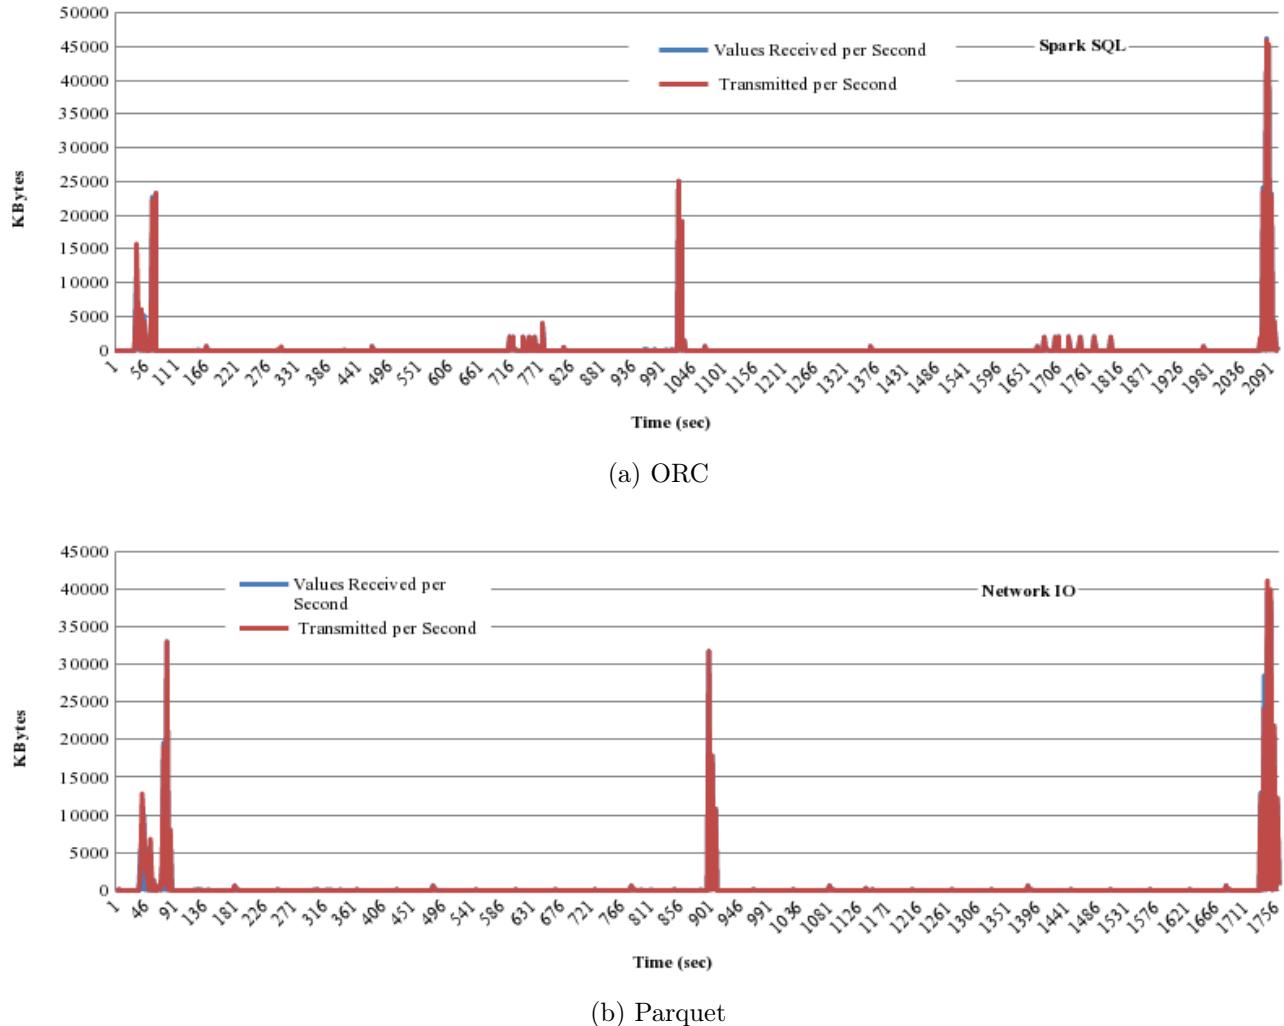


Figure 29: Network Utilization: More or less corresponding to DiskIO.

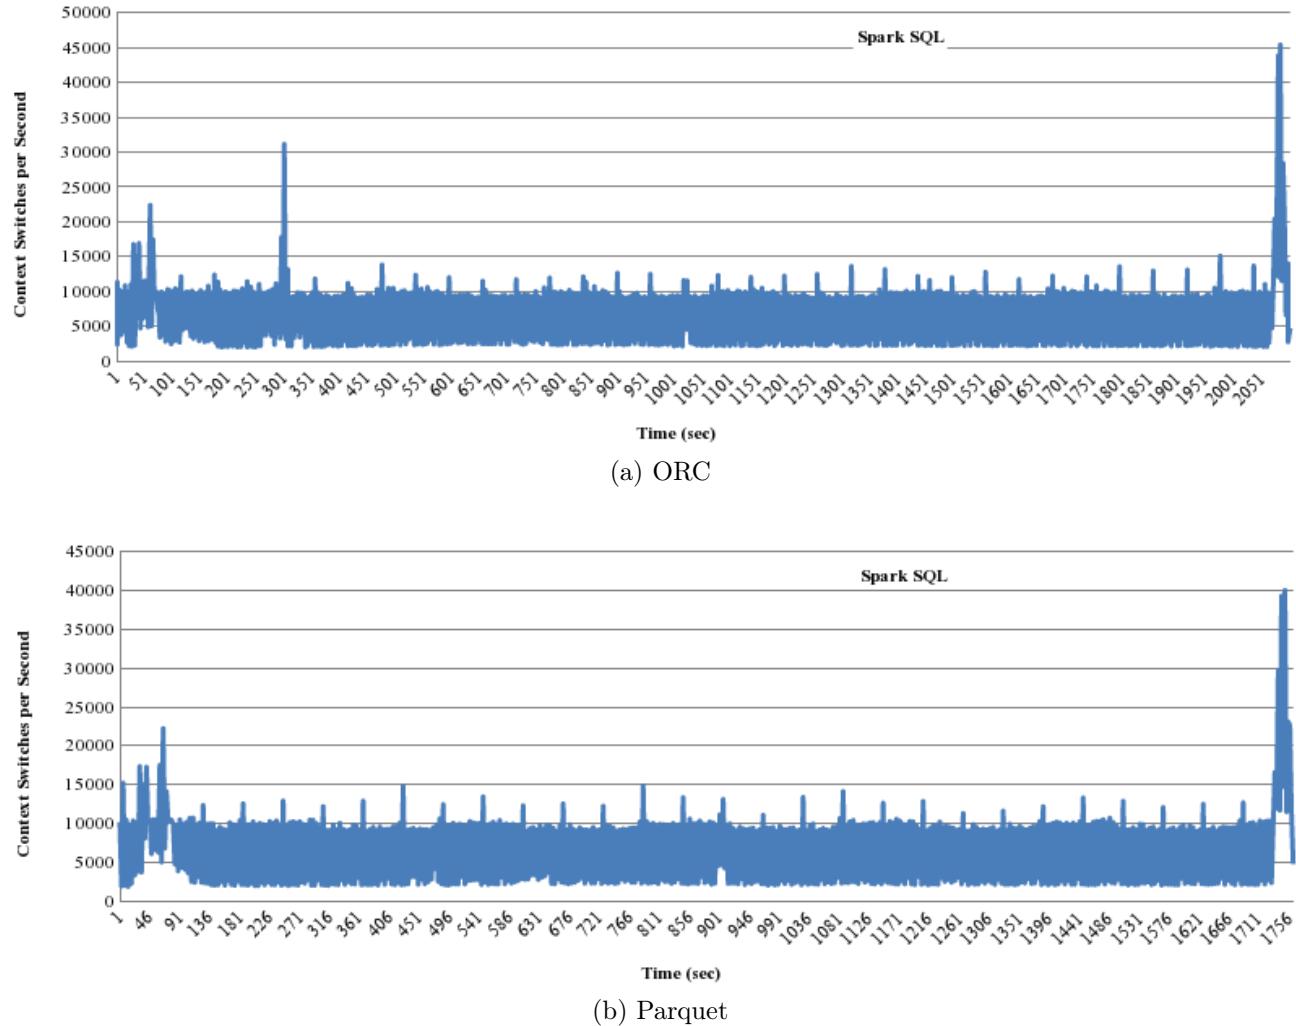


Figure 30: VMSTAT Context Switches

2.8 Snappy - ORC vs. Parquet

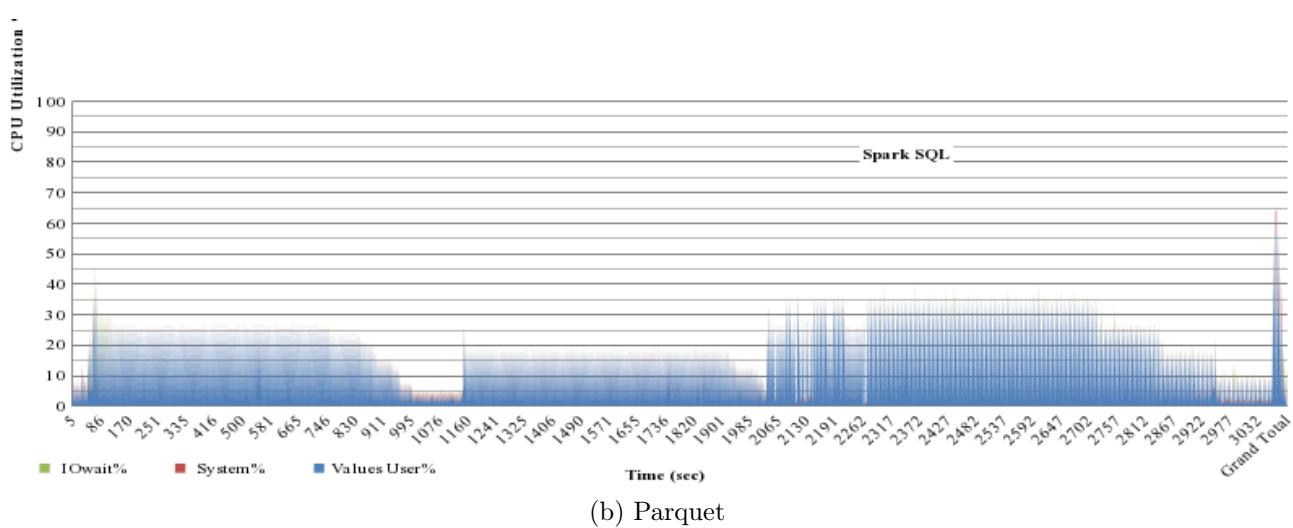
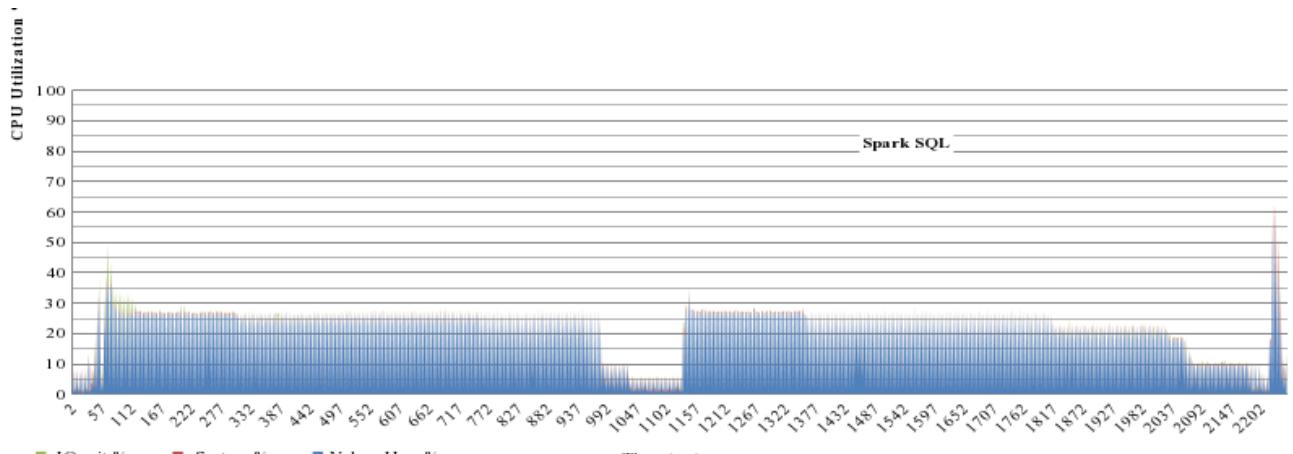


Figure 31: CPU: Low util on parquet (60%), short time for phase 1.

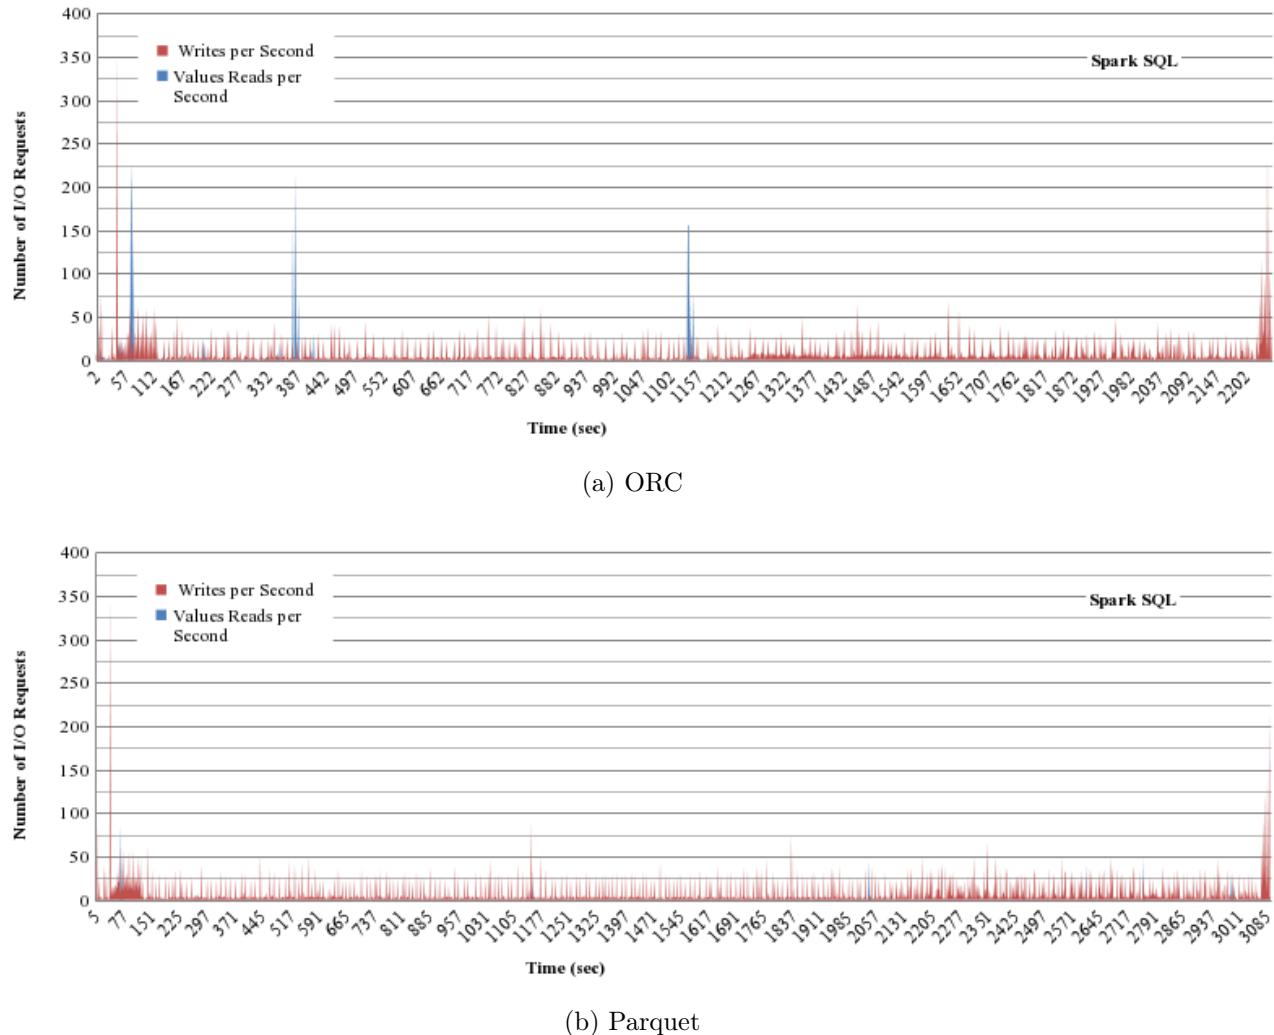


Figure 32: Disk Requests: Lower number of requests on Parquet.

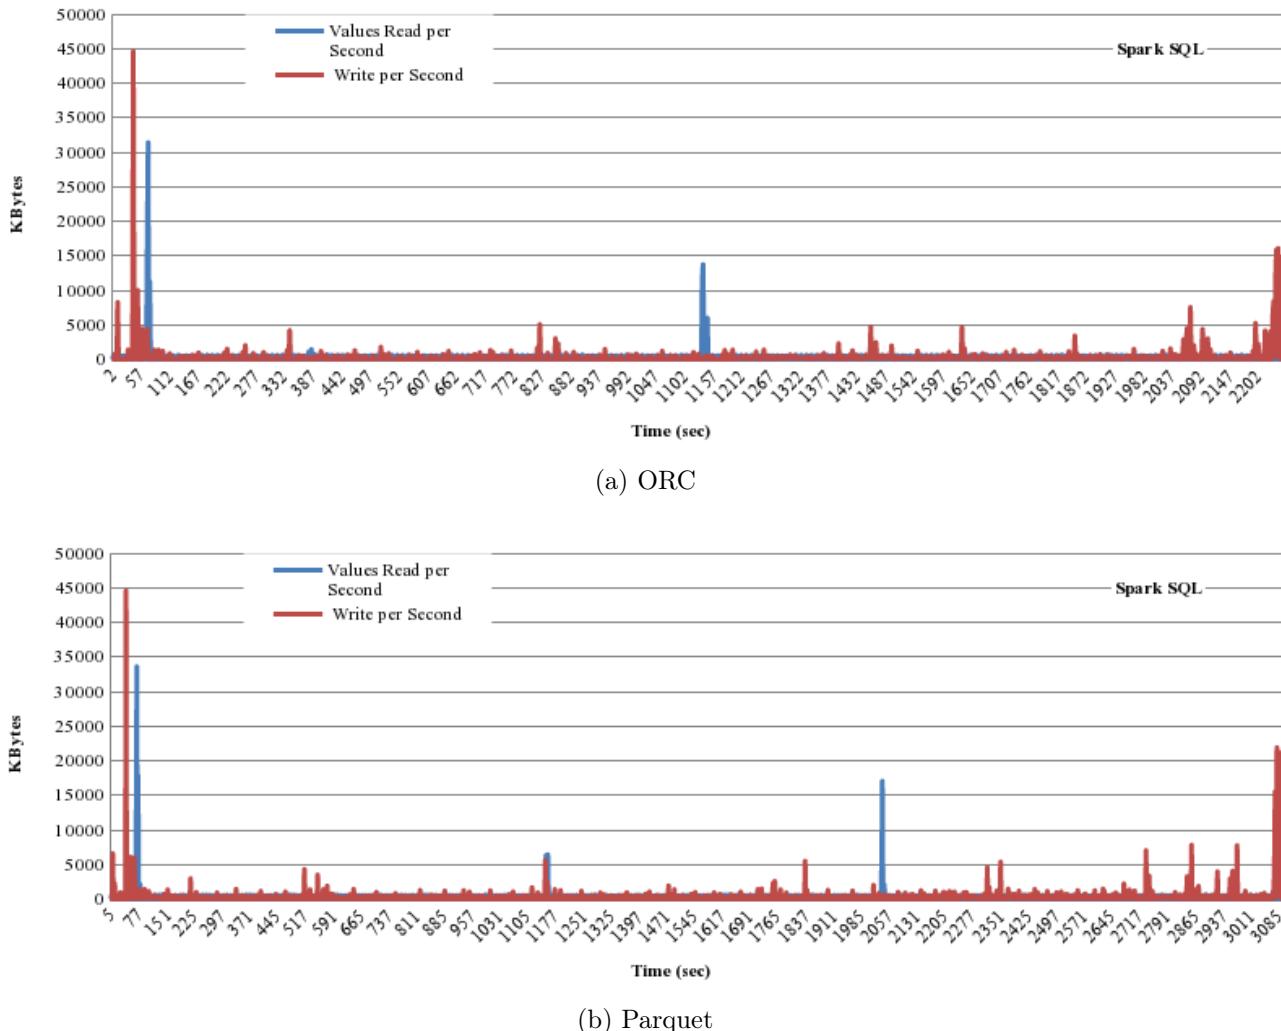


Figure 33: Disk Bandwidth: higher BW utilization on Parquet.

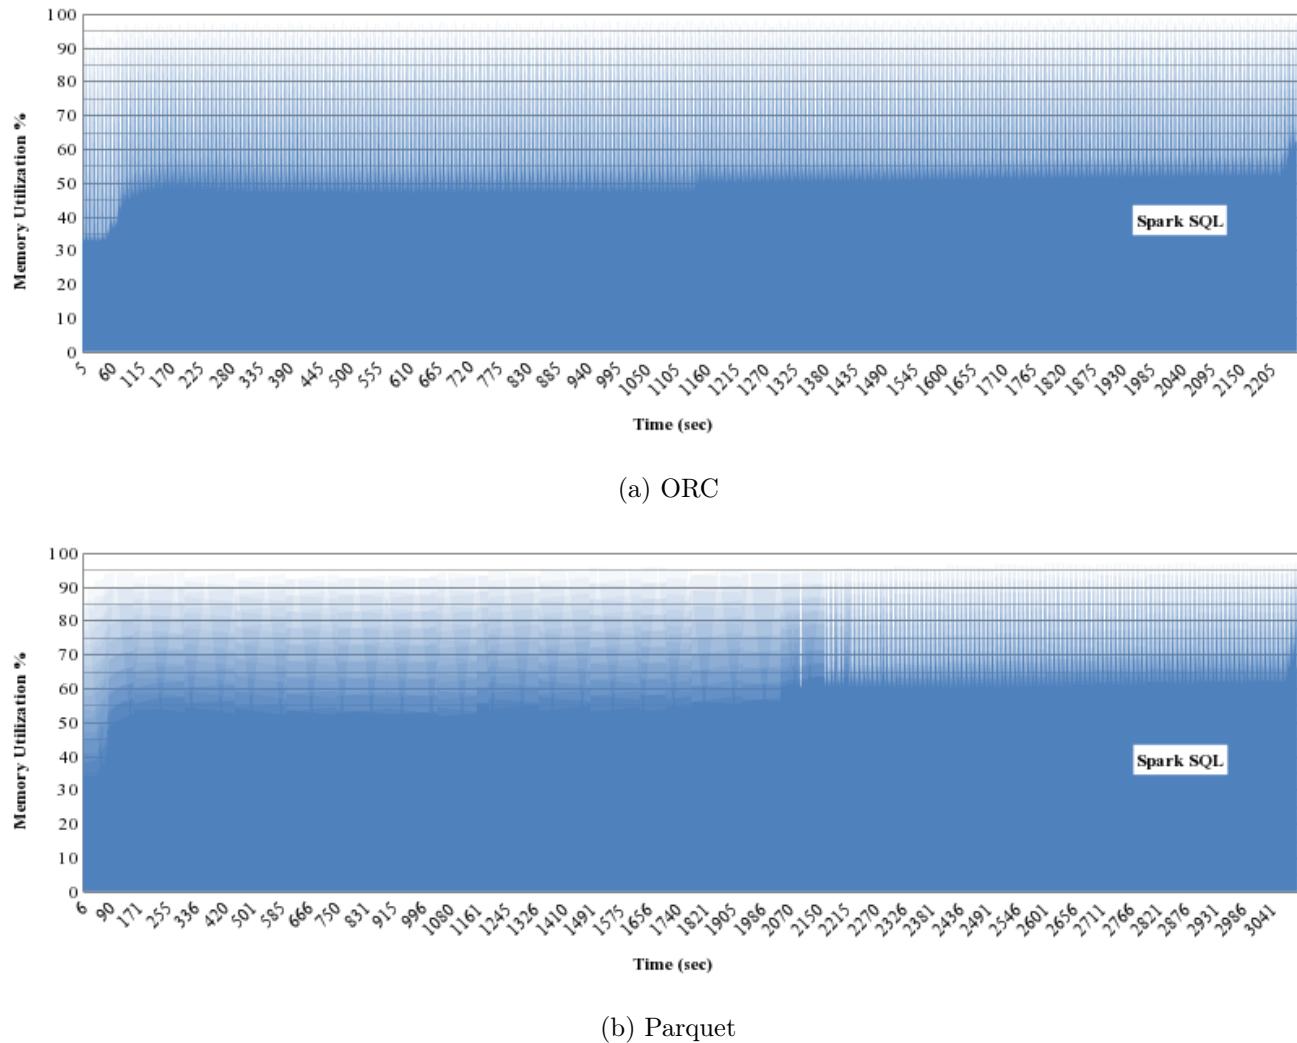


Figure 34: Free Memory.

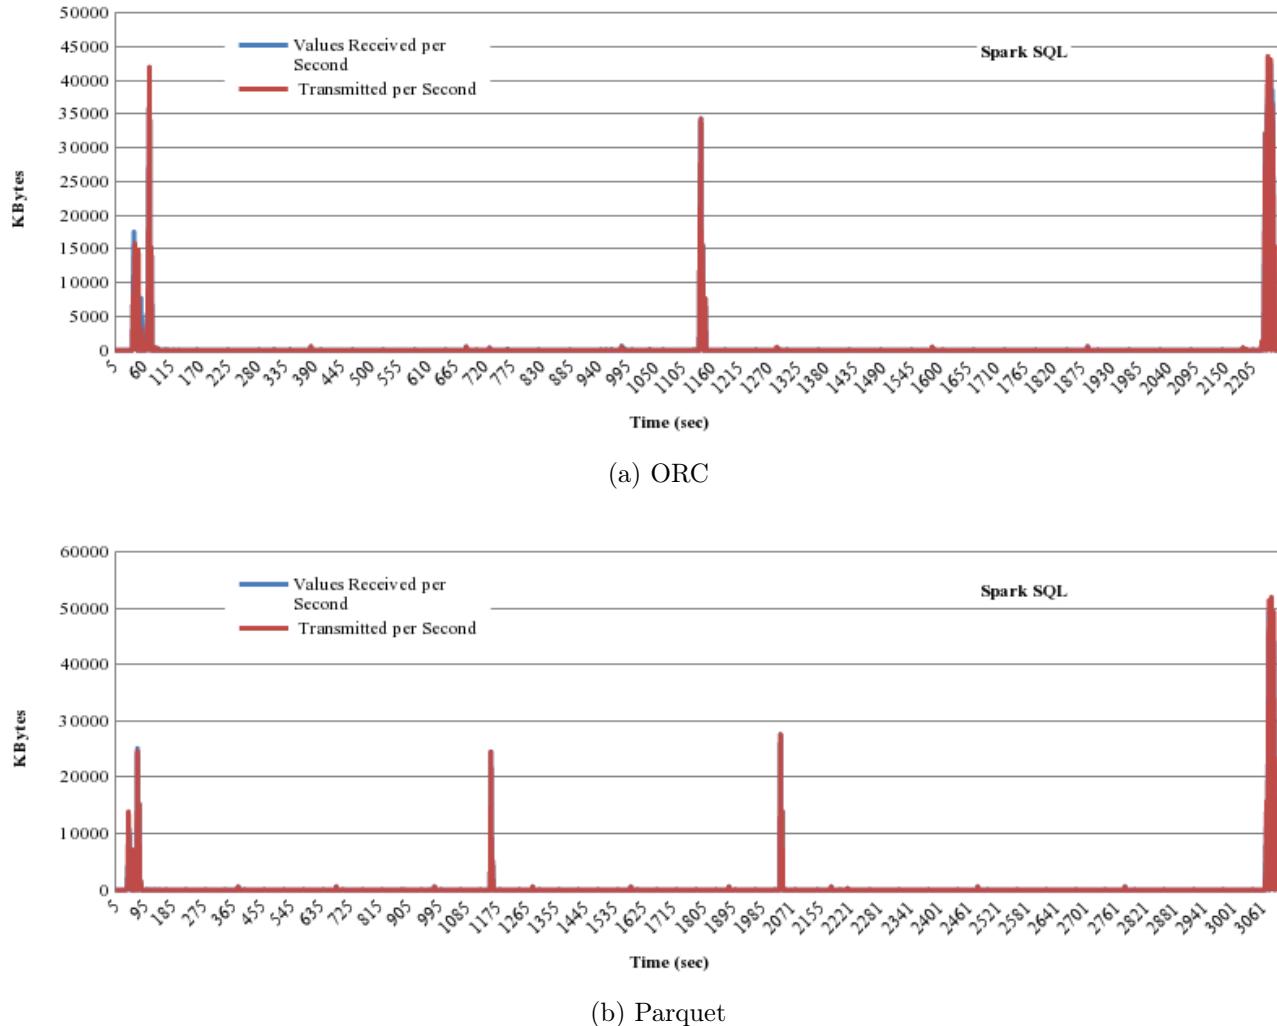


Figure 35: Network Utilization: more or less corresponding to DiskIO BW.

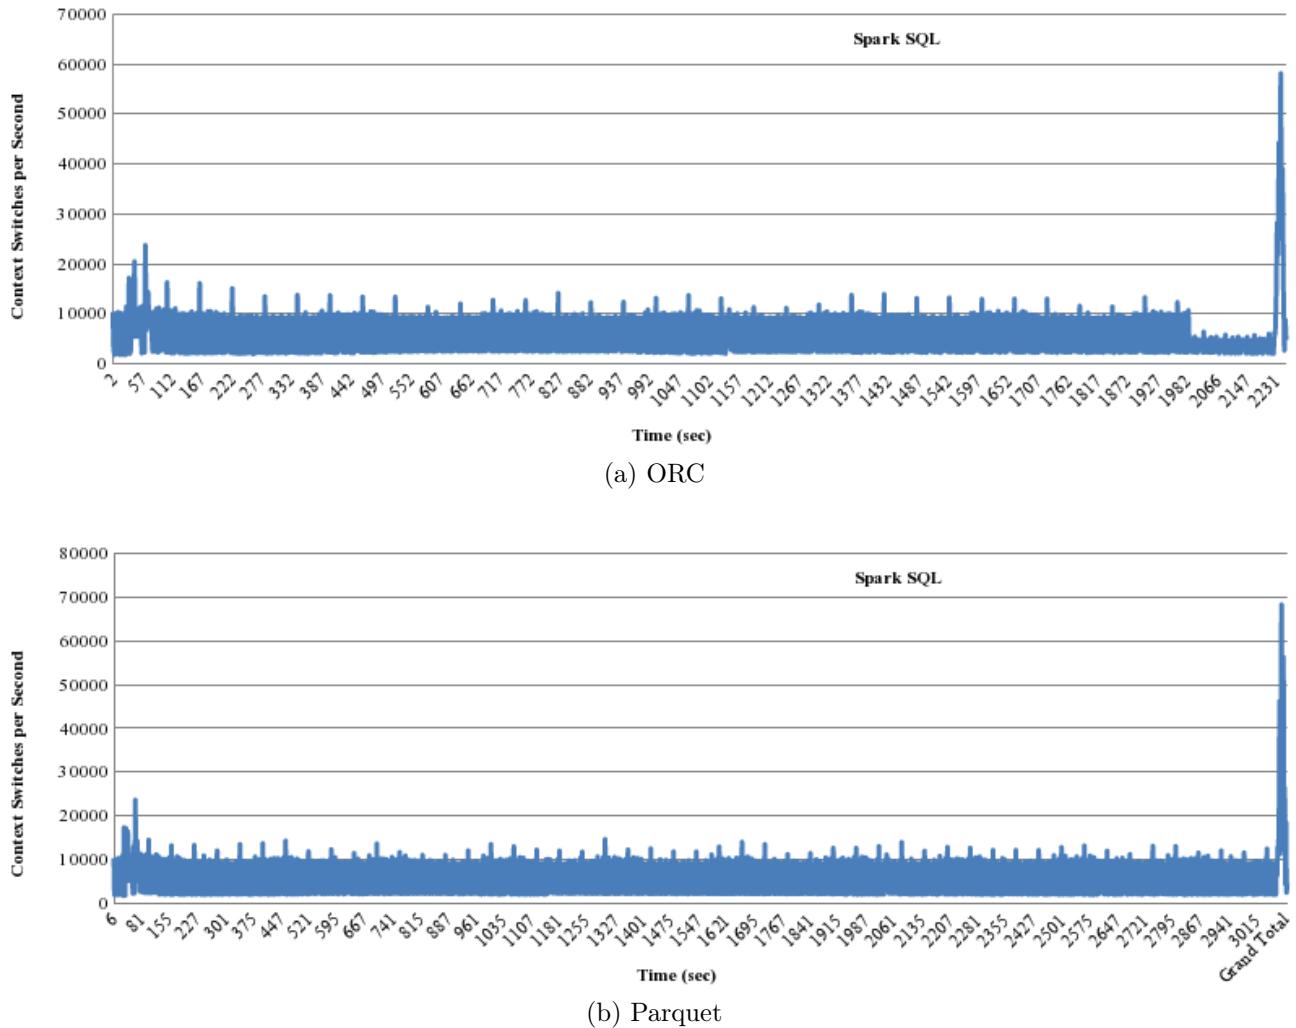


Figure 36: VMSTAT Context Switches

2.9 ORC - NoCompression vs. Snappy

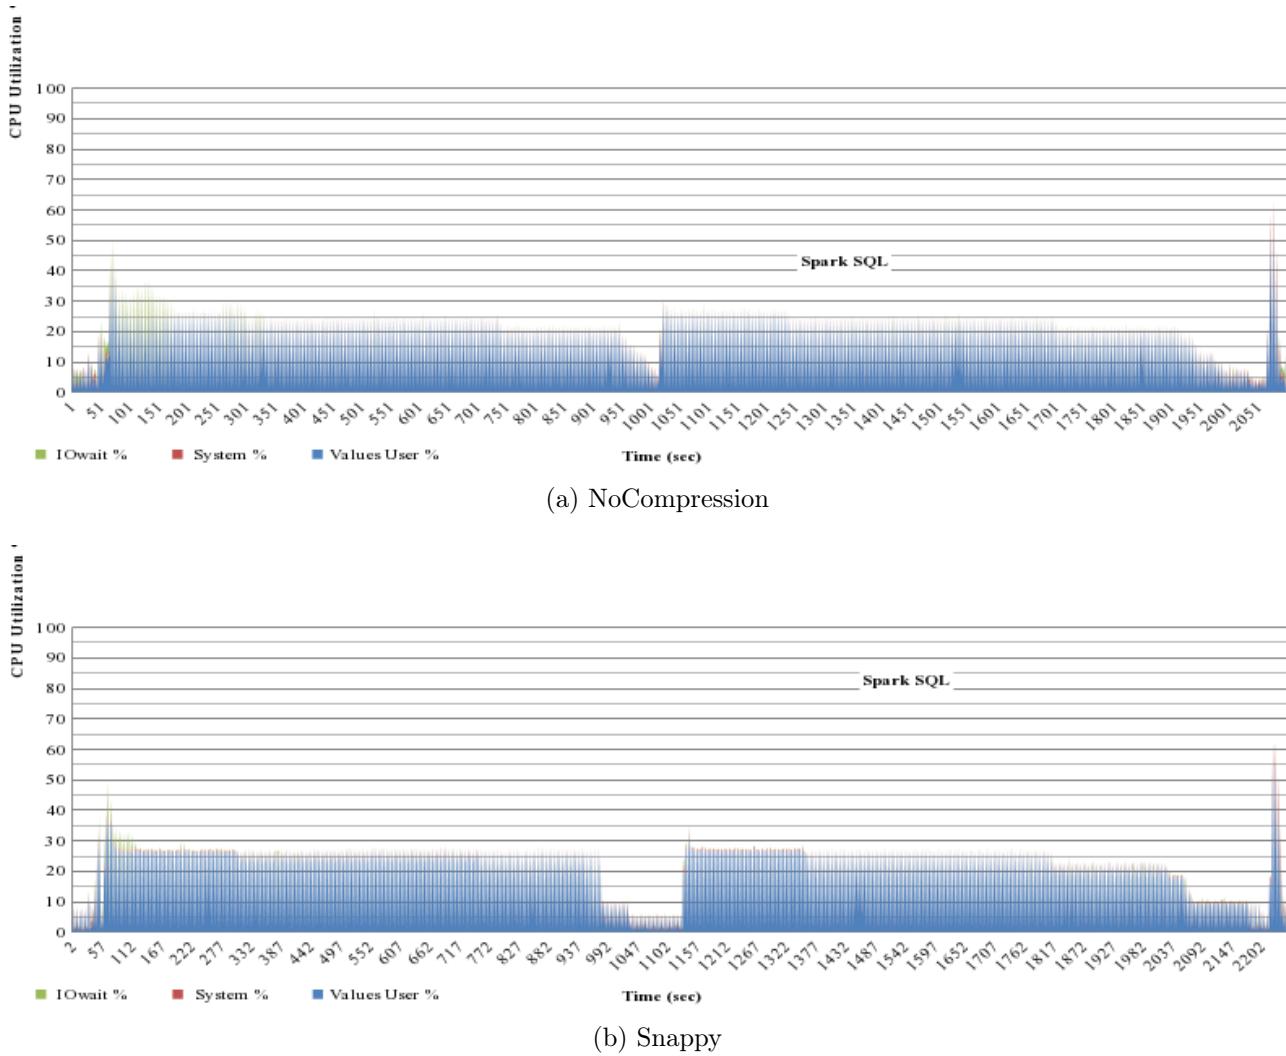


Figure 37: CPU.

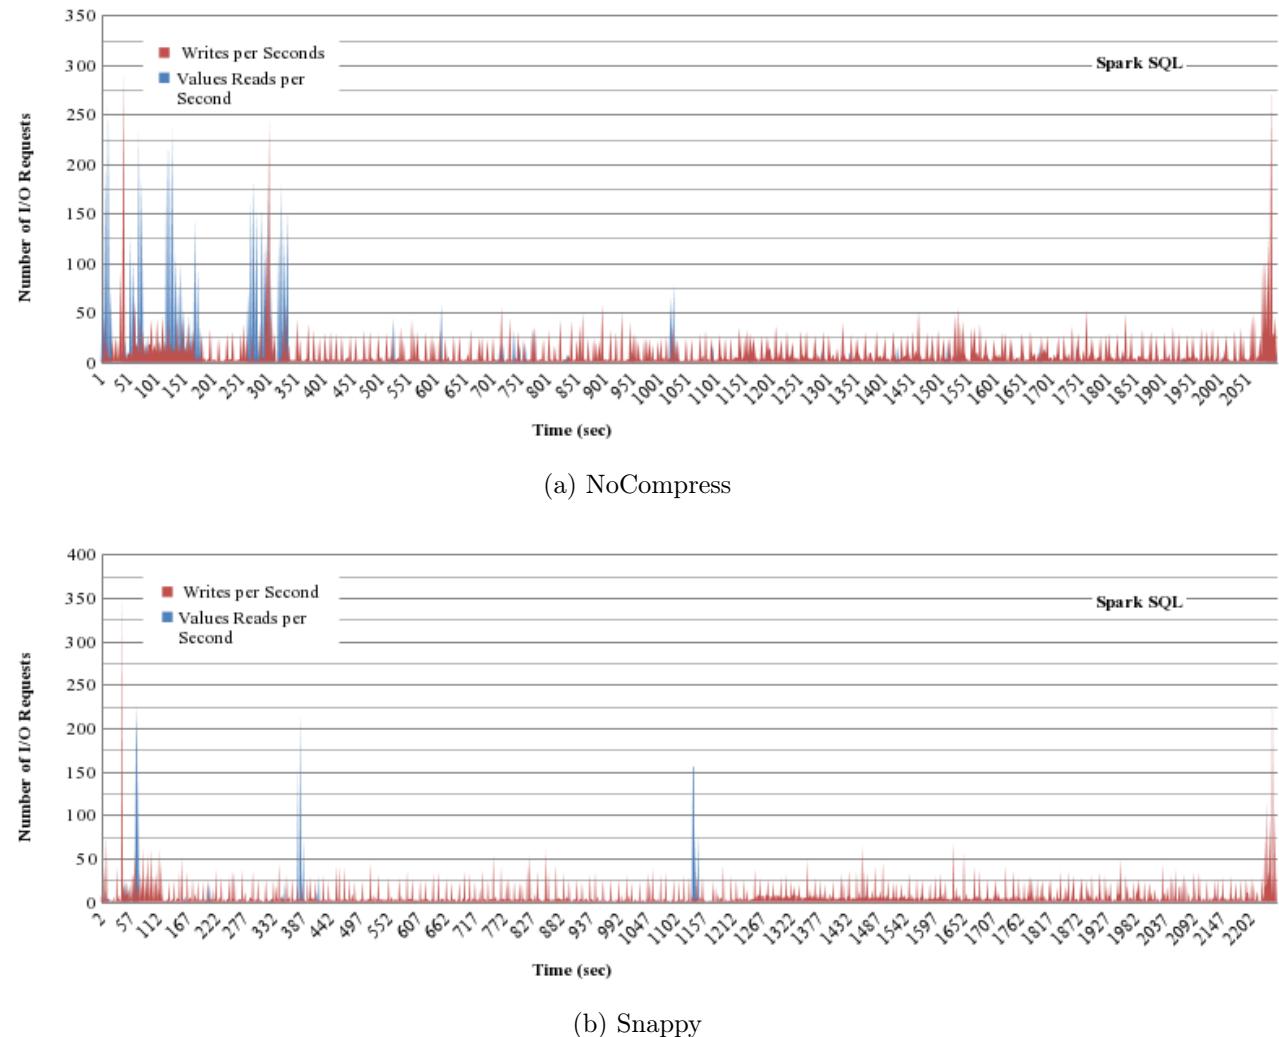


Figure 38: Disk Requests.

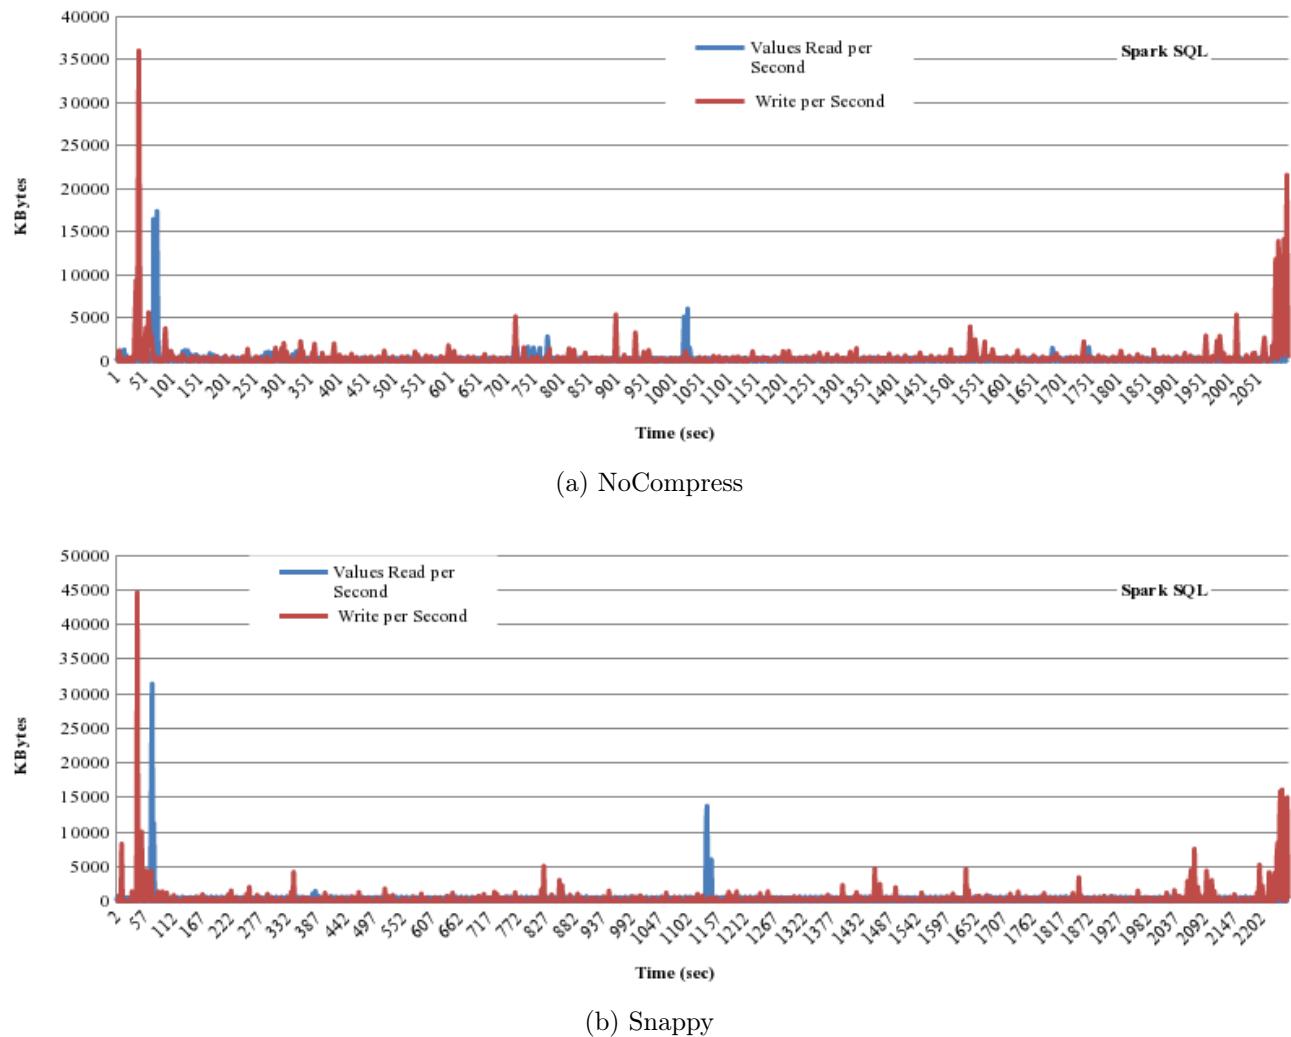


Figure 39: Disk Bandwidth.

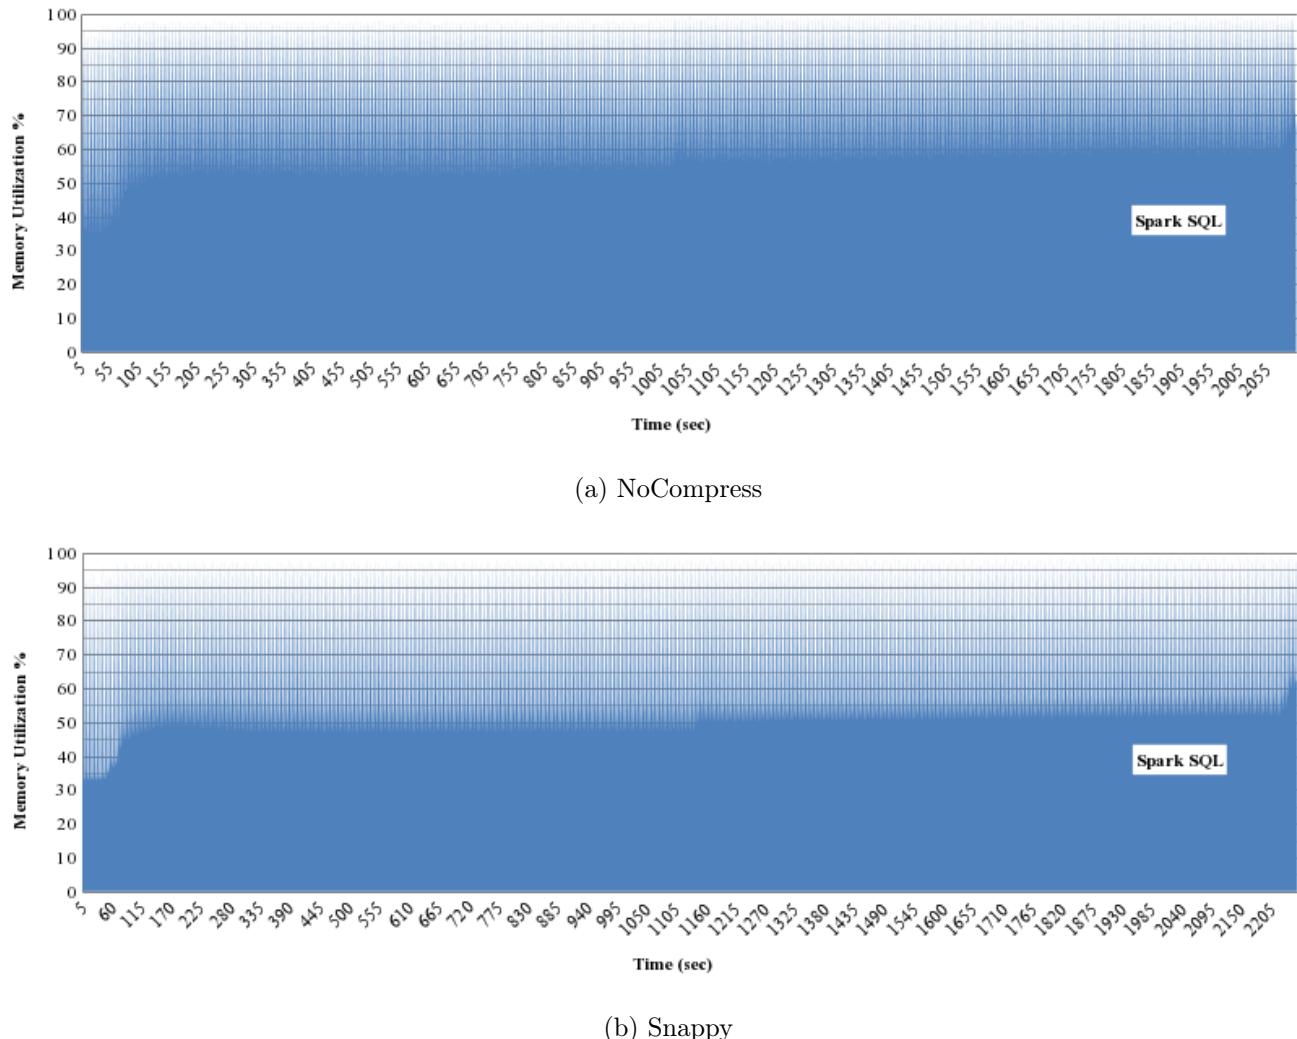


Figure 40: Free Memory.

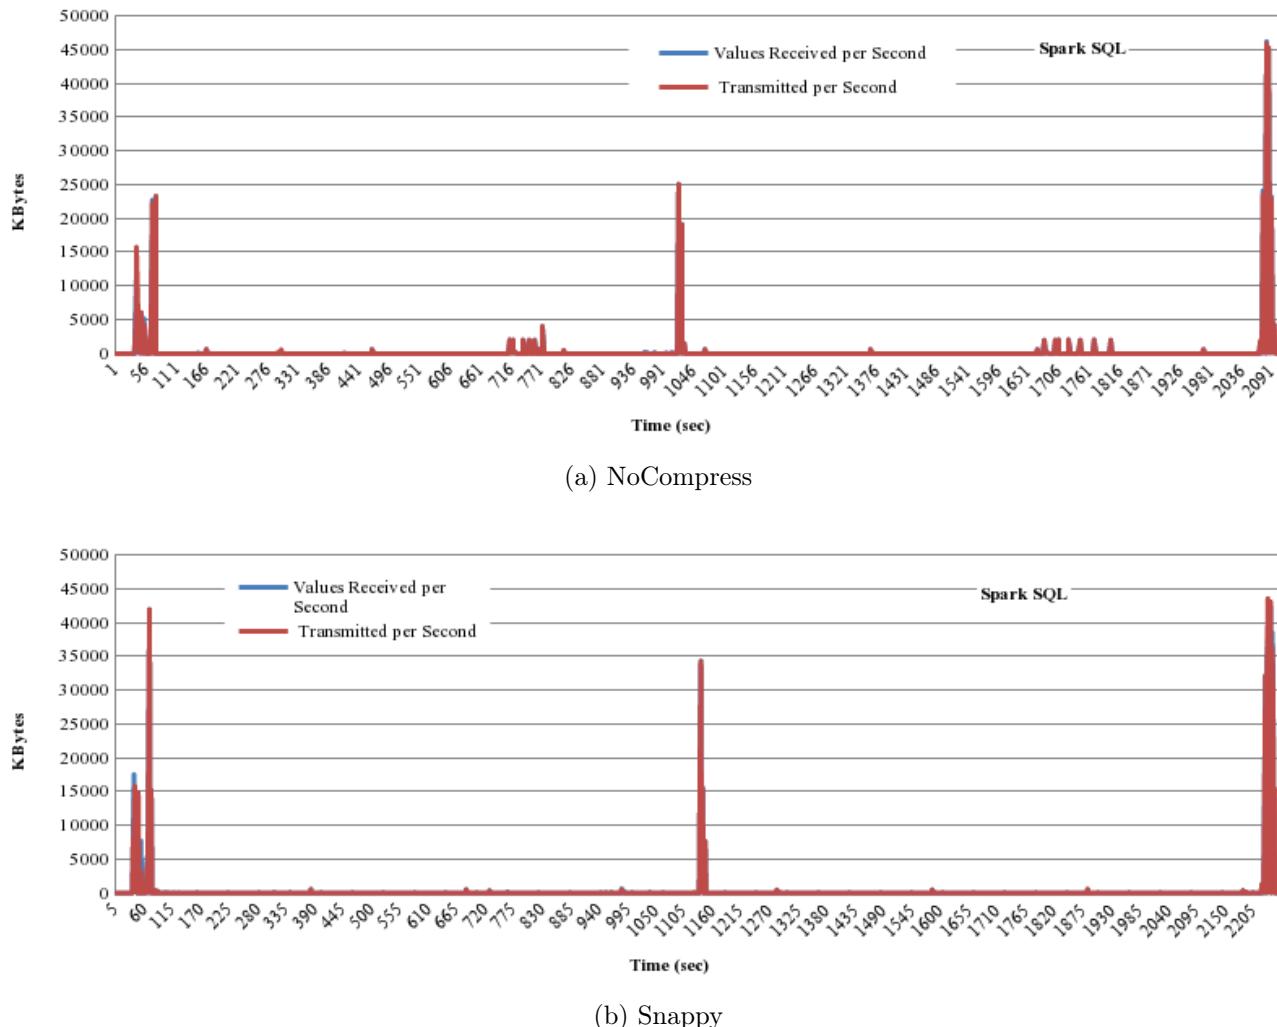


Figure 41: Network Utilization.

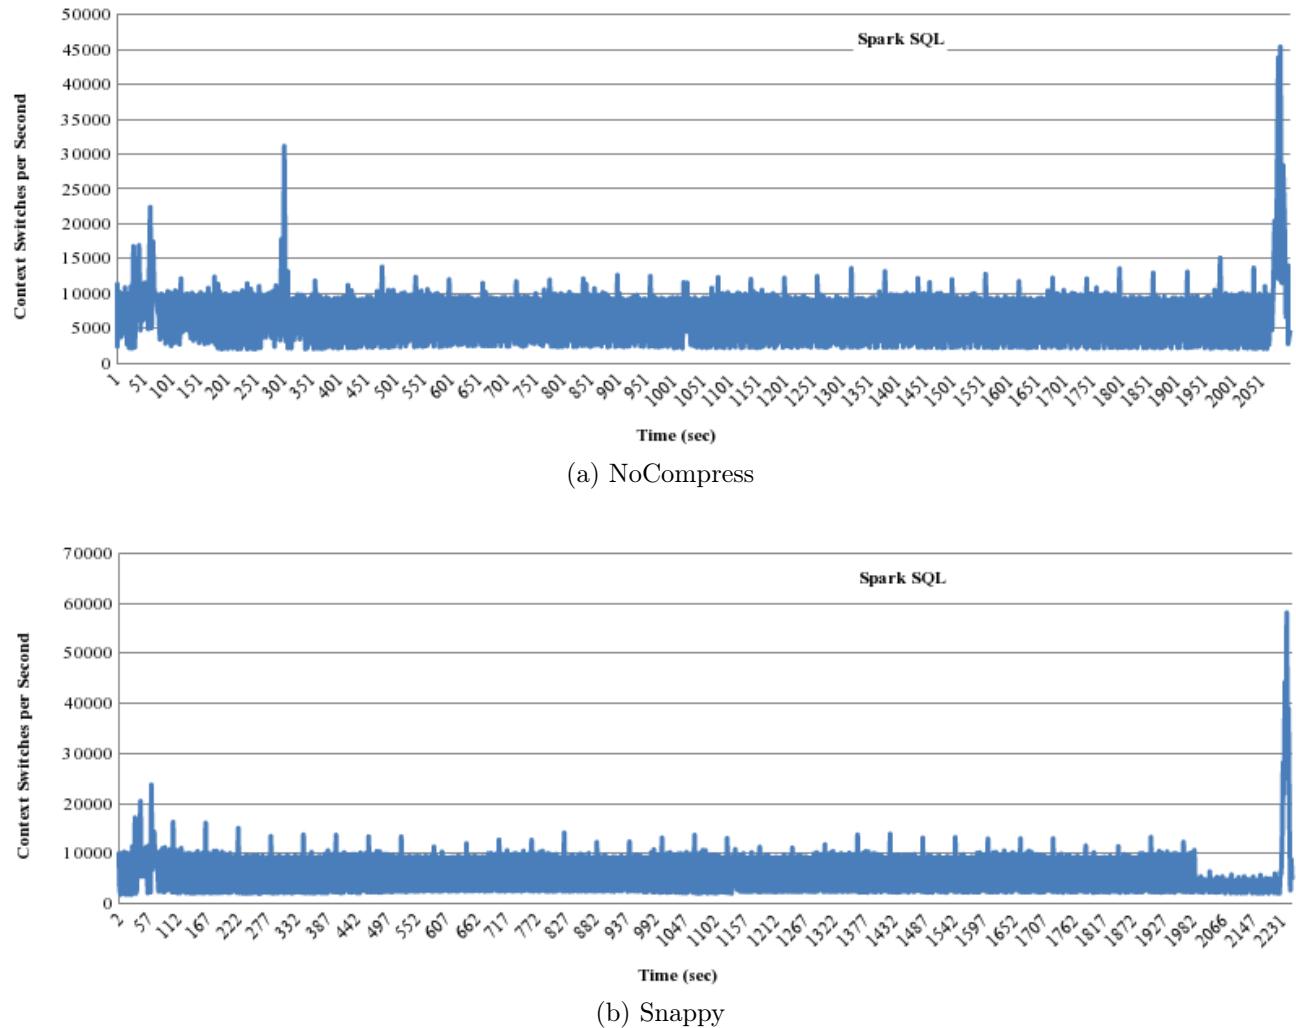


Figure 42: VMSTAT Context Switches

2.10 Parquet - NoCompression vs Snappy

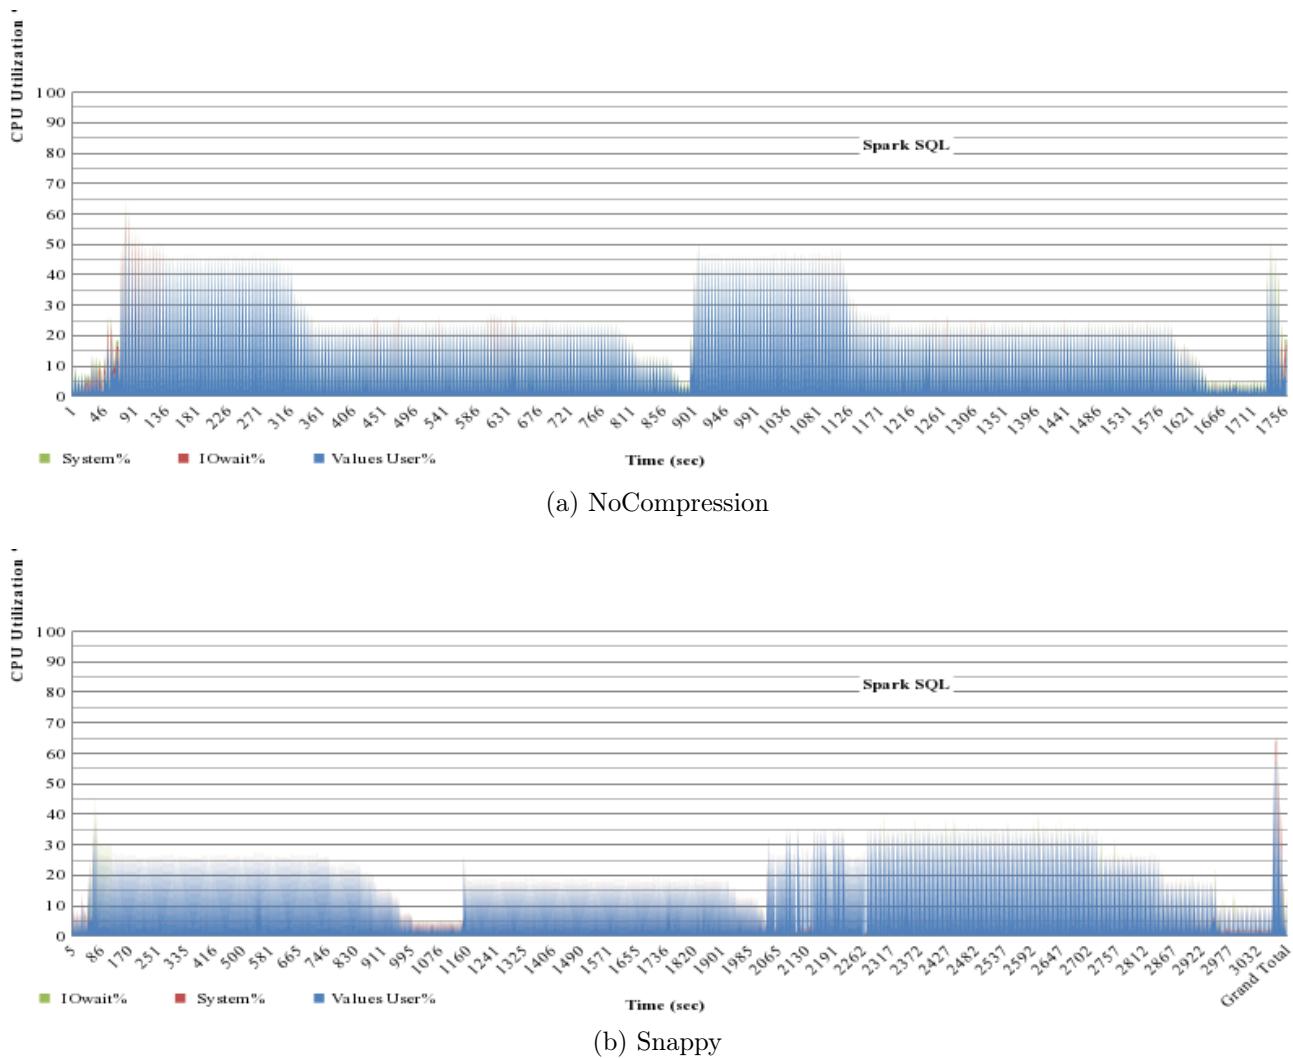


Figure 43: CPU.

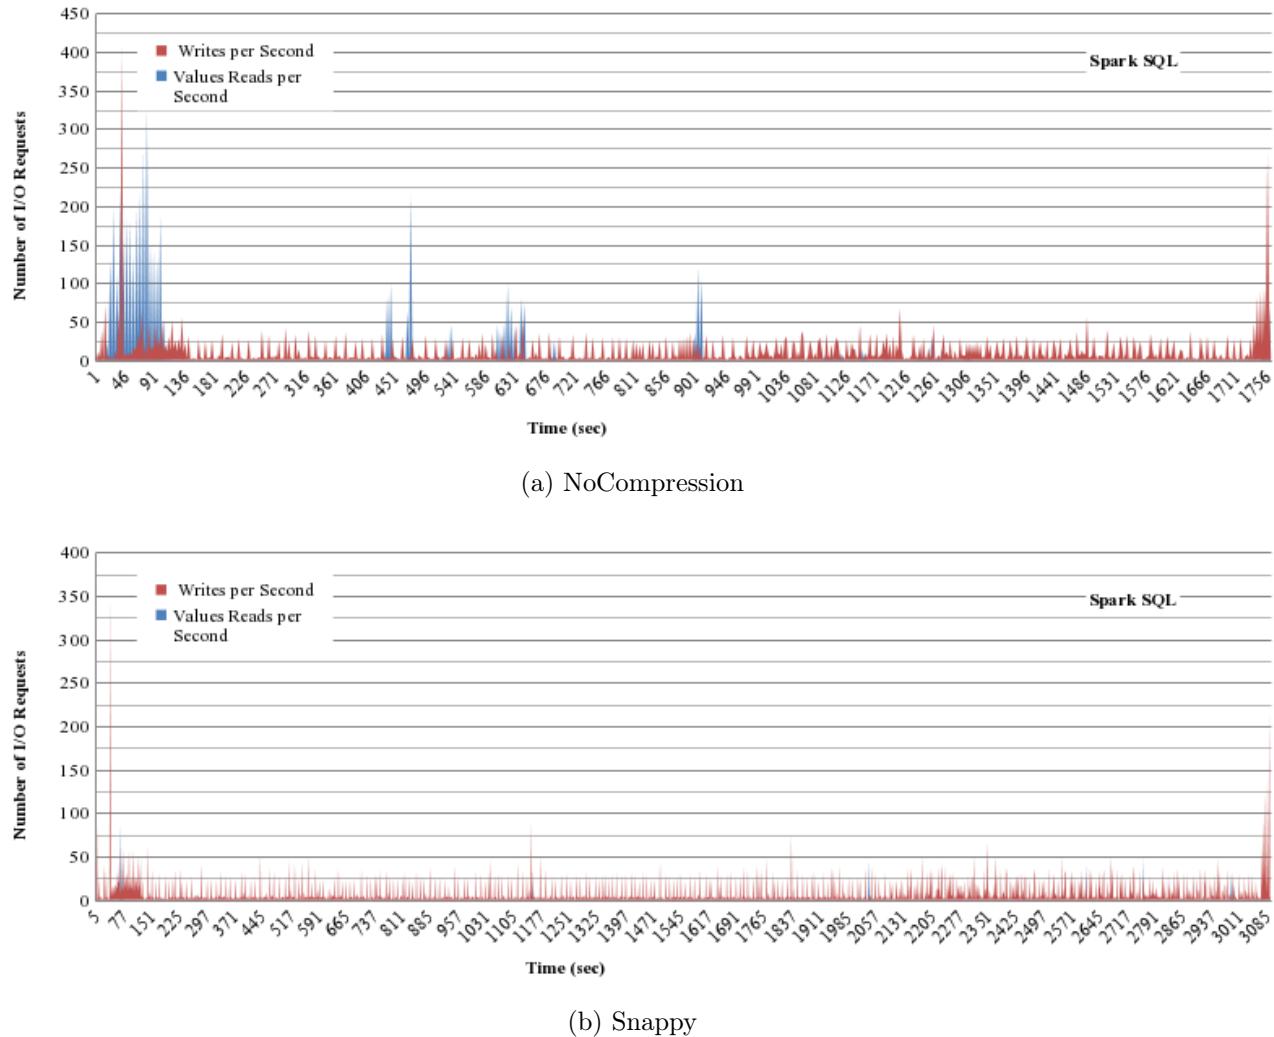


Figure 44: Disk Requests.

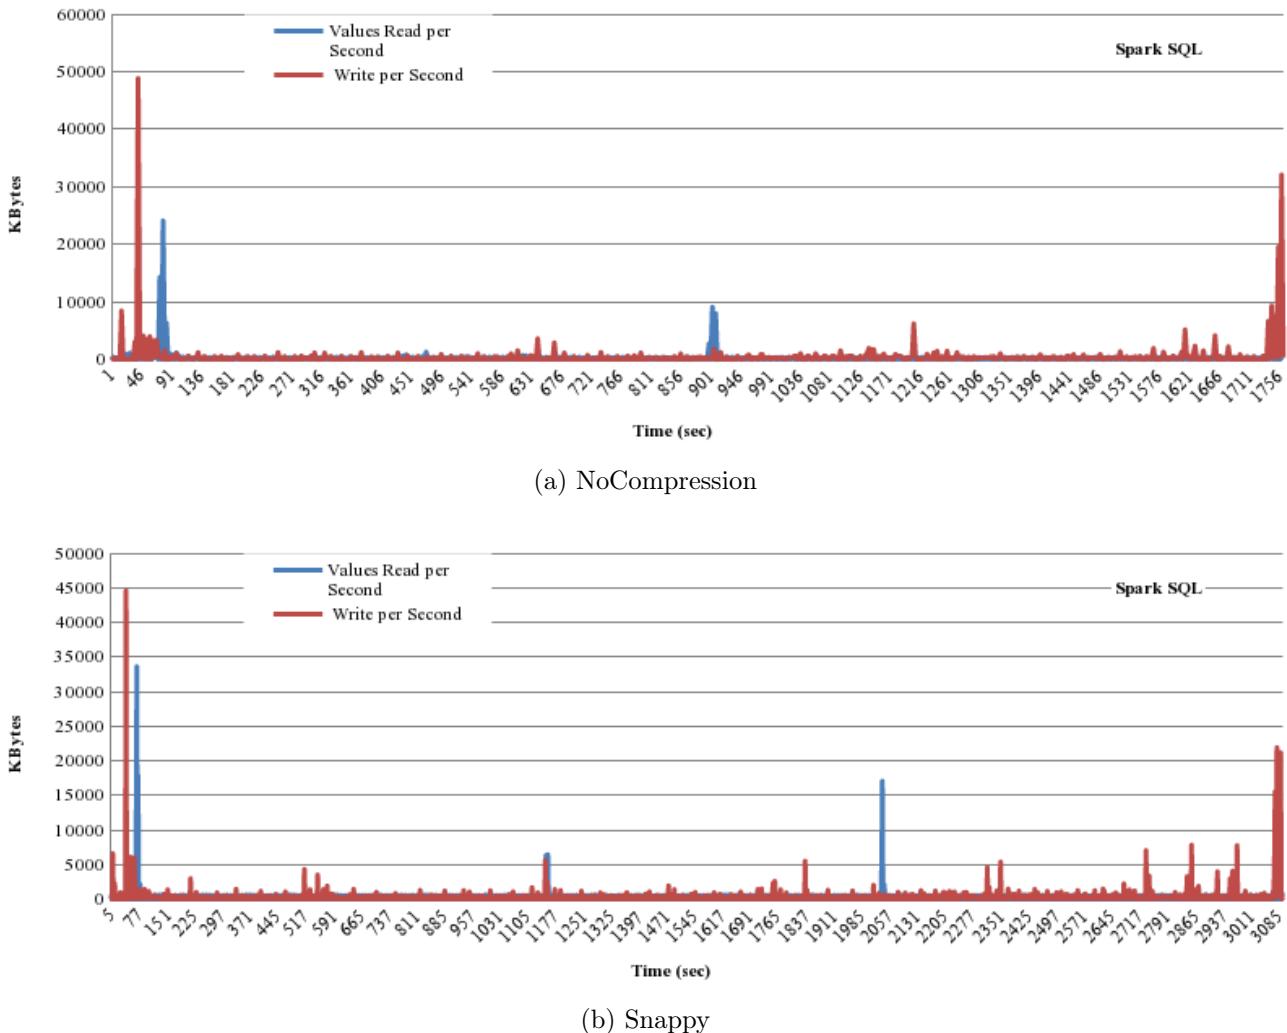


Figure 45: Disk Bandwidth.

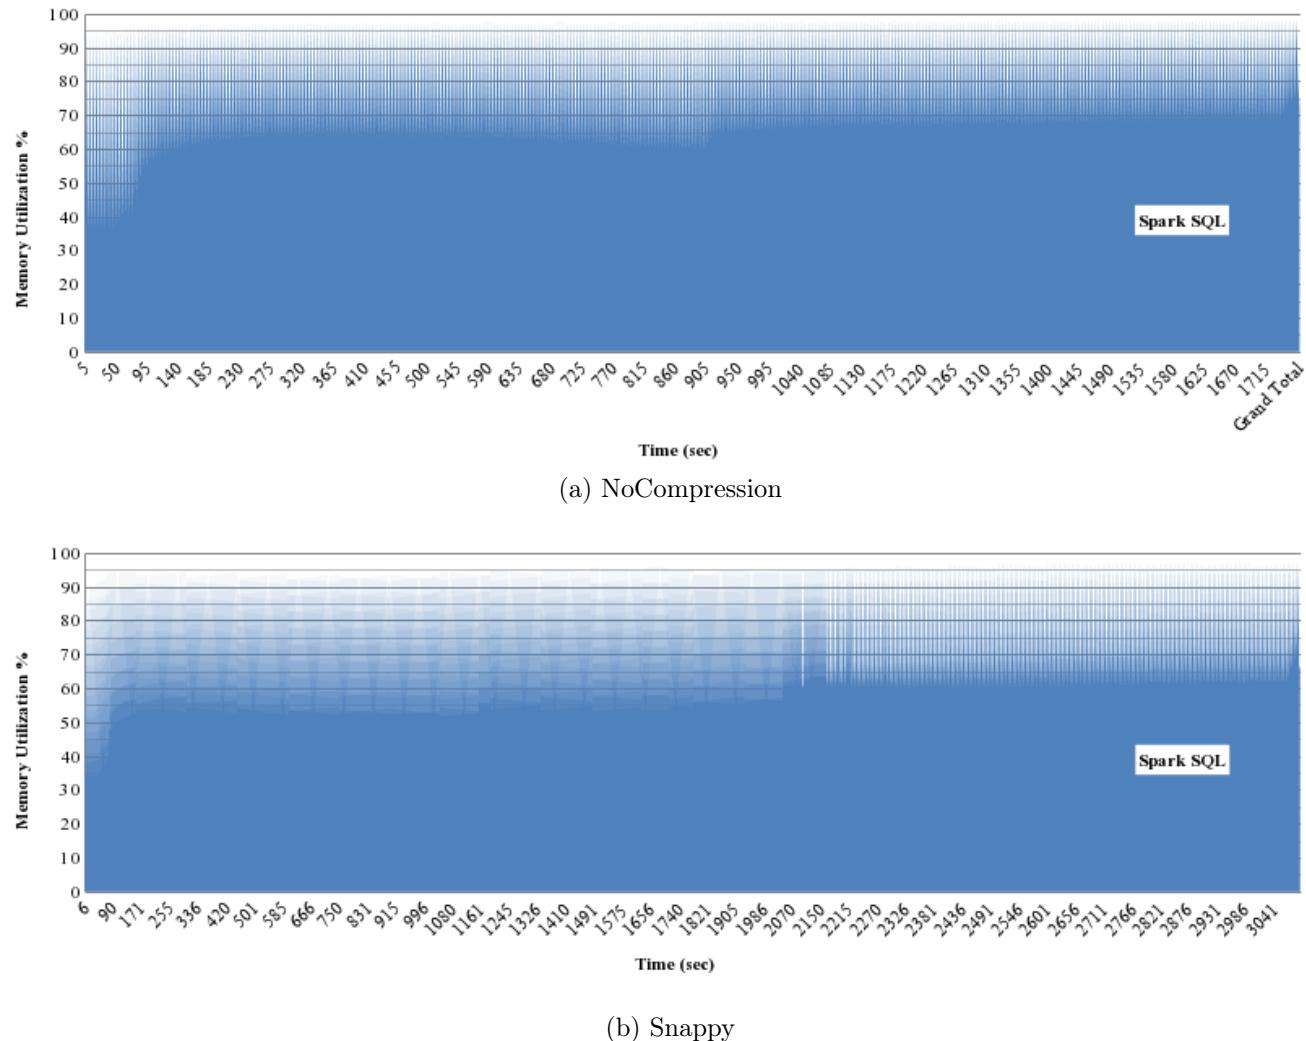


Figure 46: Free Memory.

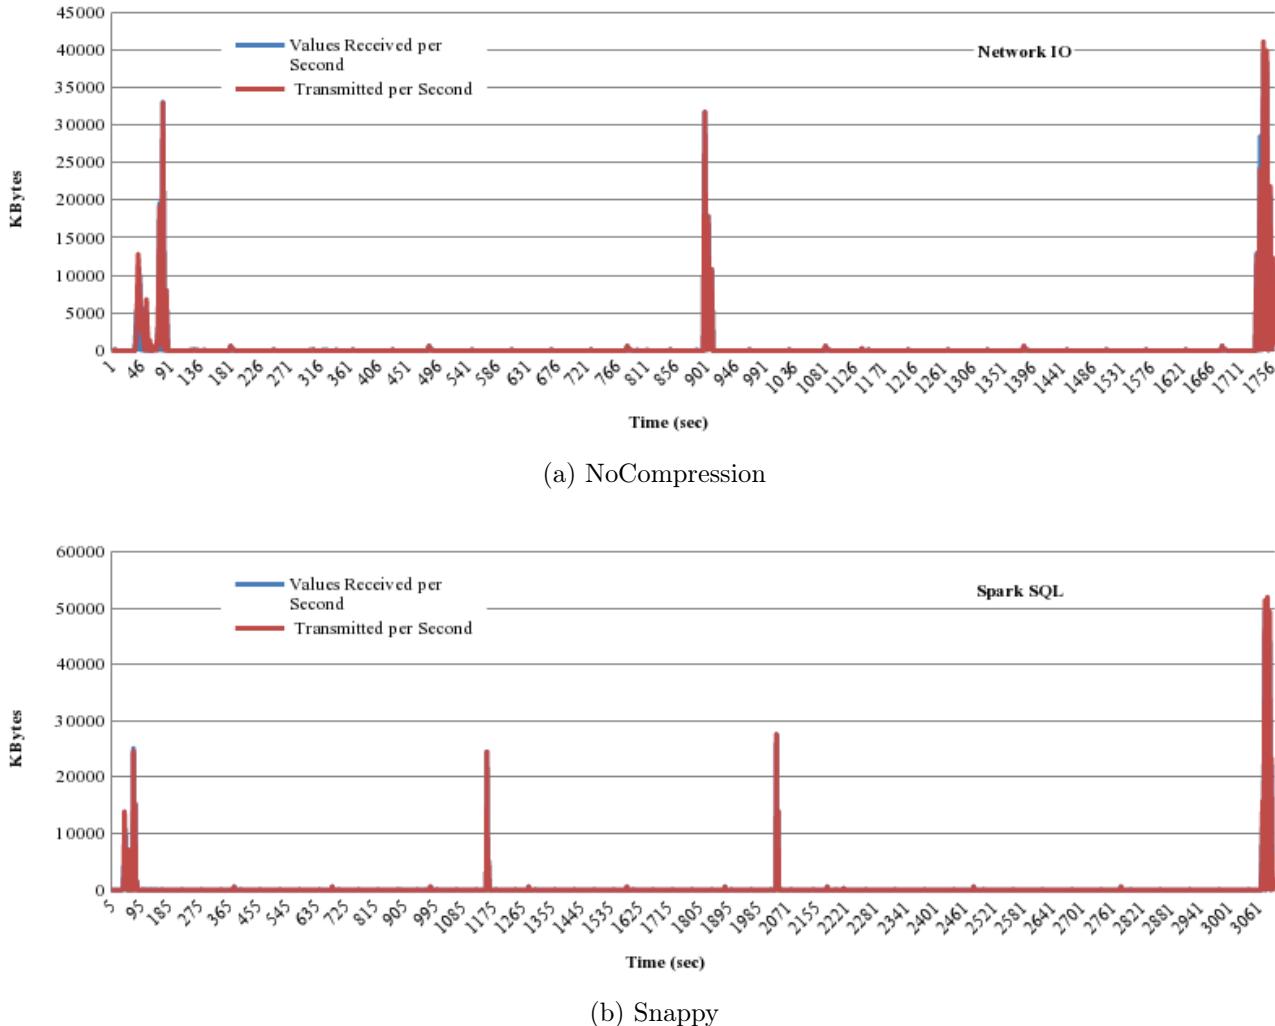


Figure 47: Network Utilization.

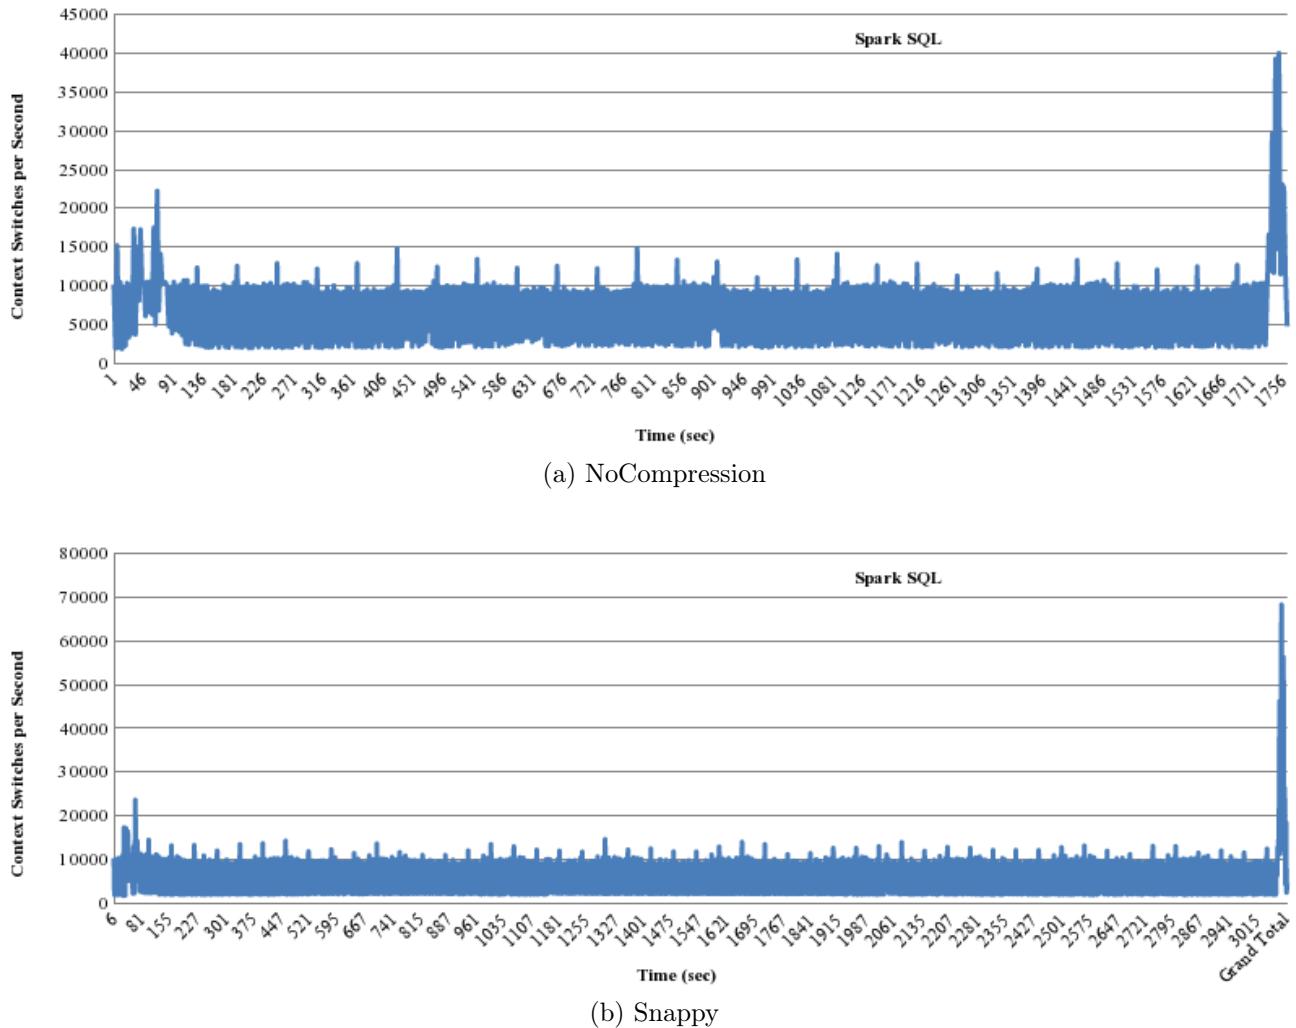


Figure 48: VMSTAT Context Switches

3 Q12

3.1 Type

Pure HiveQL

3.2 Description

Find all customers who viewed items of a given category on the web in a given month and year that was followed by an in-store purchase of an item from the same category in the three consecutive months.

3.3 Performance

	Default (sec.)		No compr. (sec.)		Snappy (sec.)	
Query	ORC	Parquet	ORC	Parquet	ORC	Parquet
Q12	331	105	331	118	314	122

3.4 Phase 1

There are no temp tables. The query gets users with the specified requirements from physical store and web store and joins them. Then, it shows the result.

3.5 Source

GitHub: <https://git.io/vpFey>

3.6 Notes

...

3.7 No compression - ORC vs Parquet

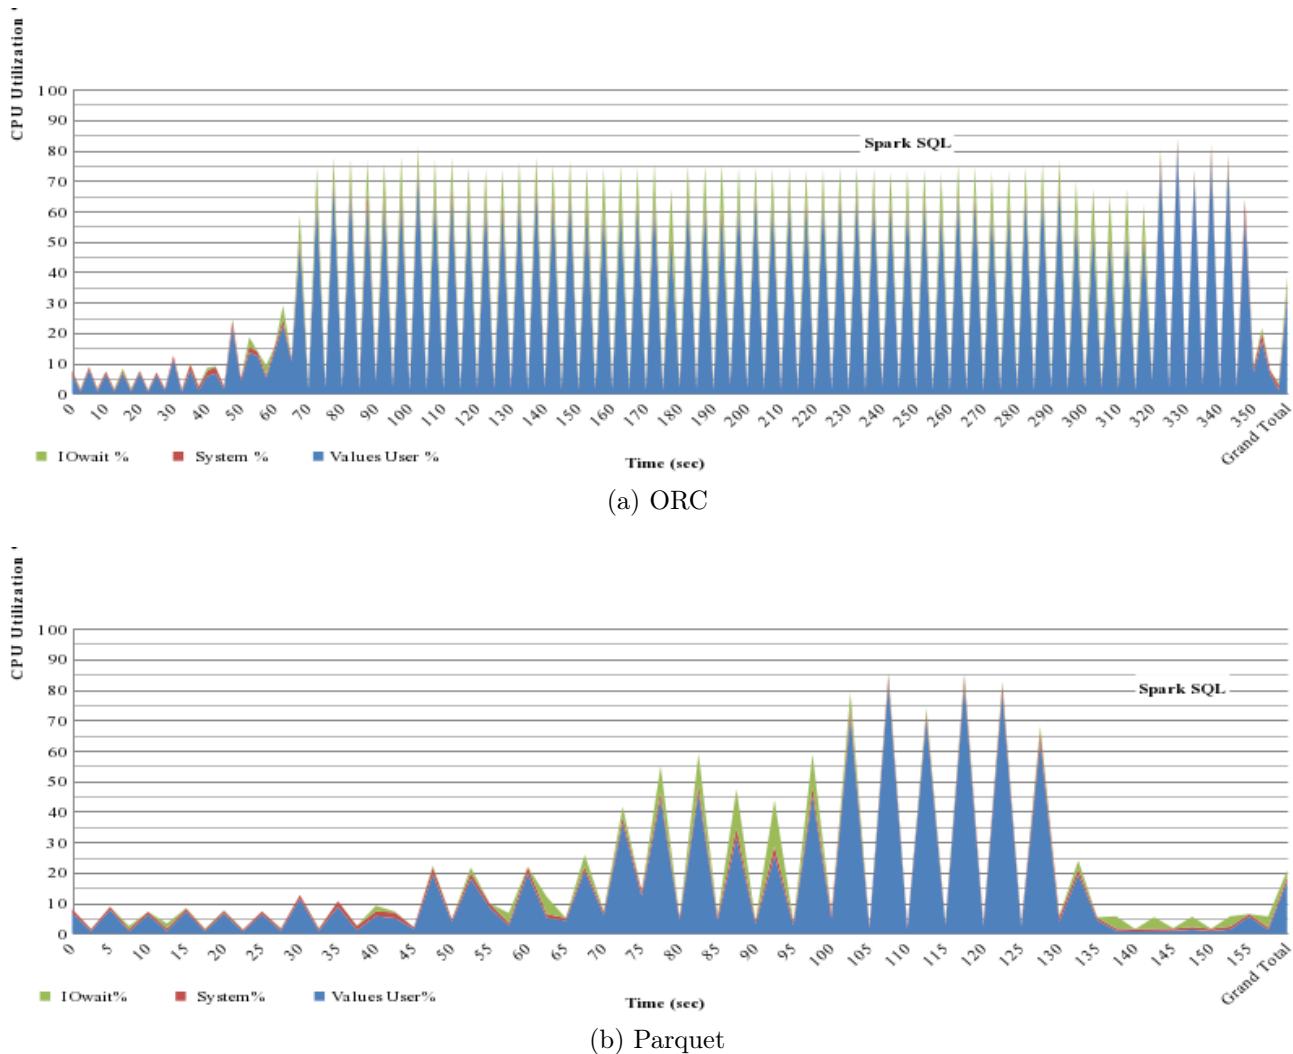


Figure 49: CPU: Lower utilization on Parquet (max is 80%)

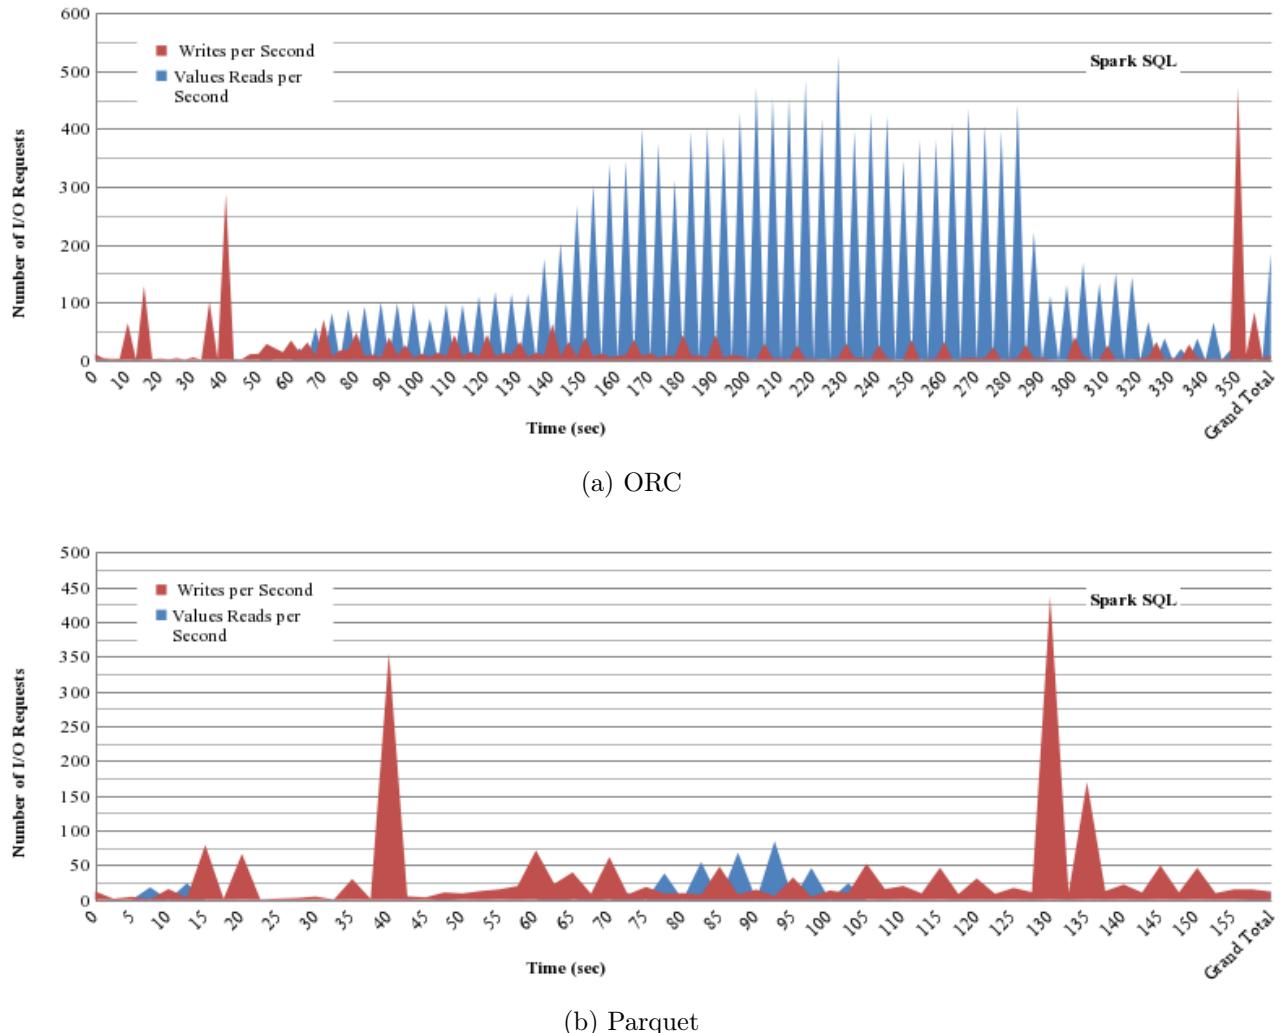


Figure 50: Disk Requests: More requests on Parquet.

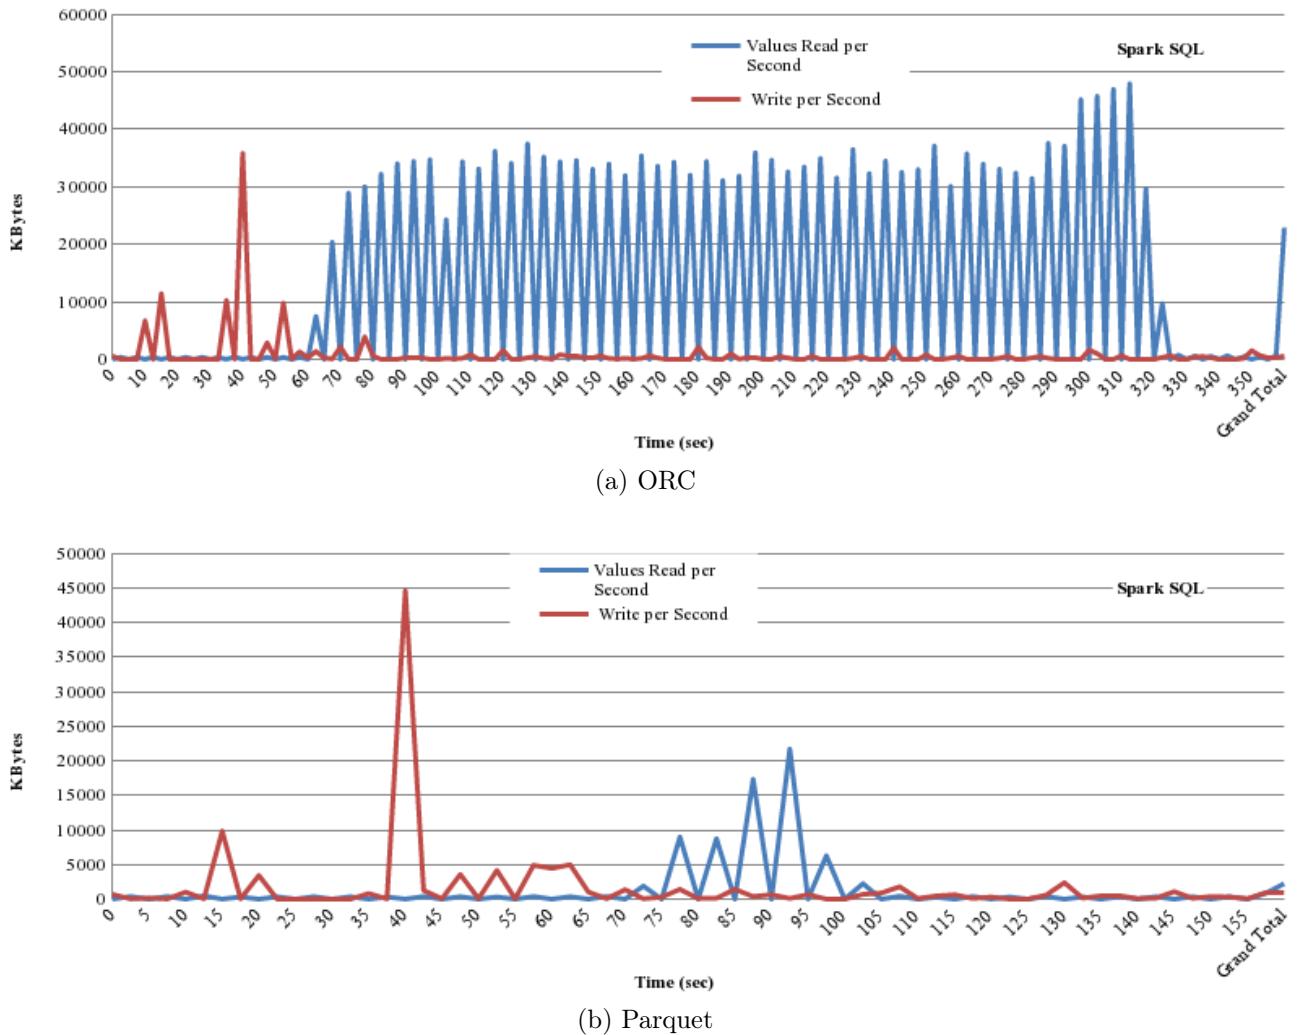
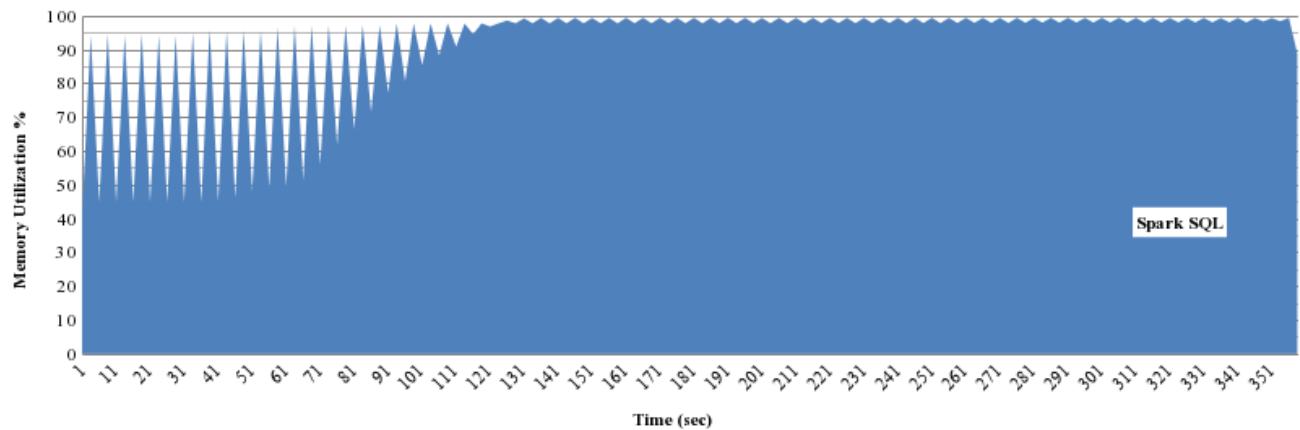
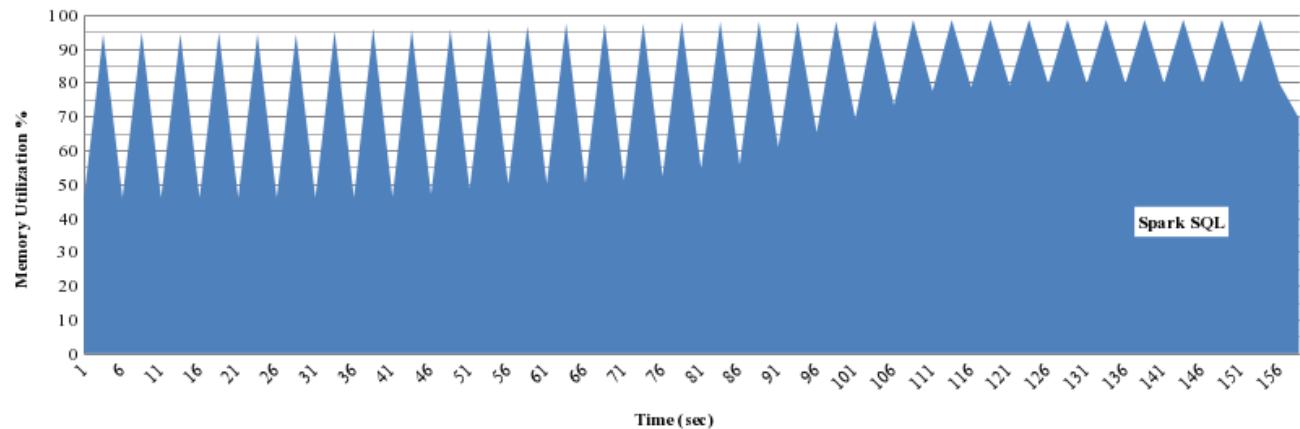


Figure 51: Disk Bandwidth: bandwidth utilization on Parquet.



(a) ORC



(b) Parquet

Figure 52: Free Memory.

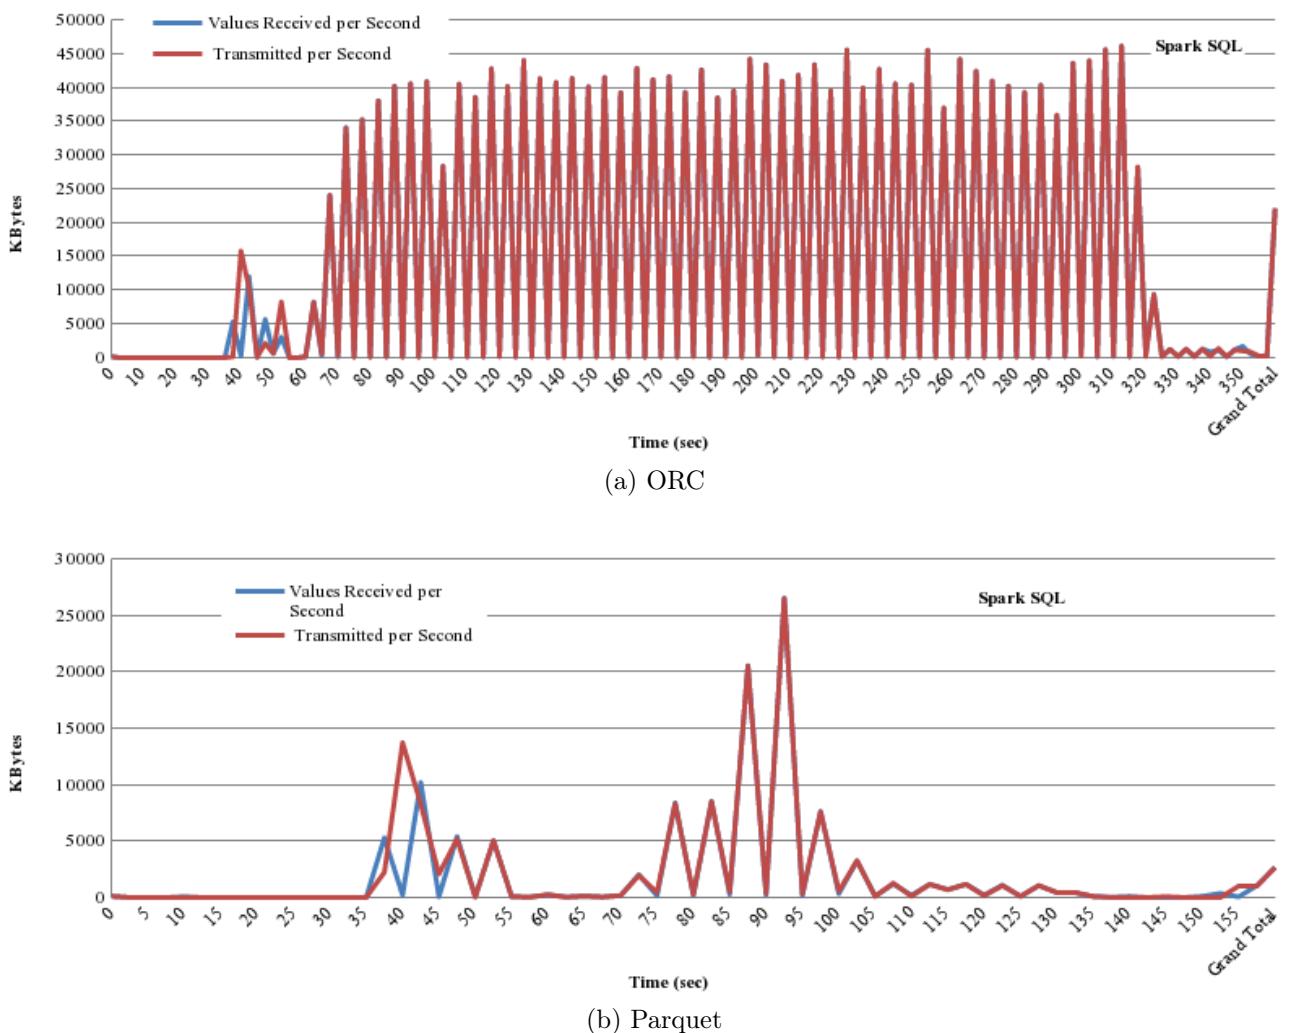


Figure 53: Network Utilization: More or less corresponding to DiskIO.

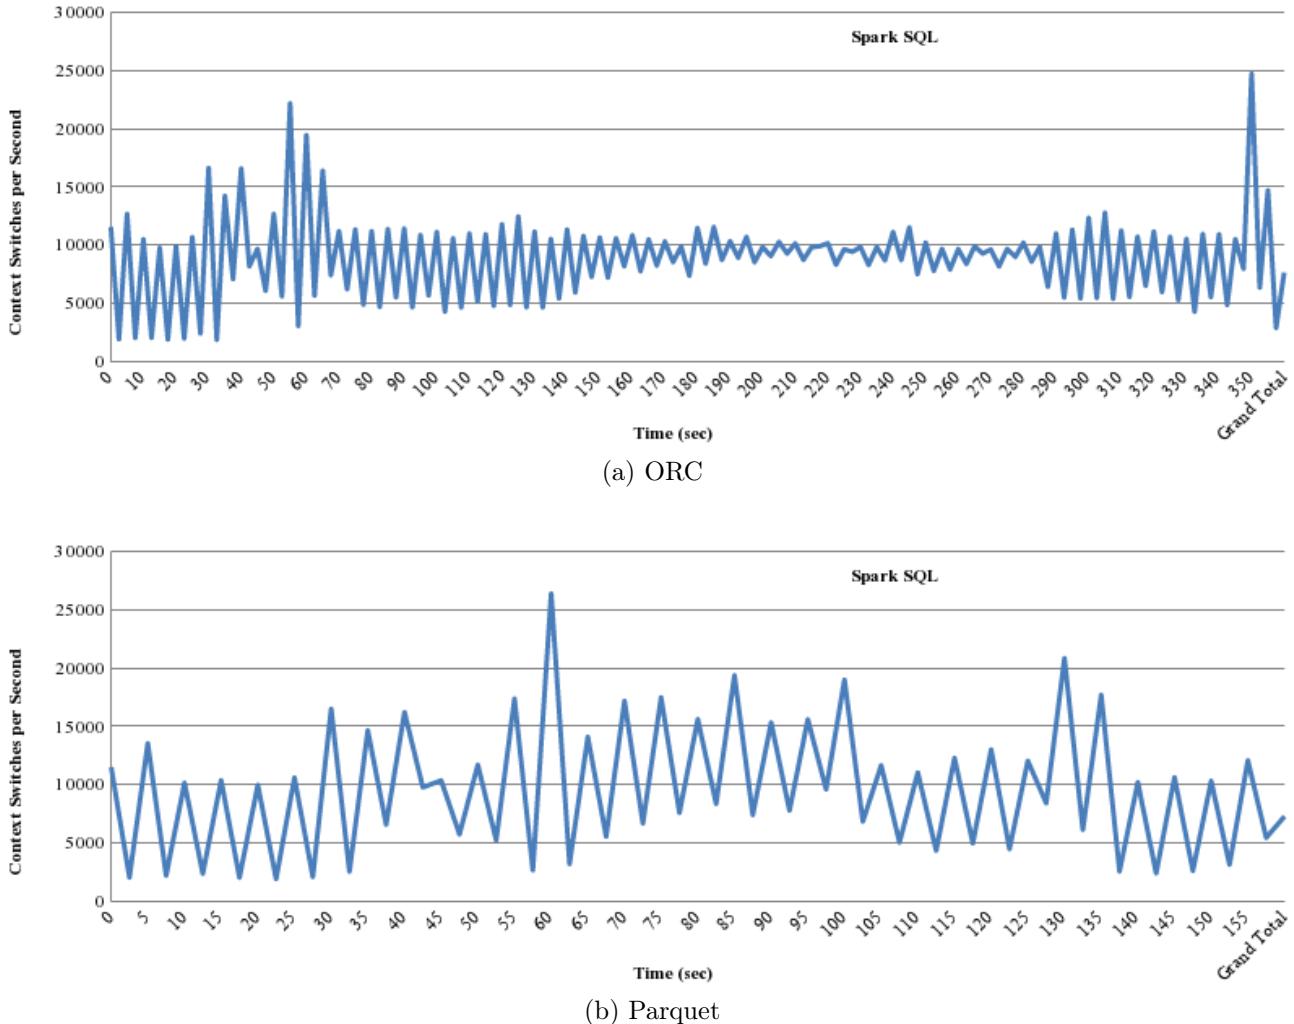


Figure 54: VMSTAT Context Switches

3.8 Snappy - ORC vs. Parquet

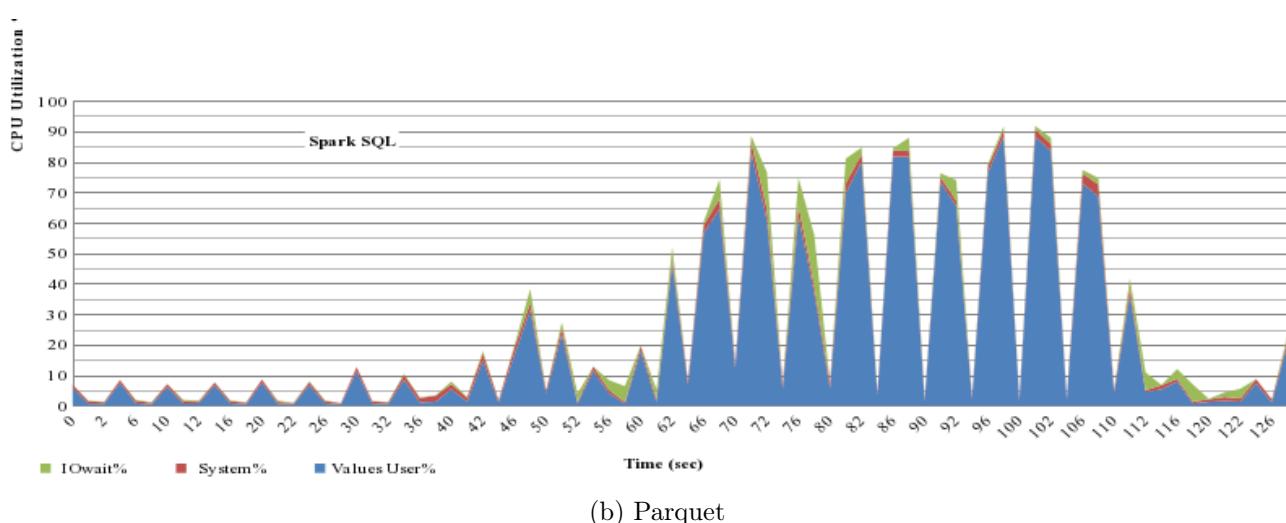
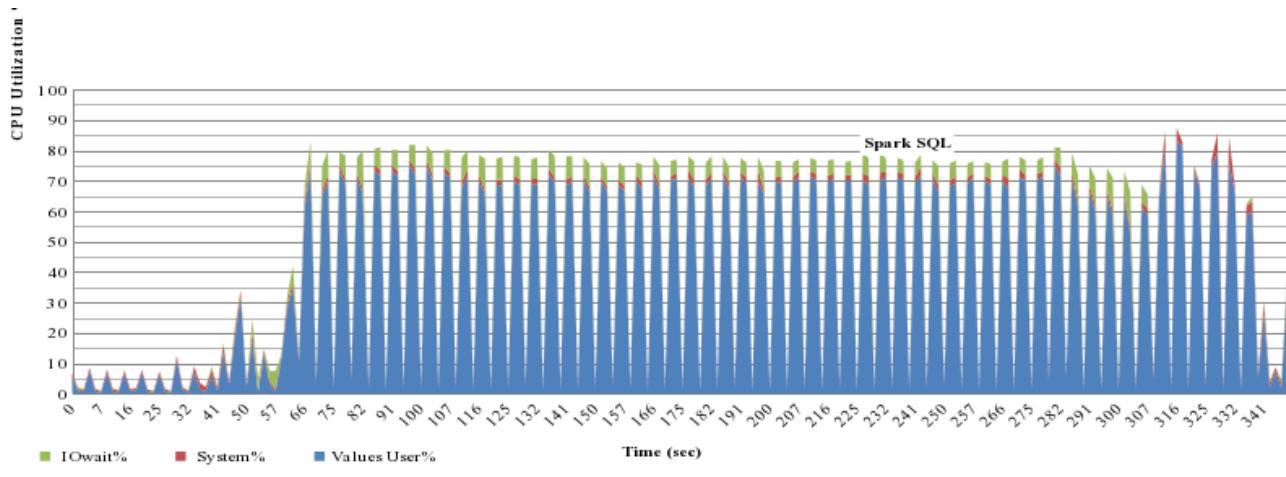


Figure 55: CPU: Low util on parquet (60%), short time for phase 1.

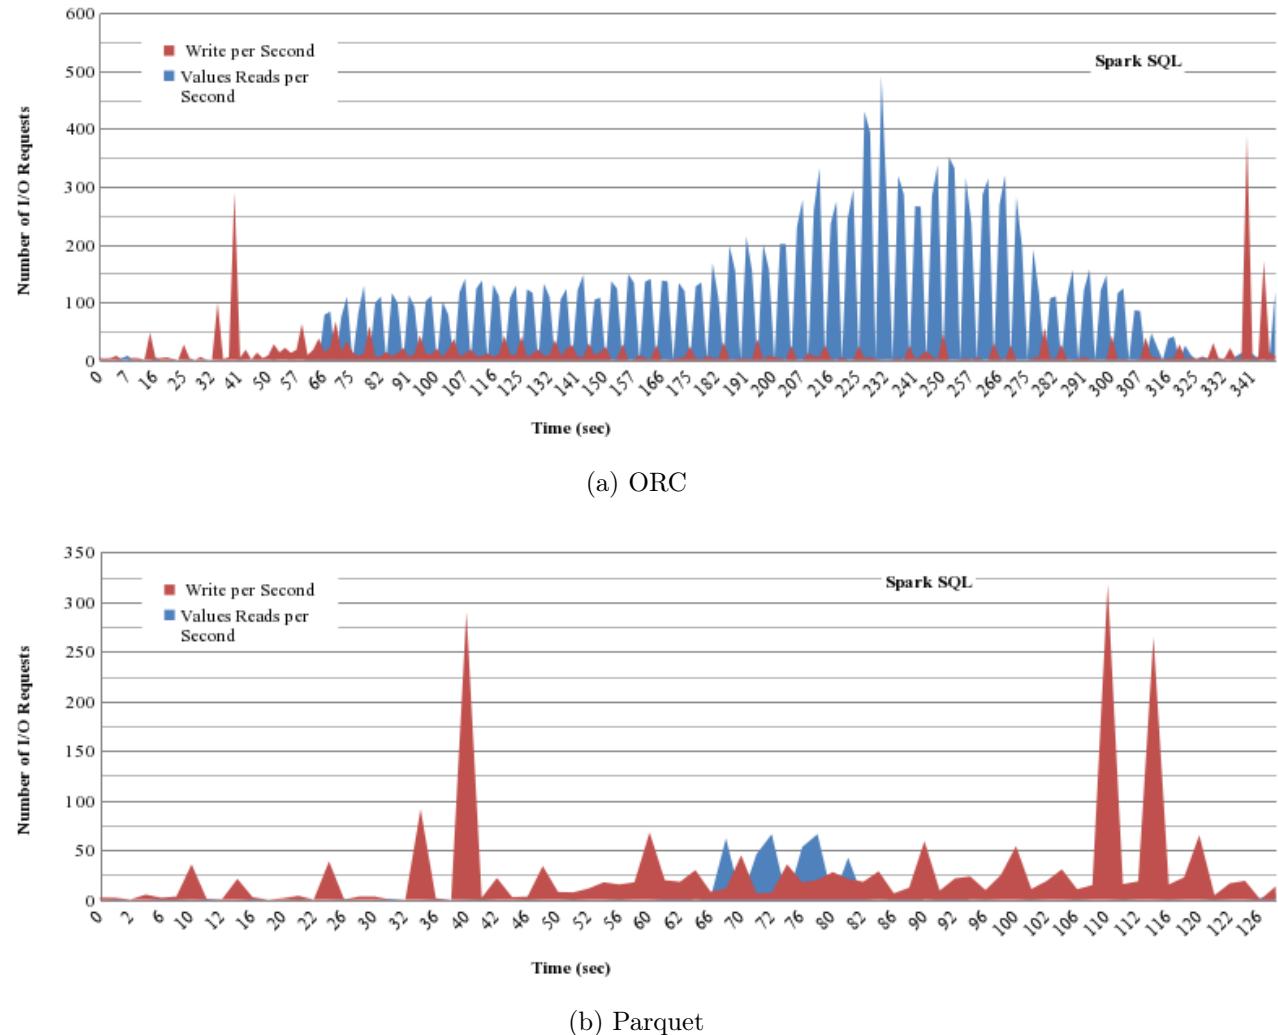


Figure 56: Disk Requests: Lower number of requests on Parquet.

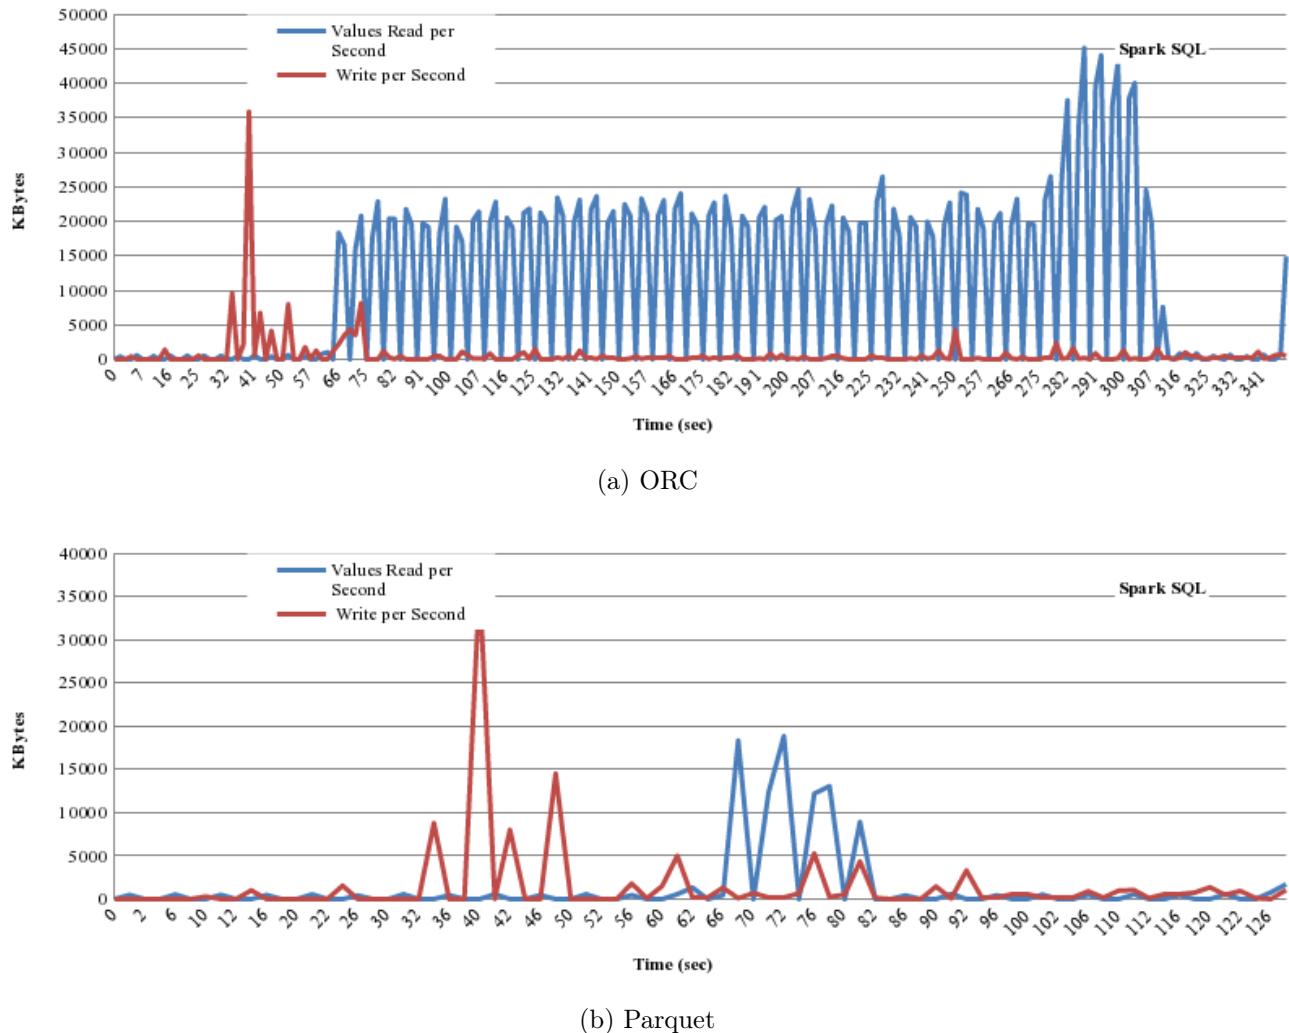


Figure 57: Disk Bandwidth: higher BW utilization on Parquet.

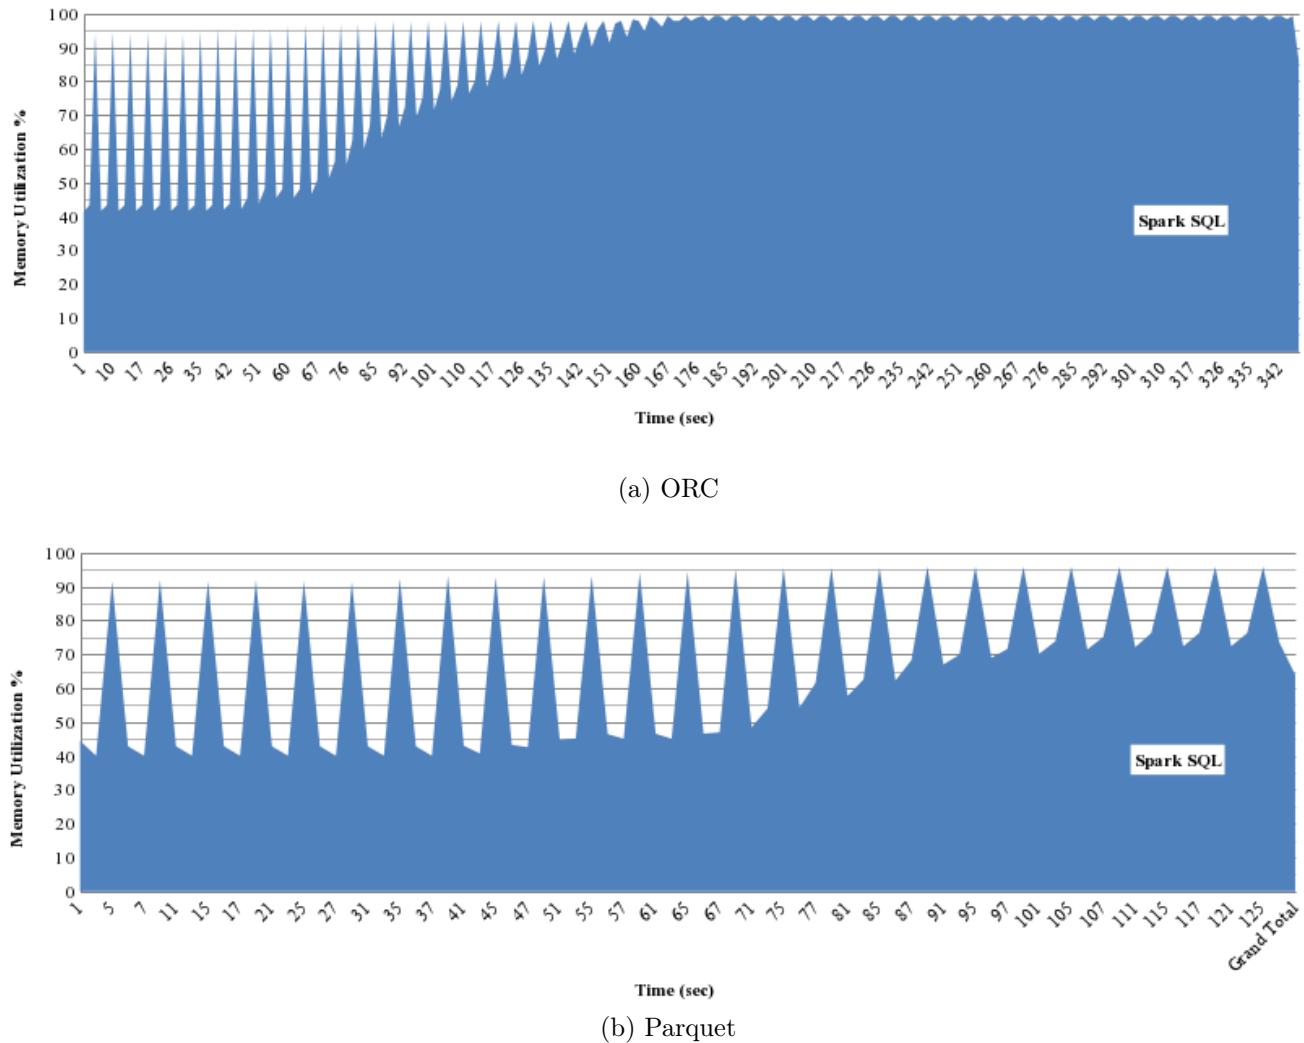


Figure 58: Free Memory.

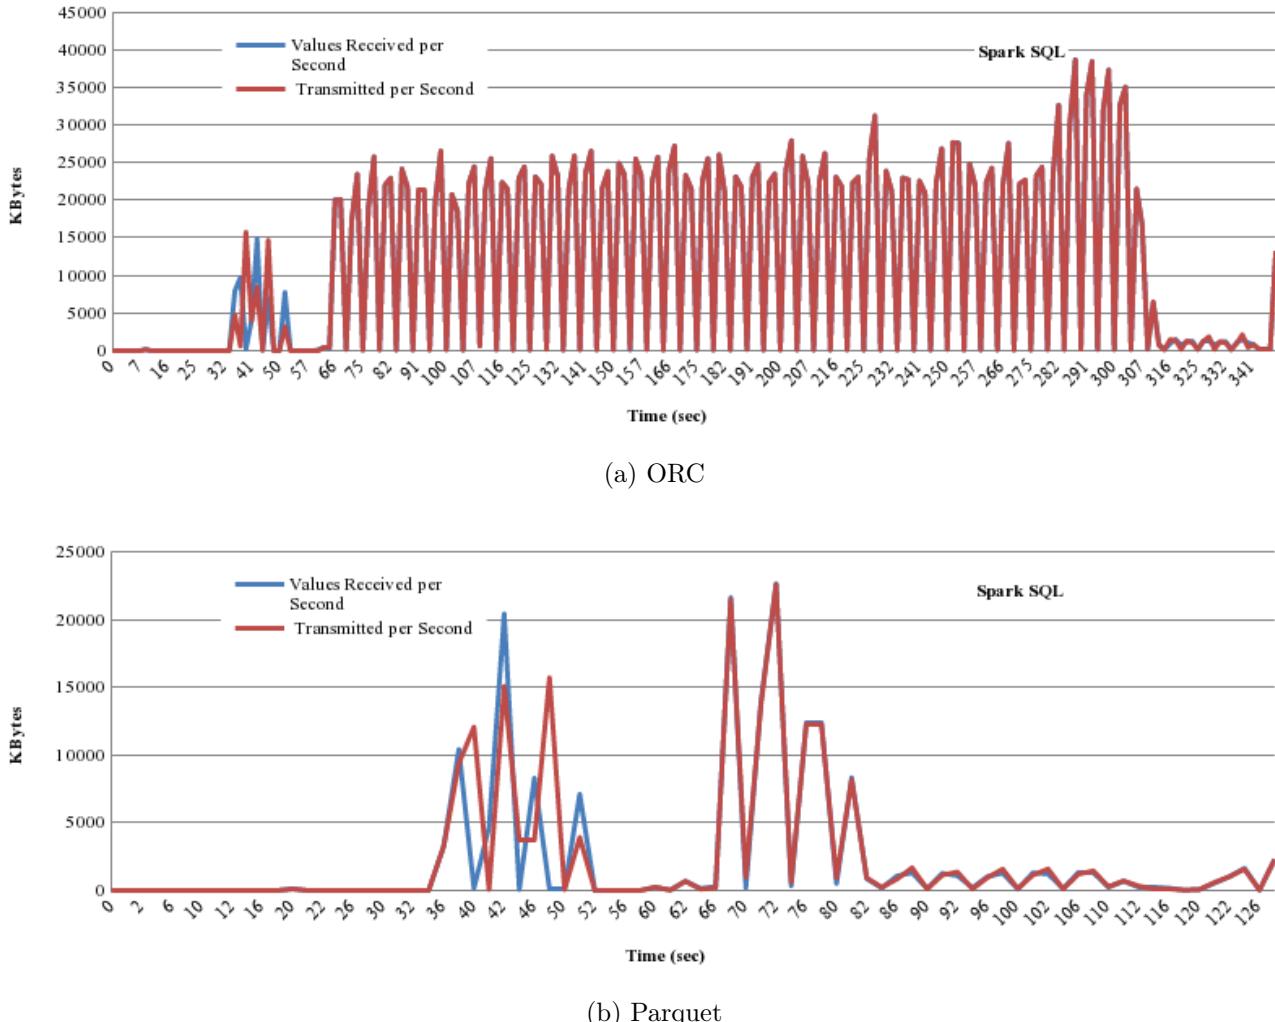


Figure 59: Network Utilization: more or less corresponding to DiskIO BW.

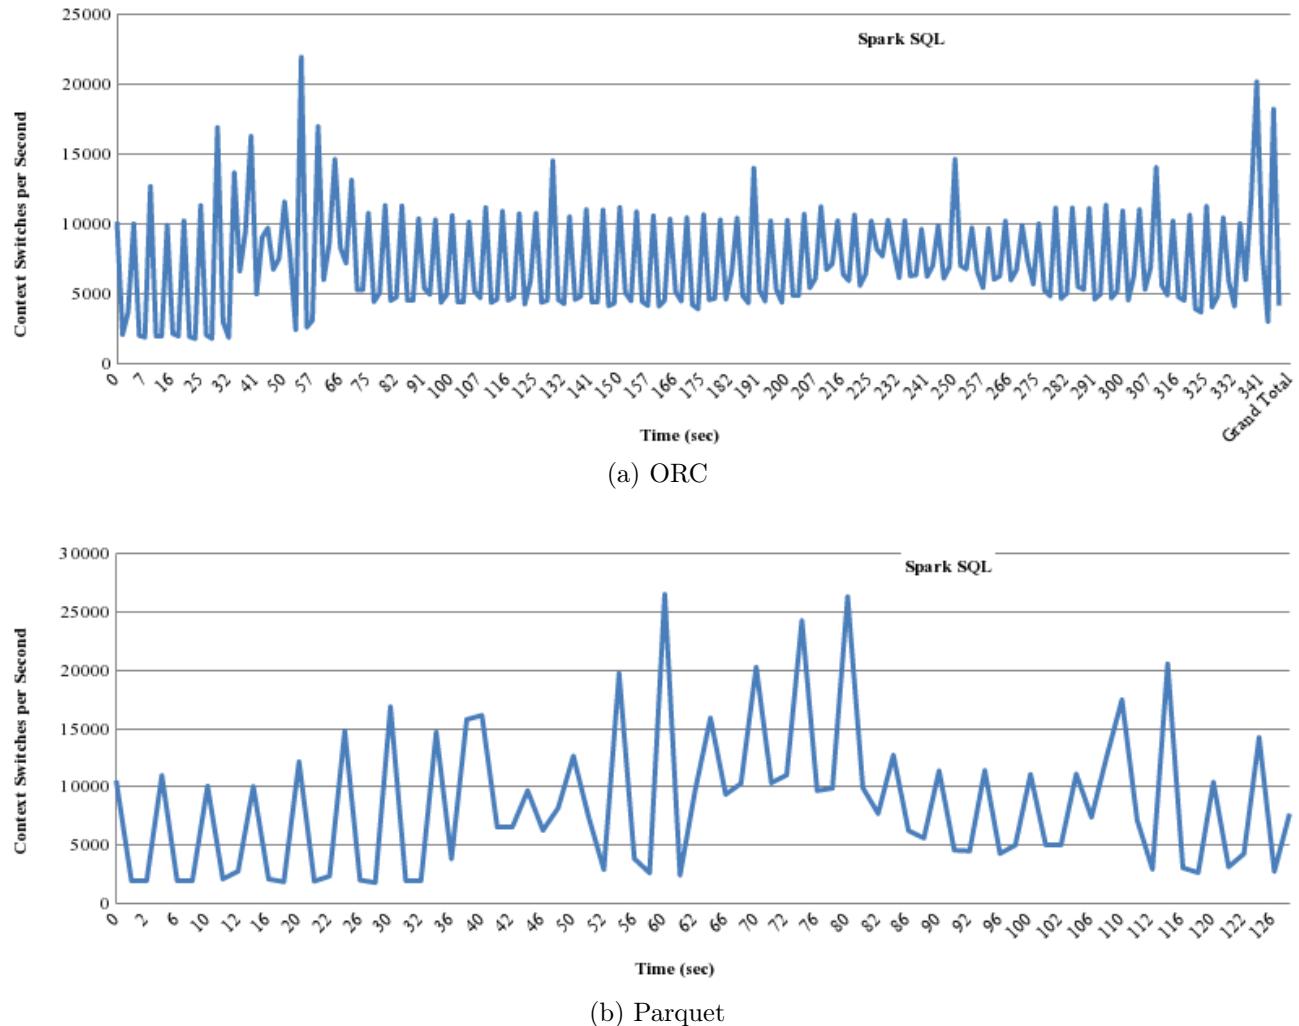


Figure 60: VMSTAT Context Switches

3.9 ORC - NoCompression vs. Snappy

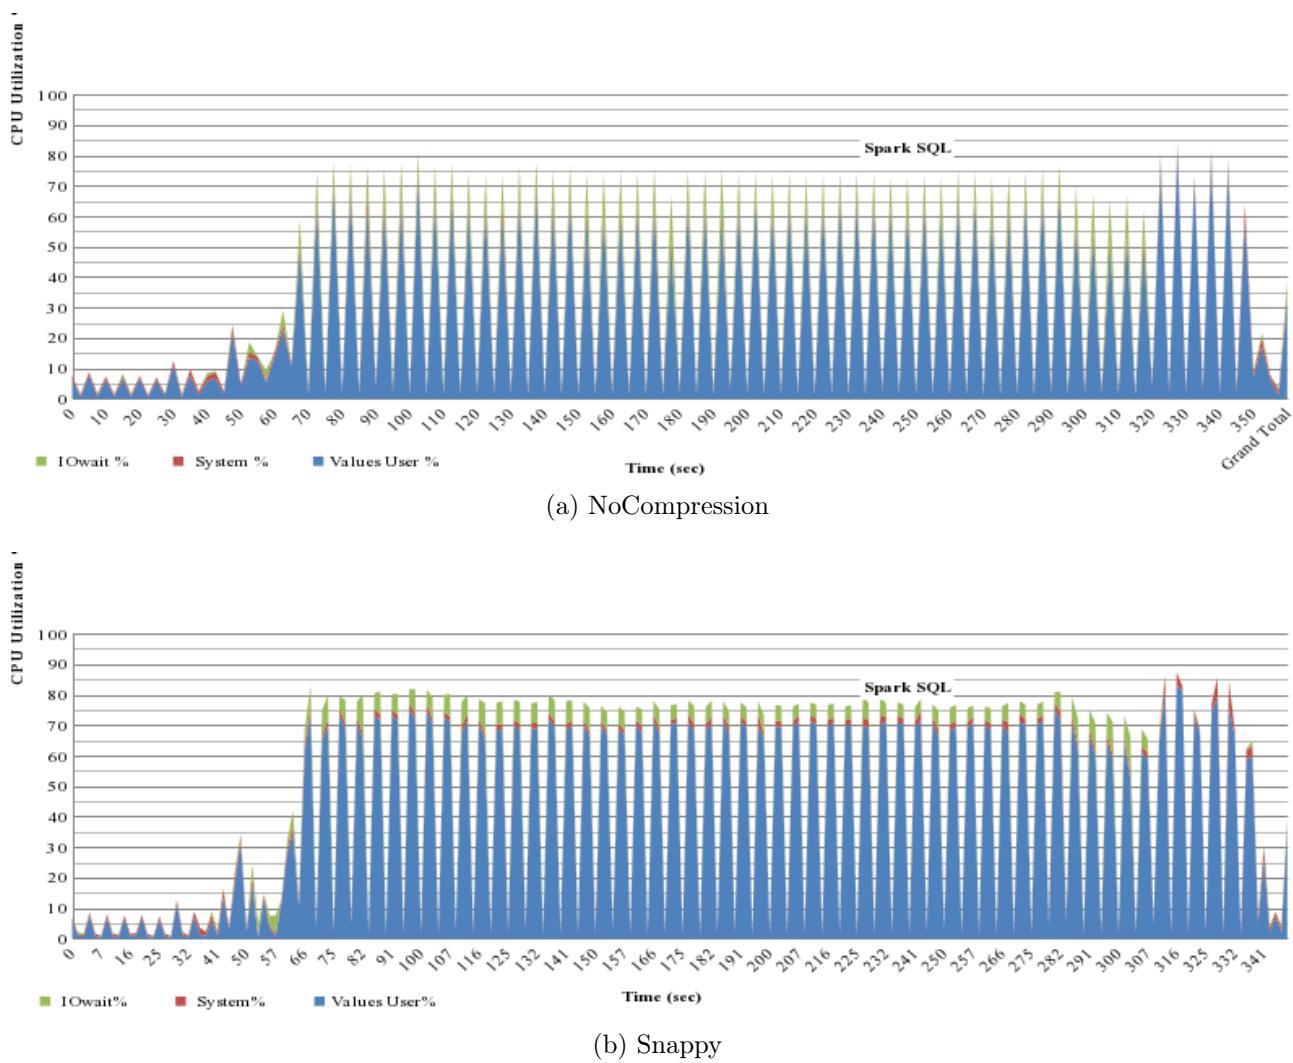


Figure 61: CPU.

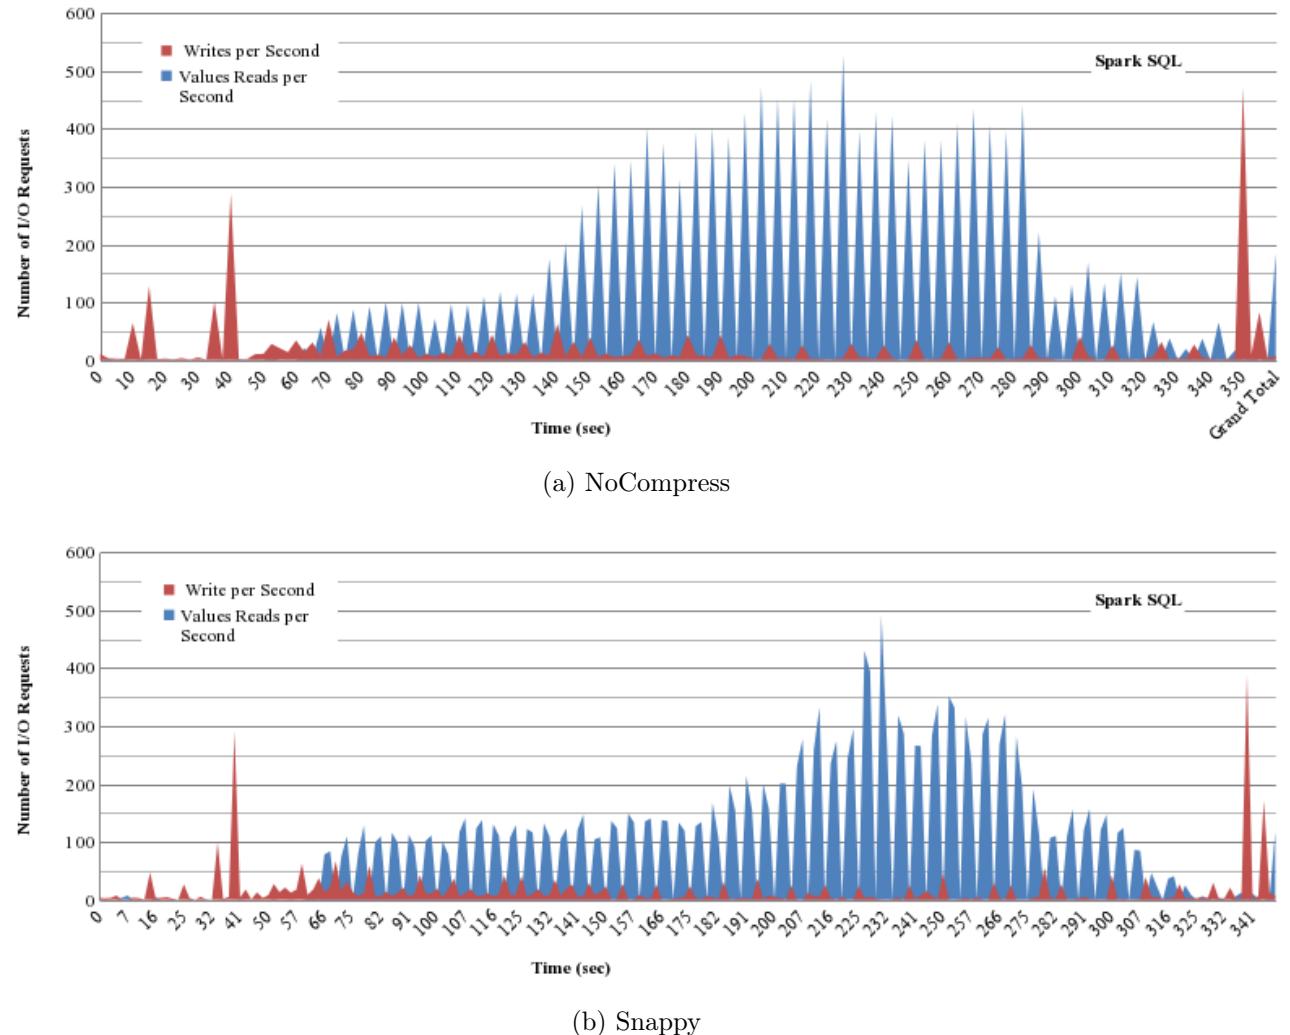


Figure 62: Disk Requests.

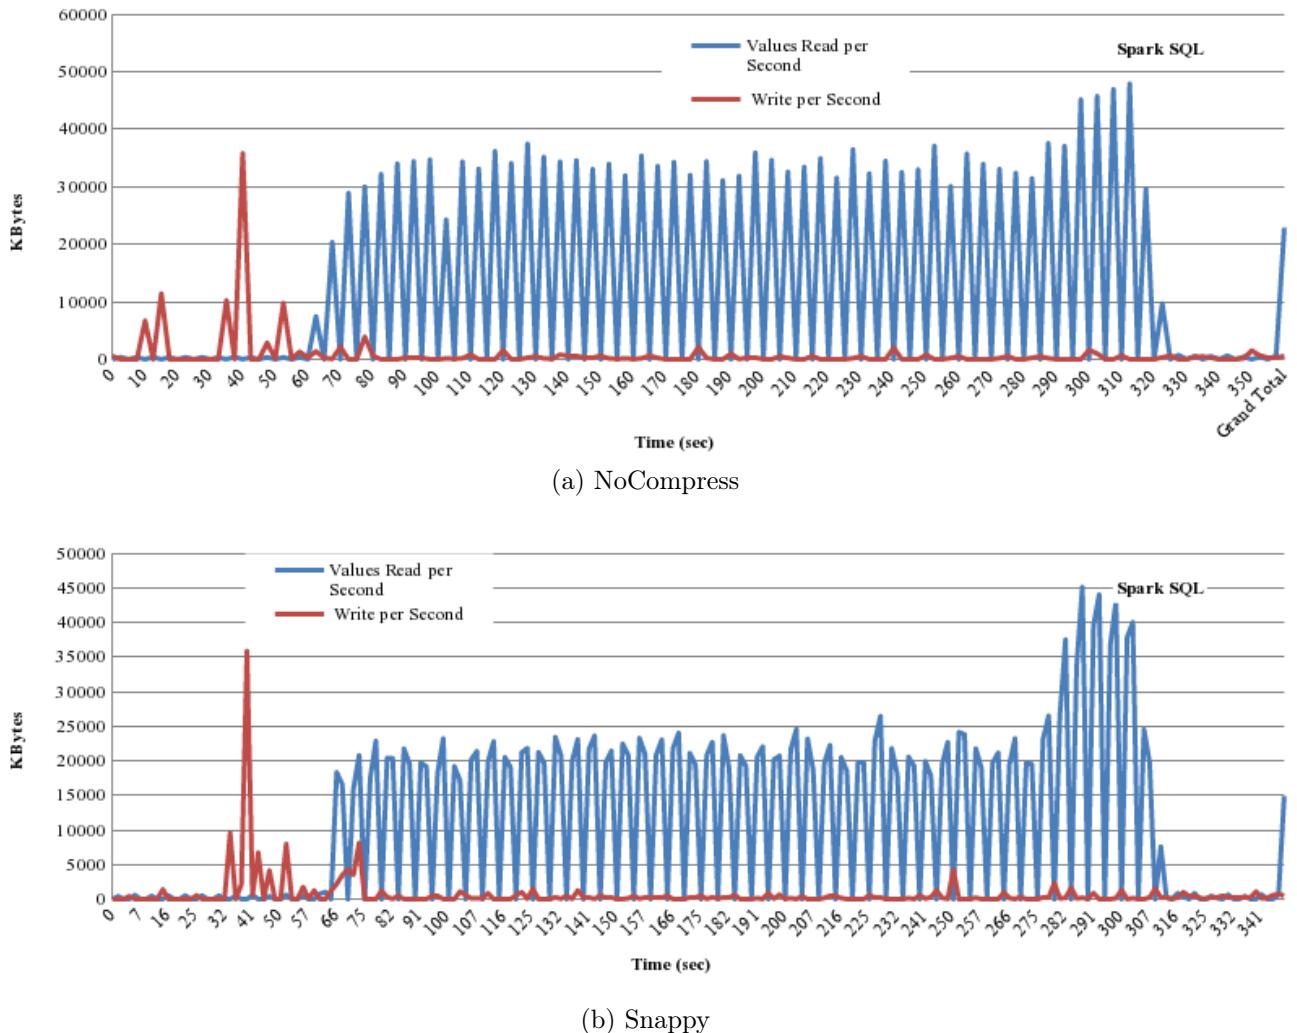


Figure 63: Disk Bandwidth.

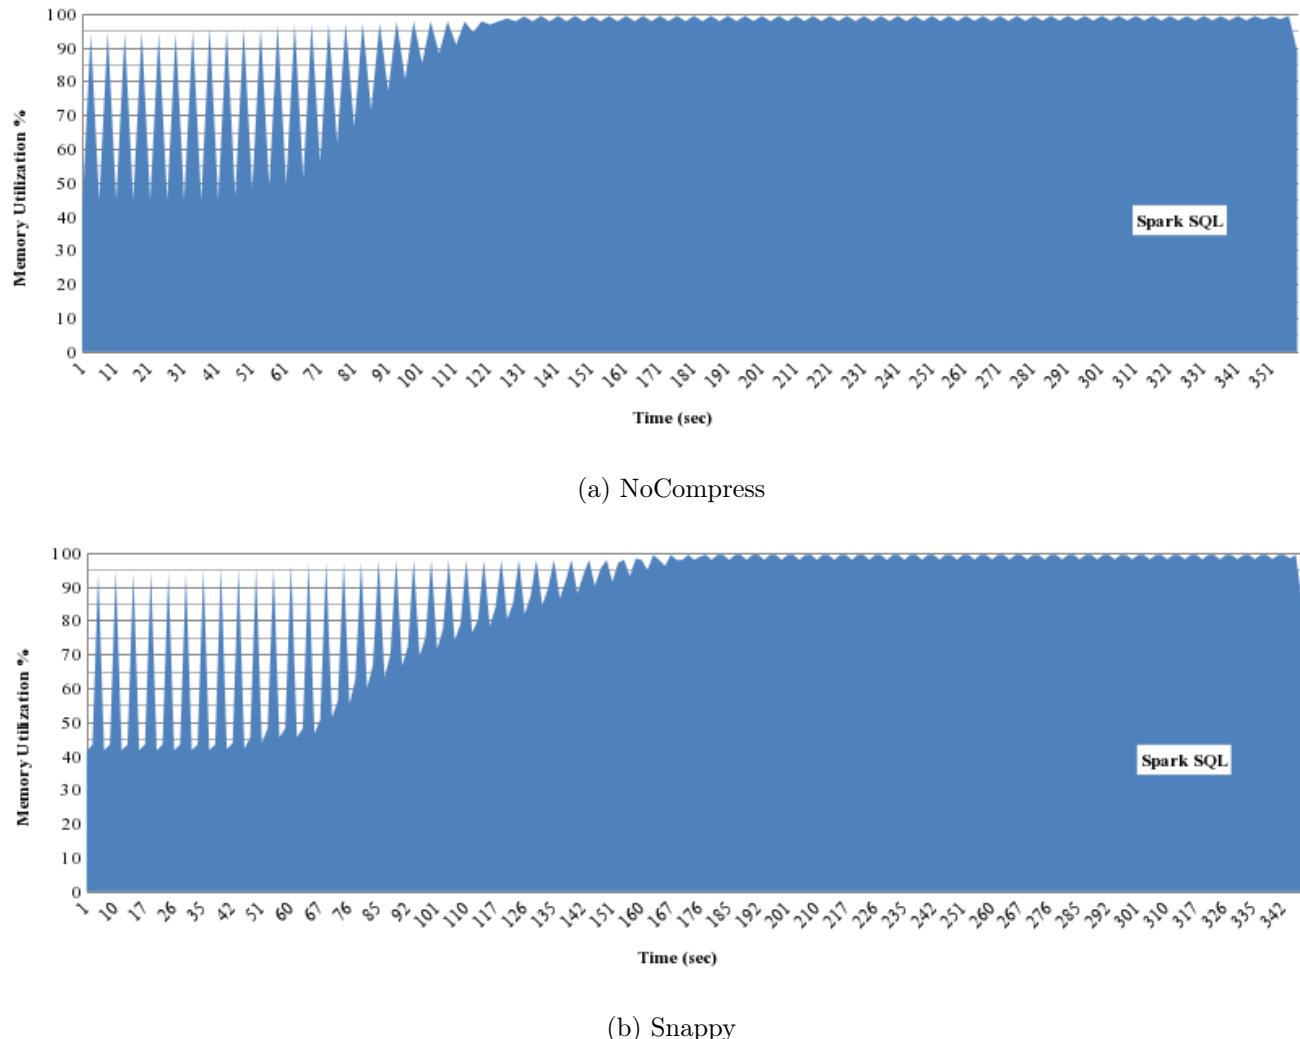


Figure 64: Free Memory.

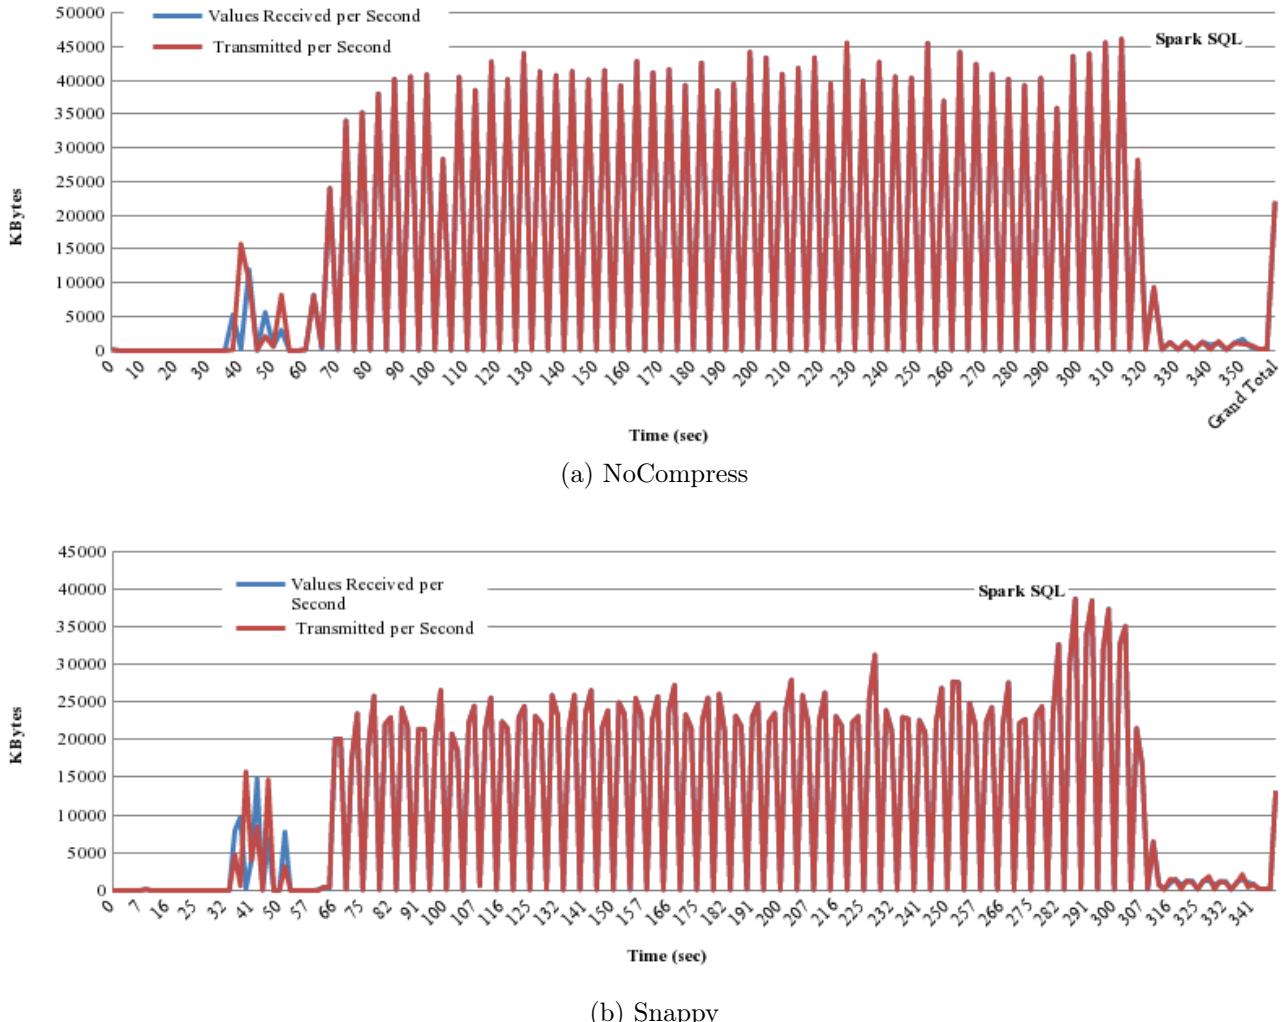


Figure 65: Network Utilization.

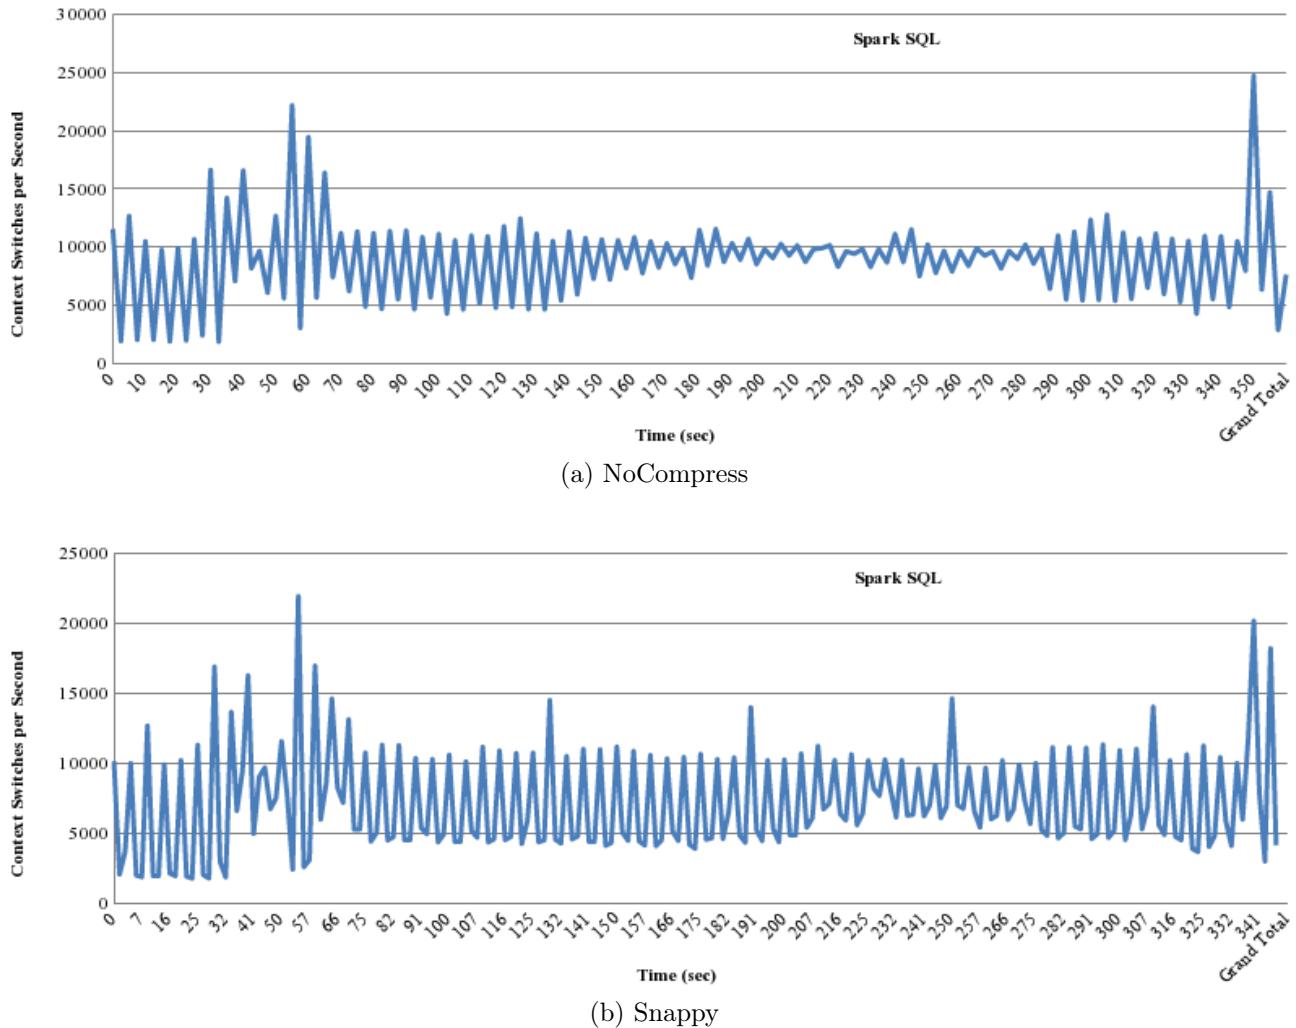


Figure 66: VMSTAT Context Switches

3.10 Parquet - NoCompression vs Snappy

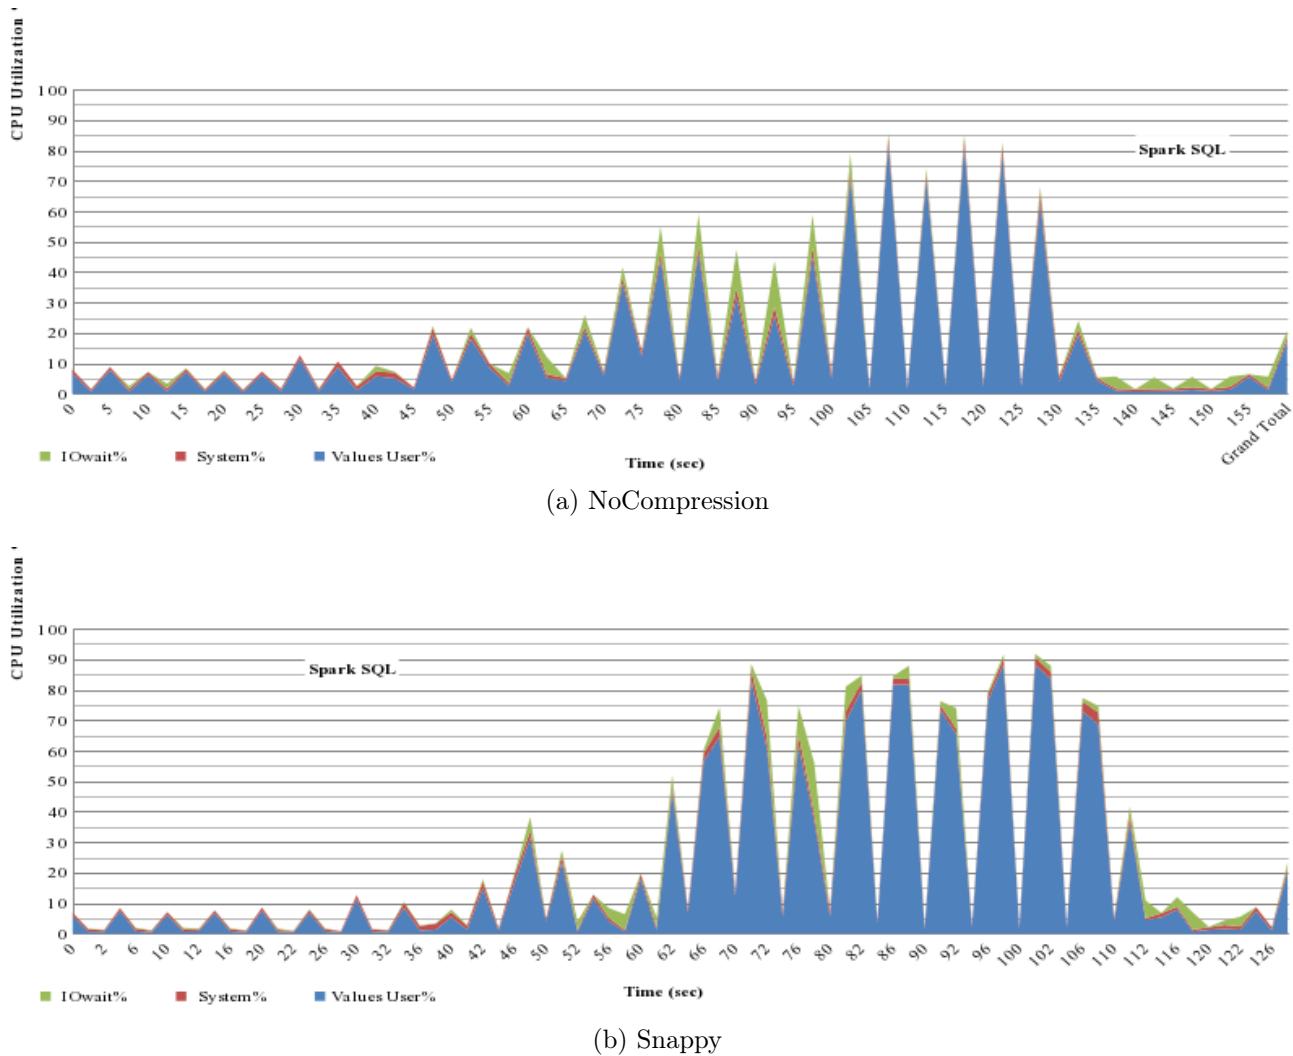


Figure 67: CPU.

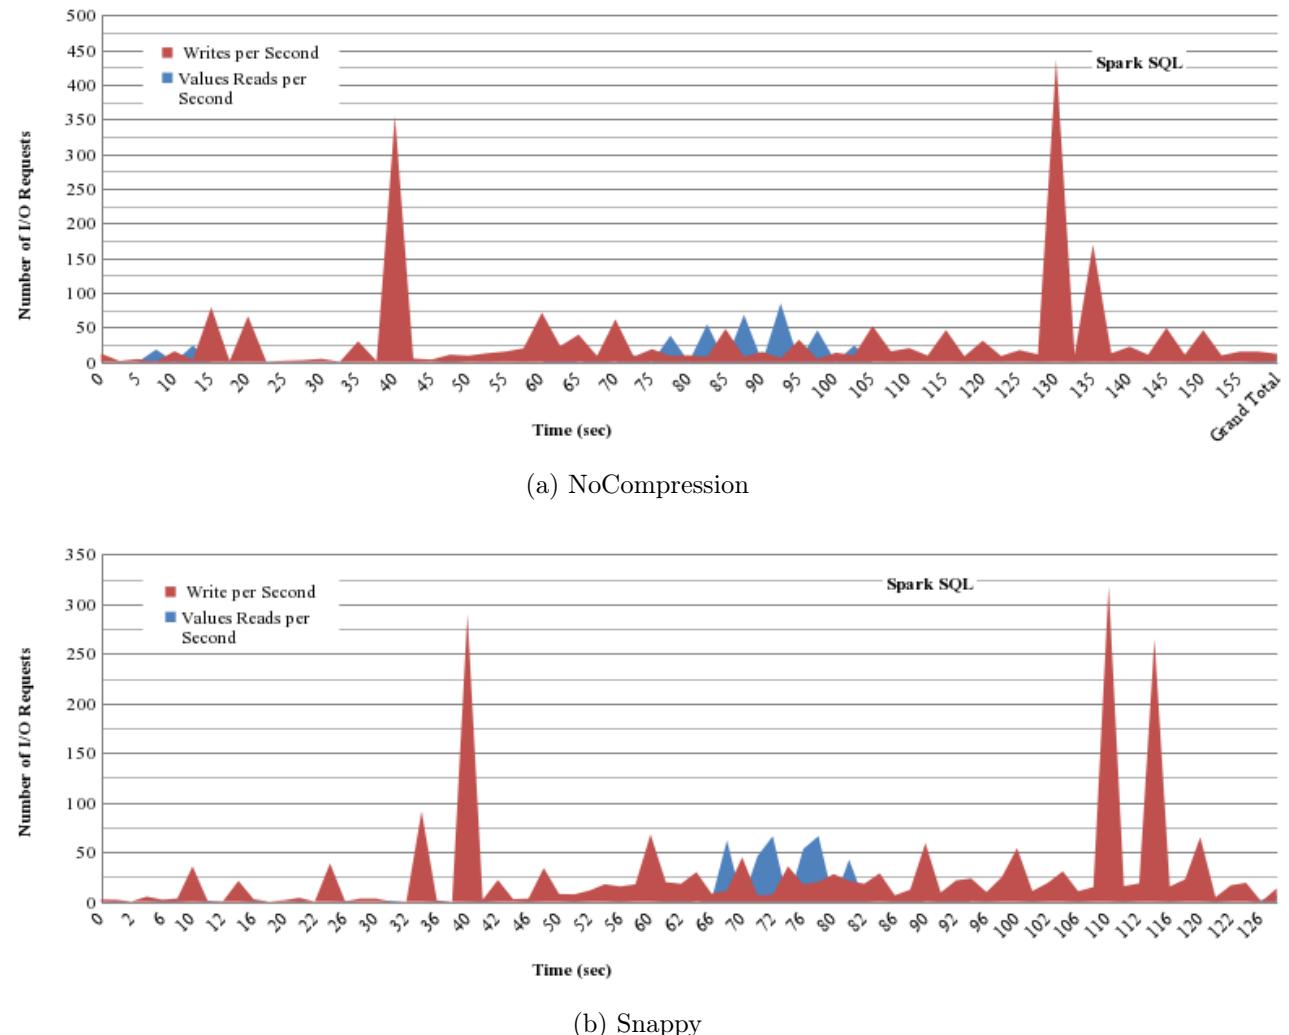


Figure 68: Disk Requests.

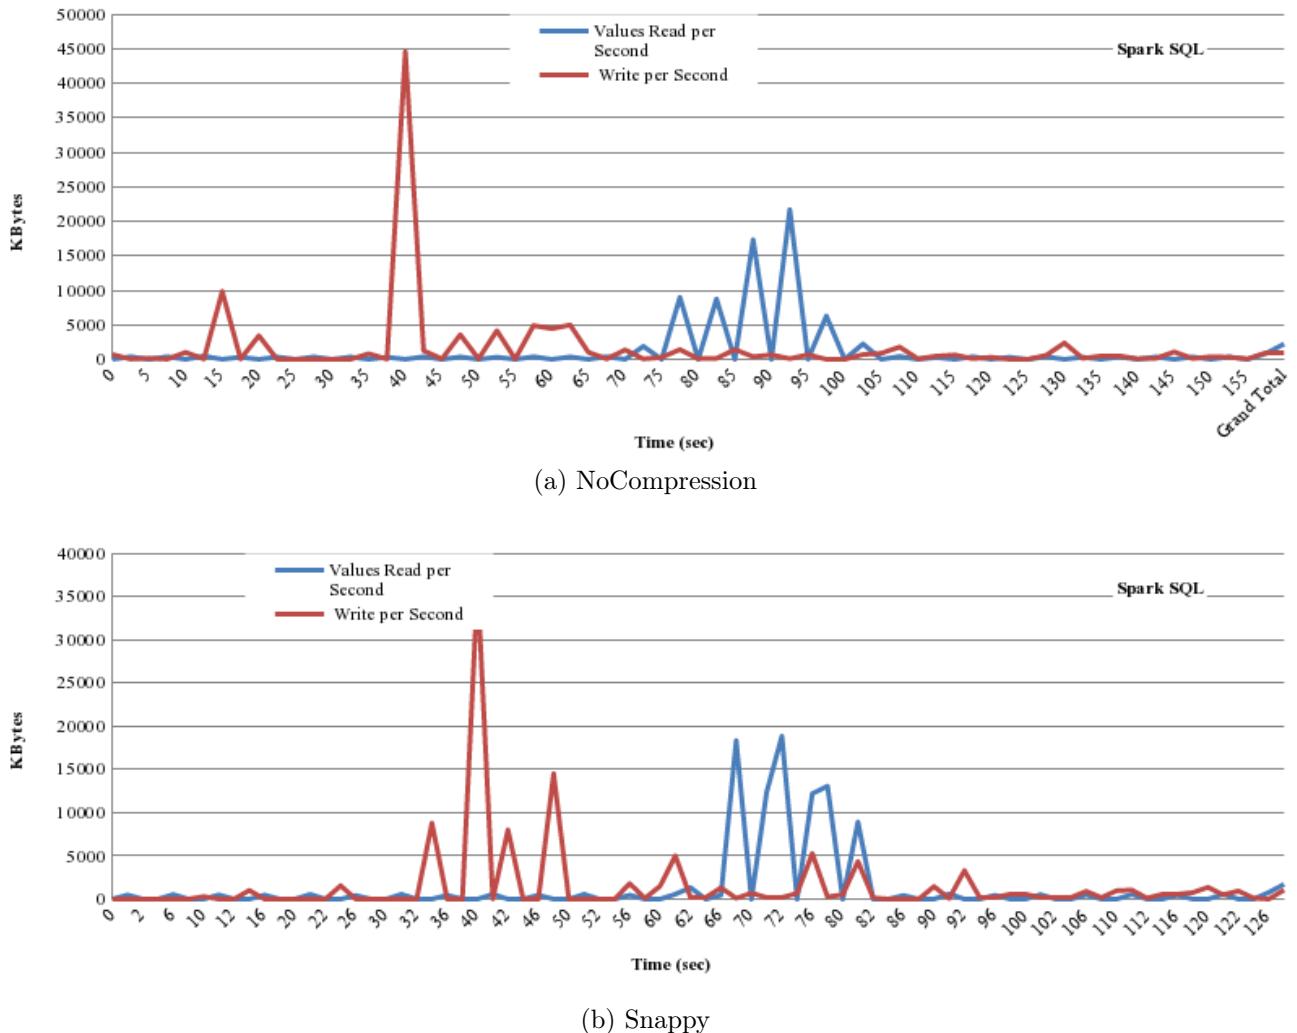


Figure 69: Disk Bandwidth.

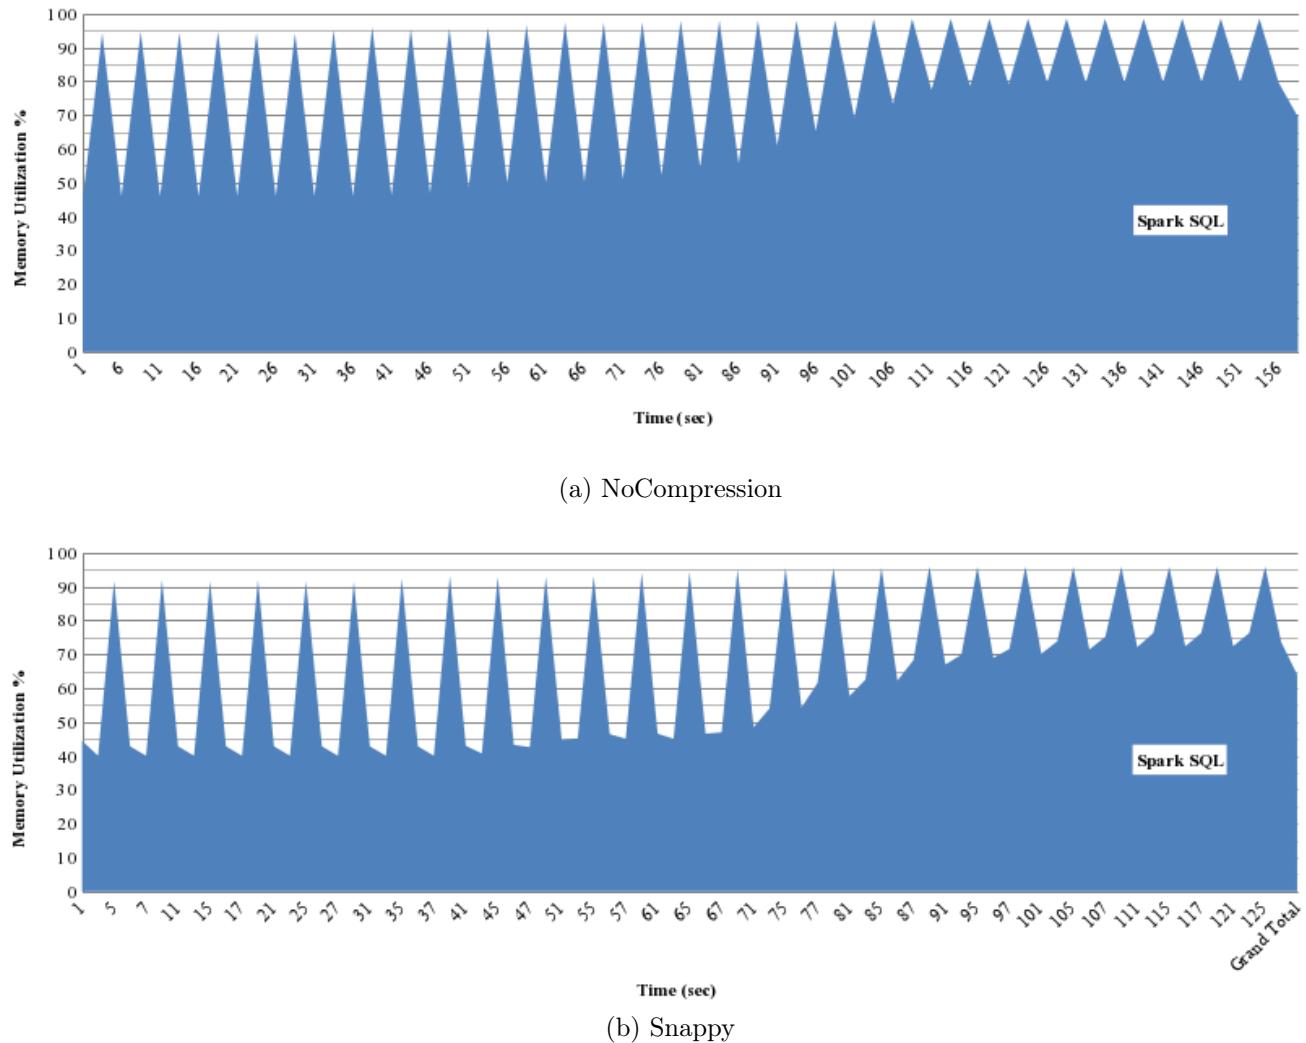


Figure 70: Free Memory.

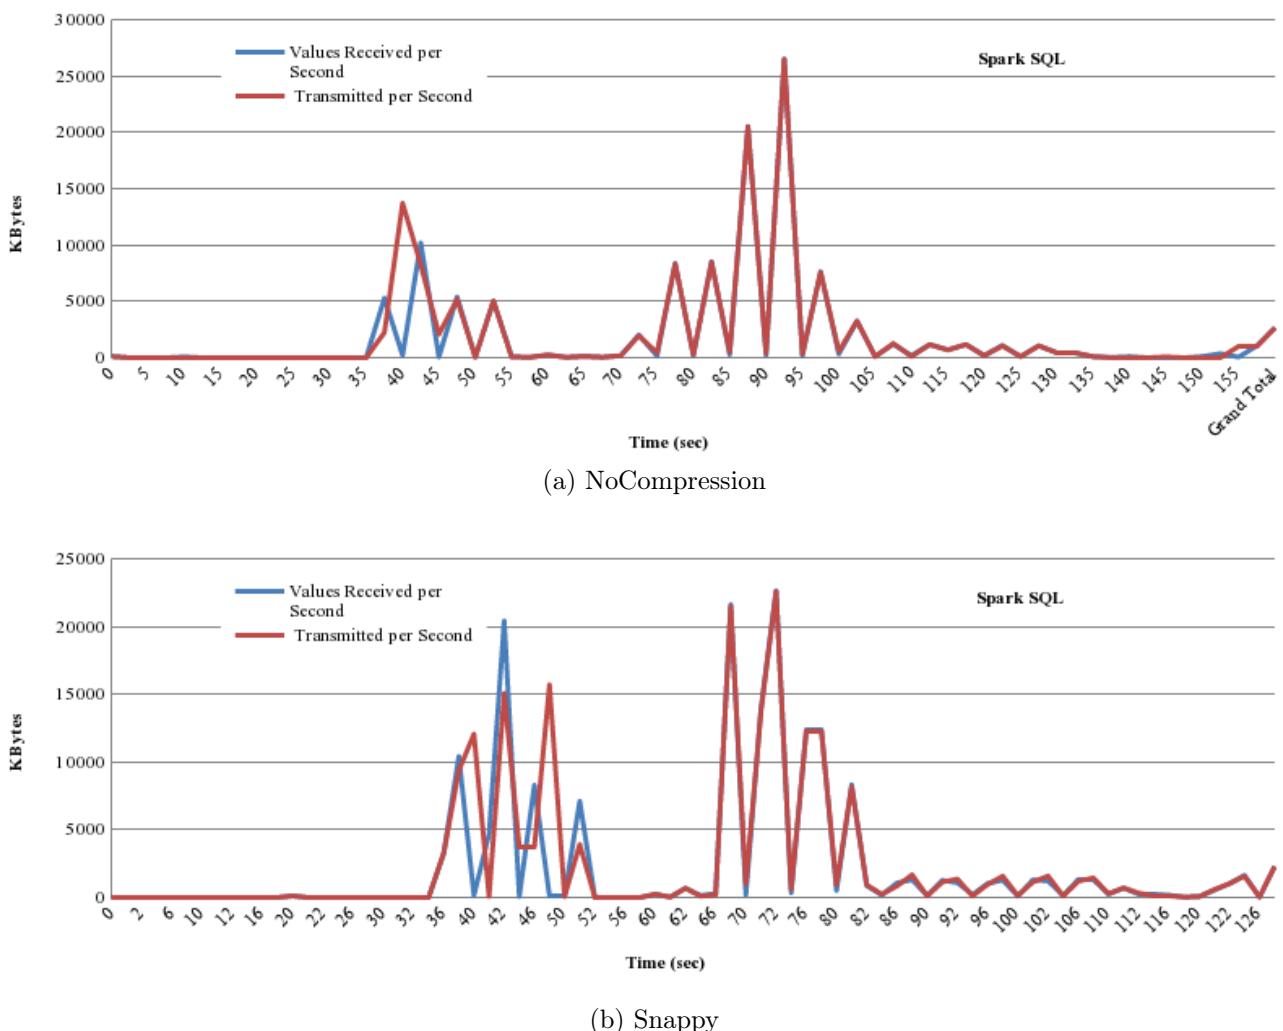


Figure 71: Network Utilization.

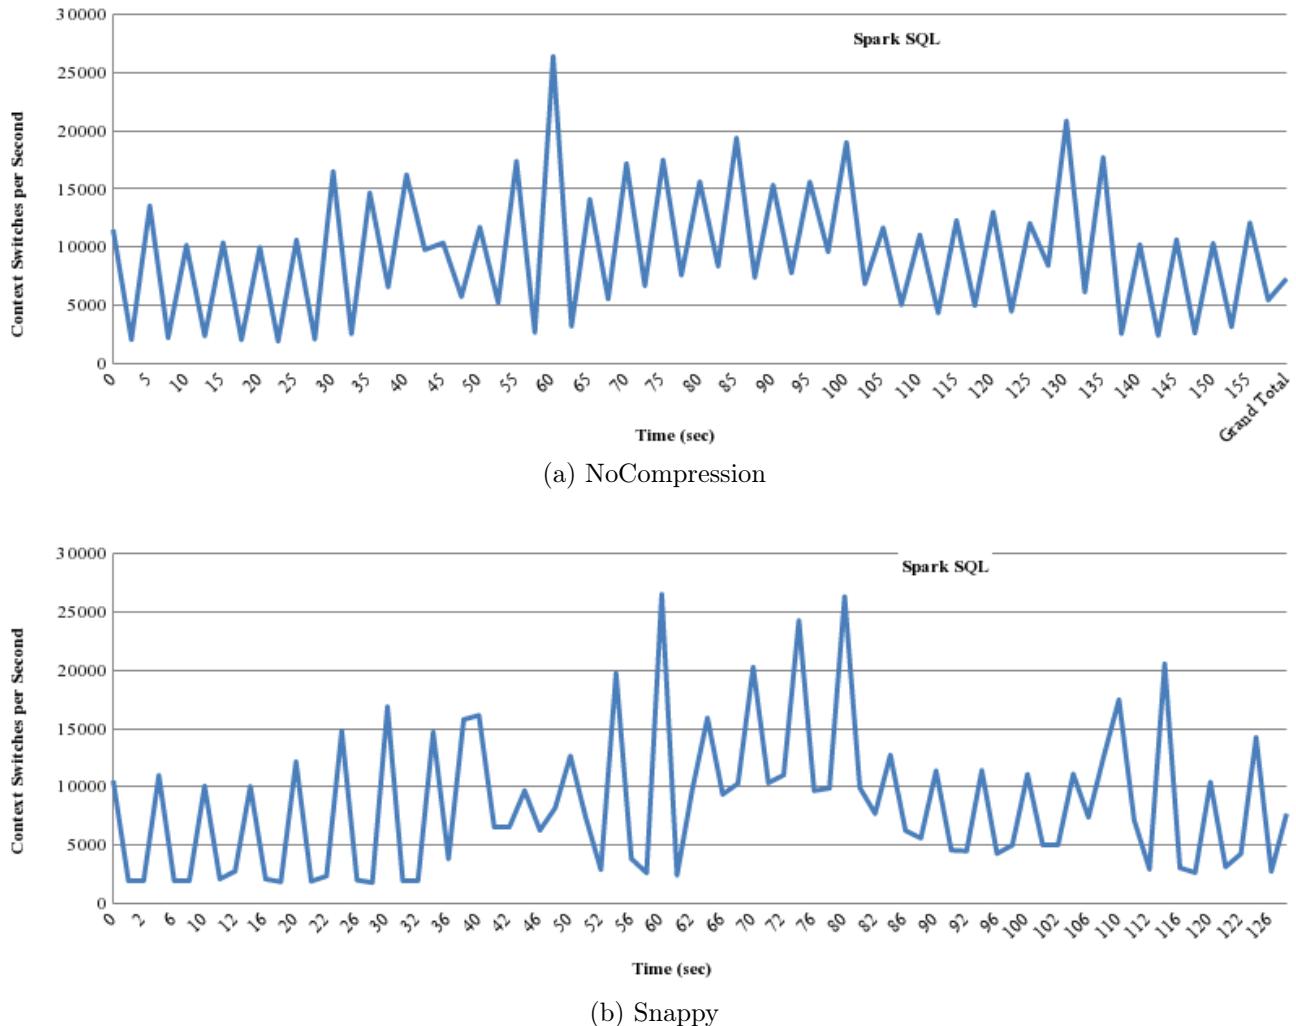


Figure 72: VMSTAT Context Switches

4 Q25

4.1 Type

HiveQL/Spark MLlib

4.2 Description

Customer segmentation analysis: Customers are separated along the following key shopping dimensions: recency of last visit, frequency of visits and monetary amount. Use the store and online purchase data during a given year to compute. After model of separation is build, report for the analysed customers to which "group" they where assigned

4.3 Performance

	Default (sec.)		No compr. (sec.)		Snappy (sec.)	
Query	ORC	Parquet	ORC	Parquet	ORC	Parquet
Q25	462	359	453	363	440	346

4.4 Phase 1

Selects desired customers from store and web sales and puts them in a TEMP table.

4.5 Phase 2

Users are grouped in a second TEMP table with the correct format value for the ML algorithm.

4.6 Phase 3

The KMeansClustering Spark class runs with the aforementioned table as input. The final result is written in a new table.

4.7 Source

GitHub: <https://git.io/vpFve>

4.8 Notes

...

4.9 No compression - ORC vs Parquet

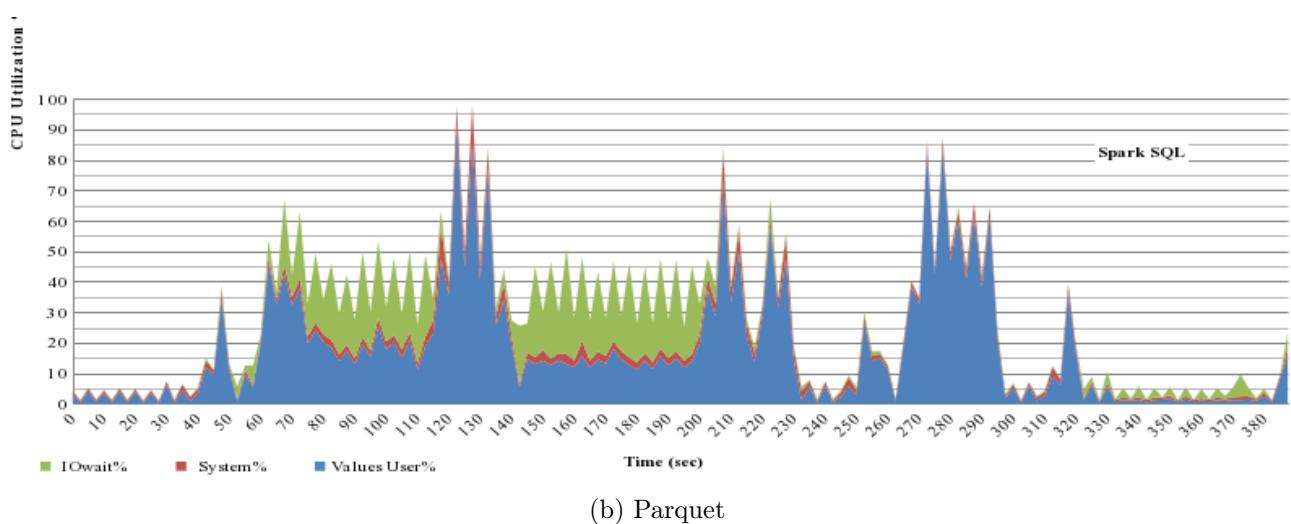
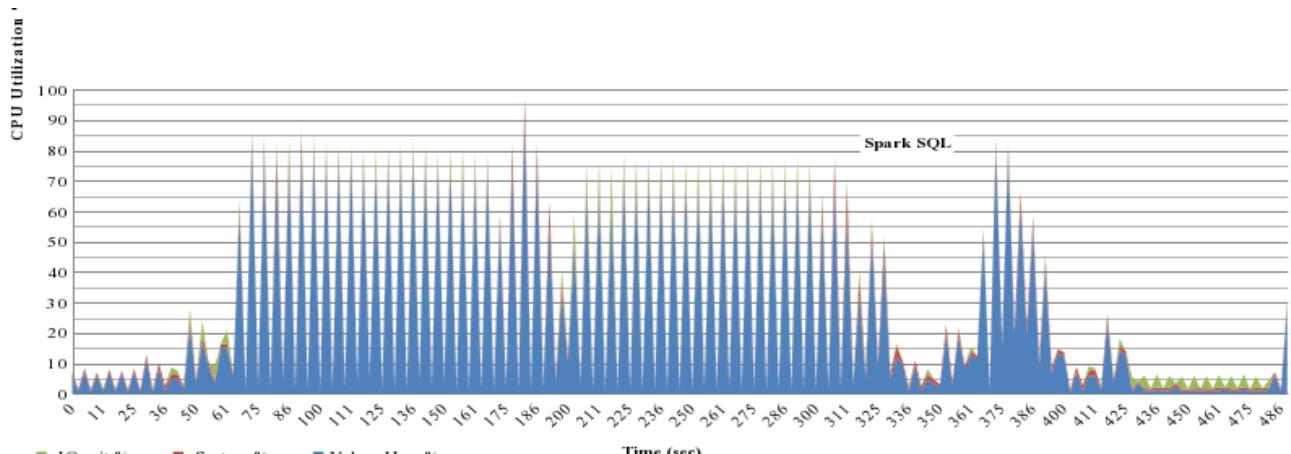


Figure 73: CPU: Lower utilization on Parquet (max is 80%)

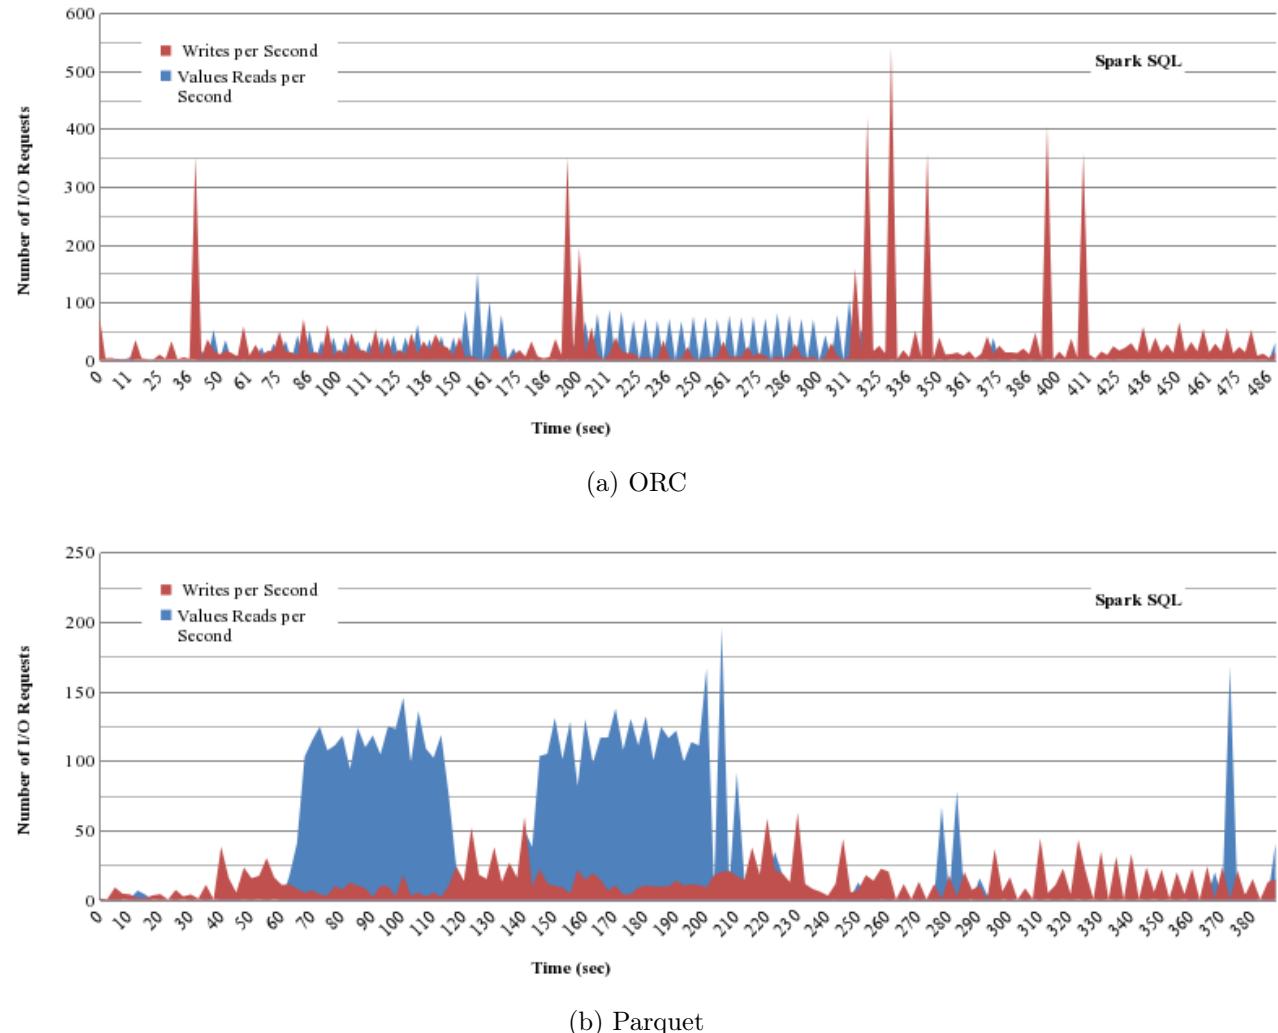


Figure 74: Disk Requests: More requests on Parquet.

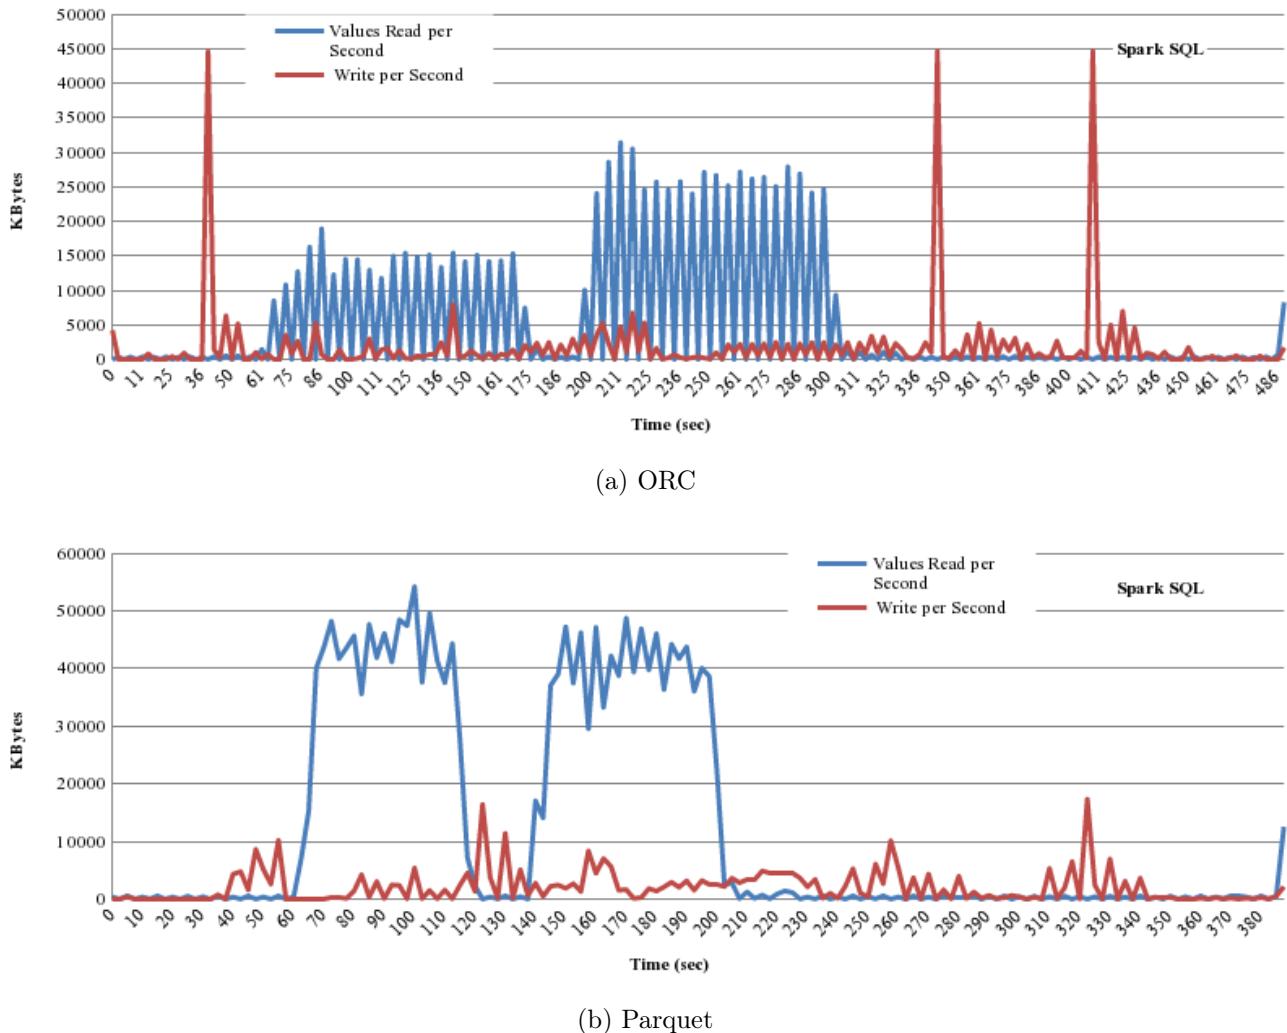


Figure 75: Disk Bandwidth: bandwidth utilization on Parquet.

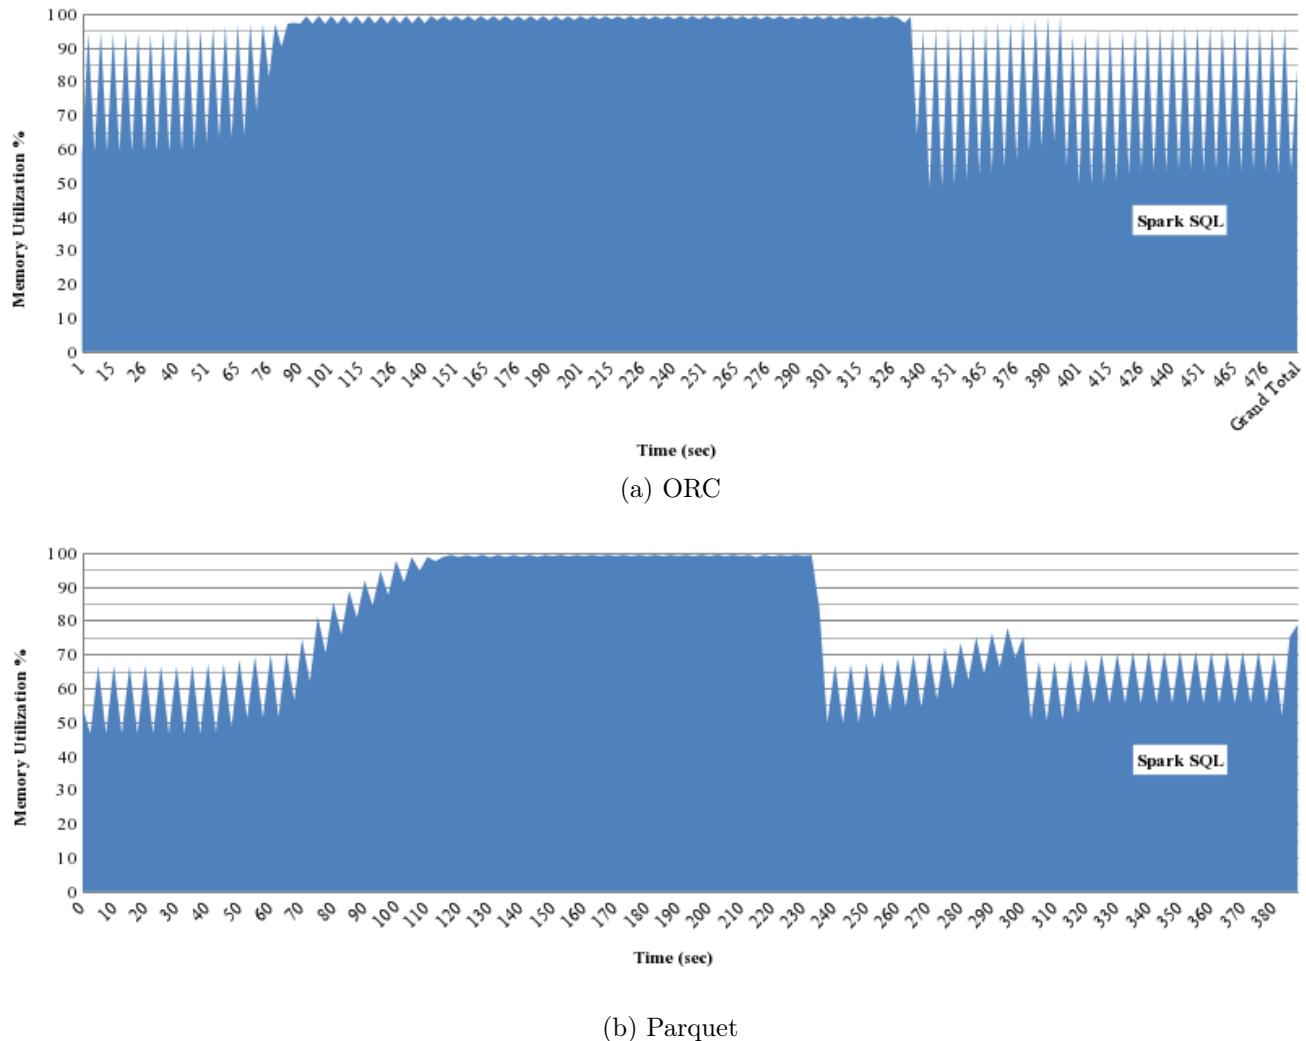


Figure 76: Free Memory.

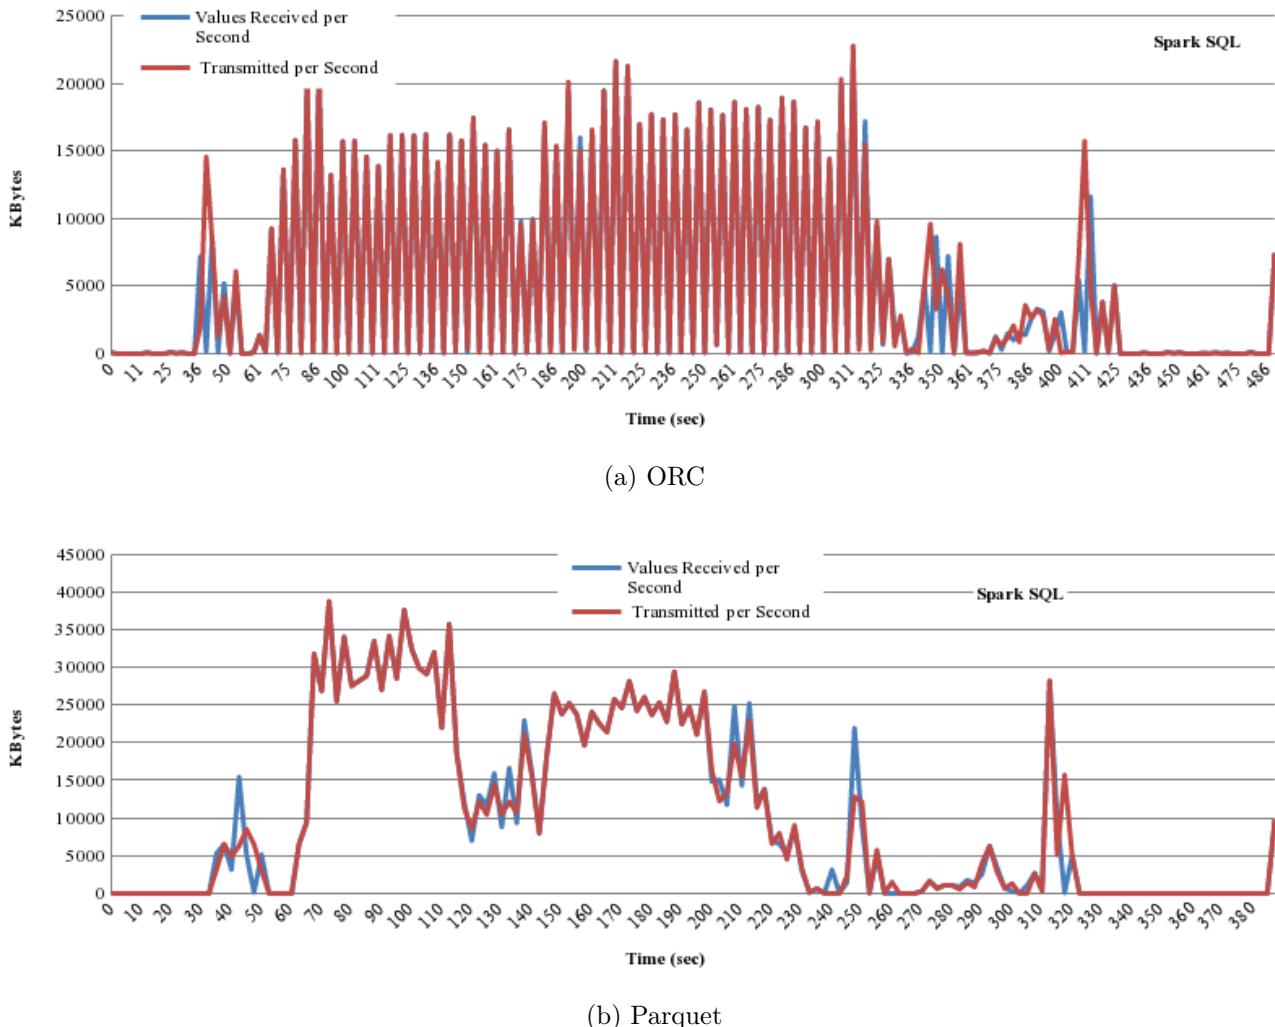


Figure 77: Network Utilization: More or less corresponding to DiskIO.

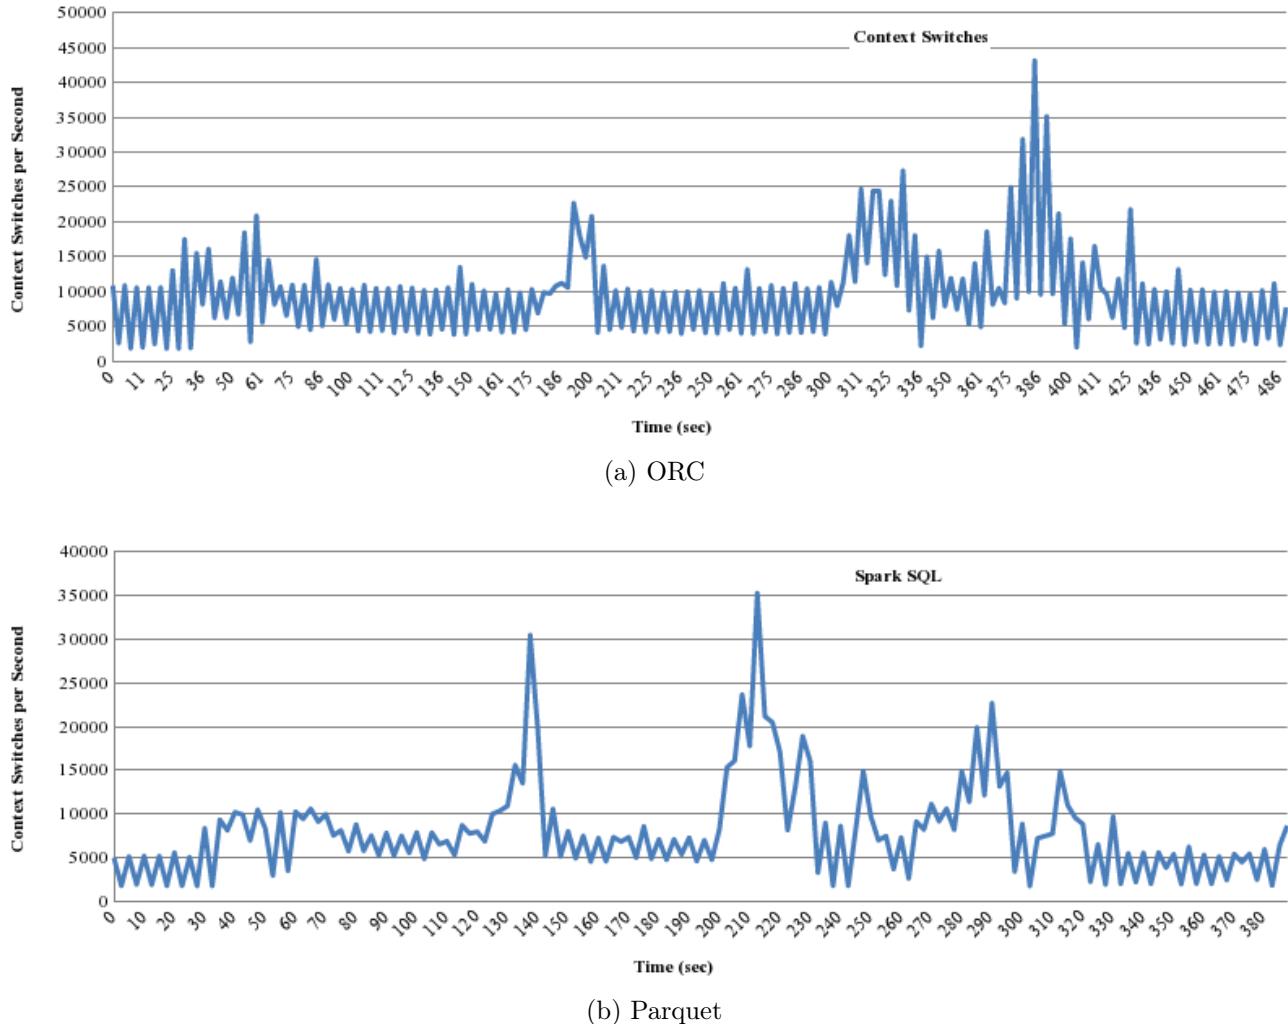


Figure 78: VMSTAT Context Switches

4.10 Snappy - ORC vs. Parquet

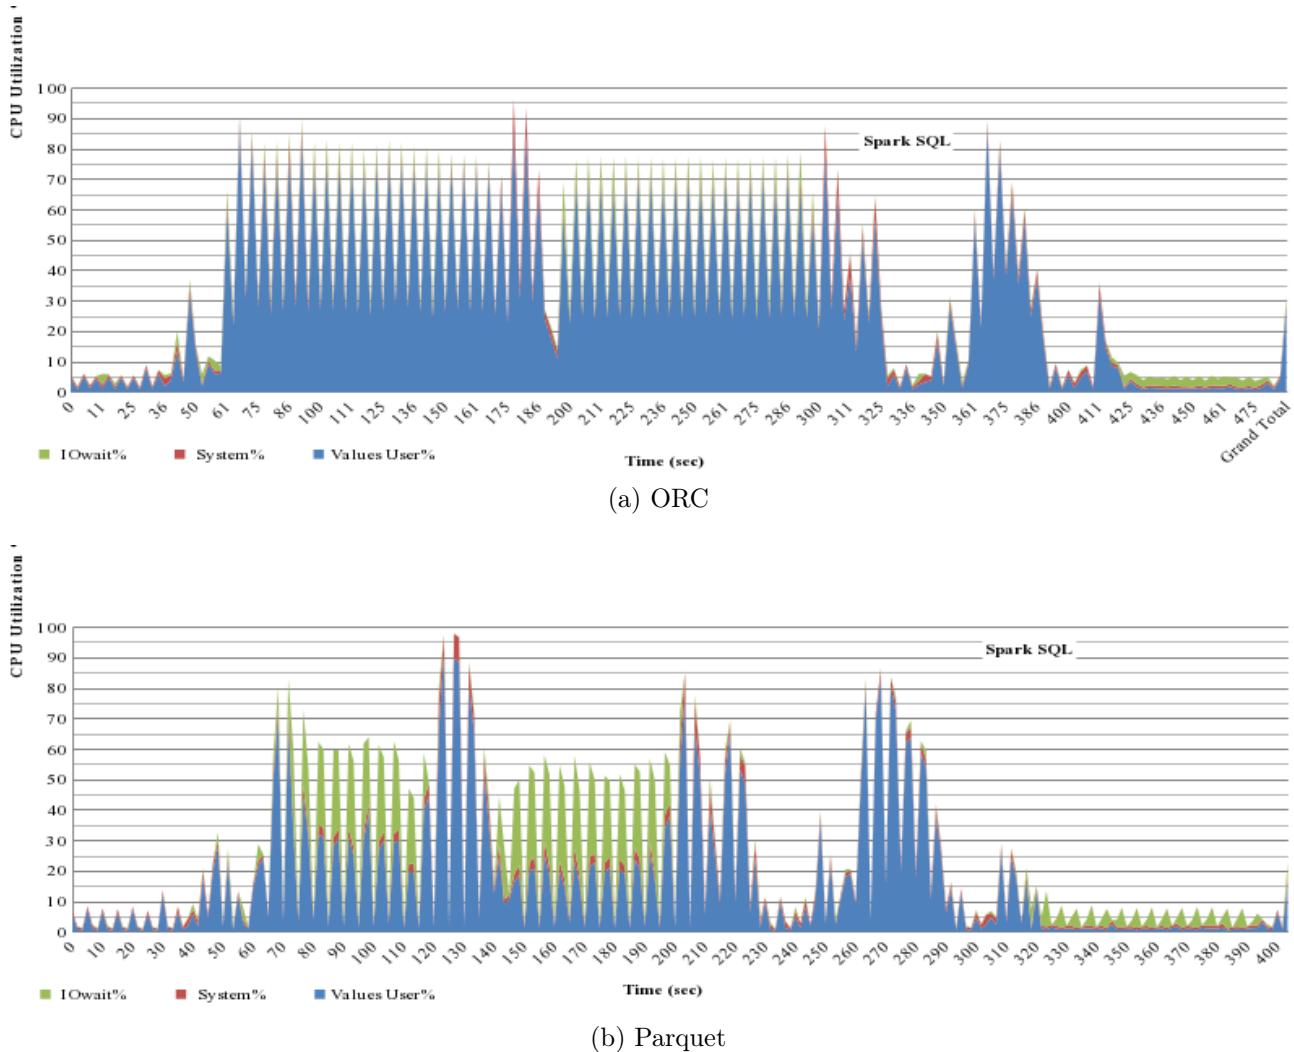


Figure 79: CPU: Low util on parquet (60%), short time for phase 1.

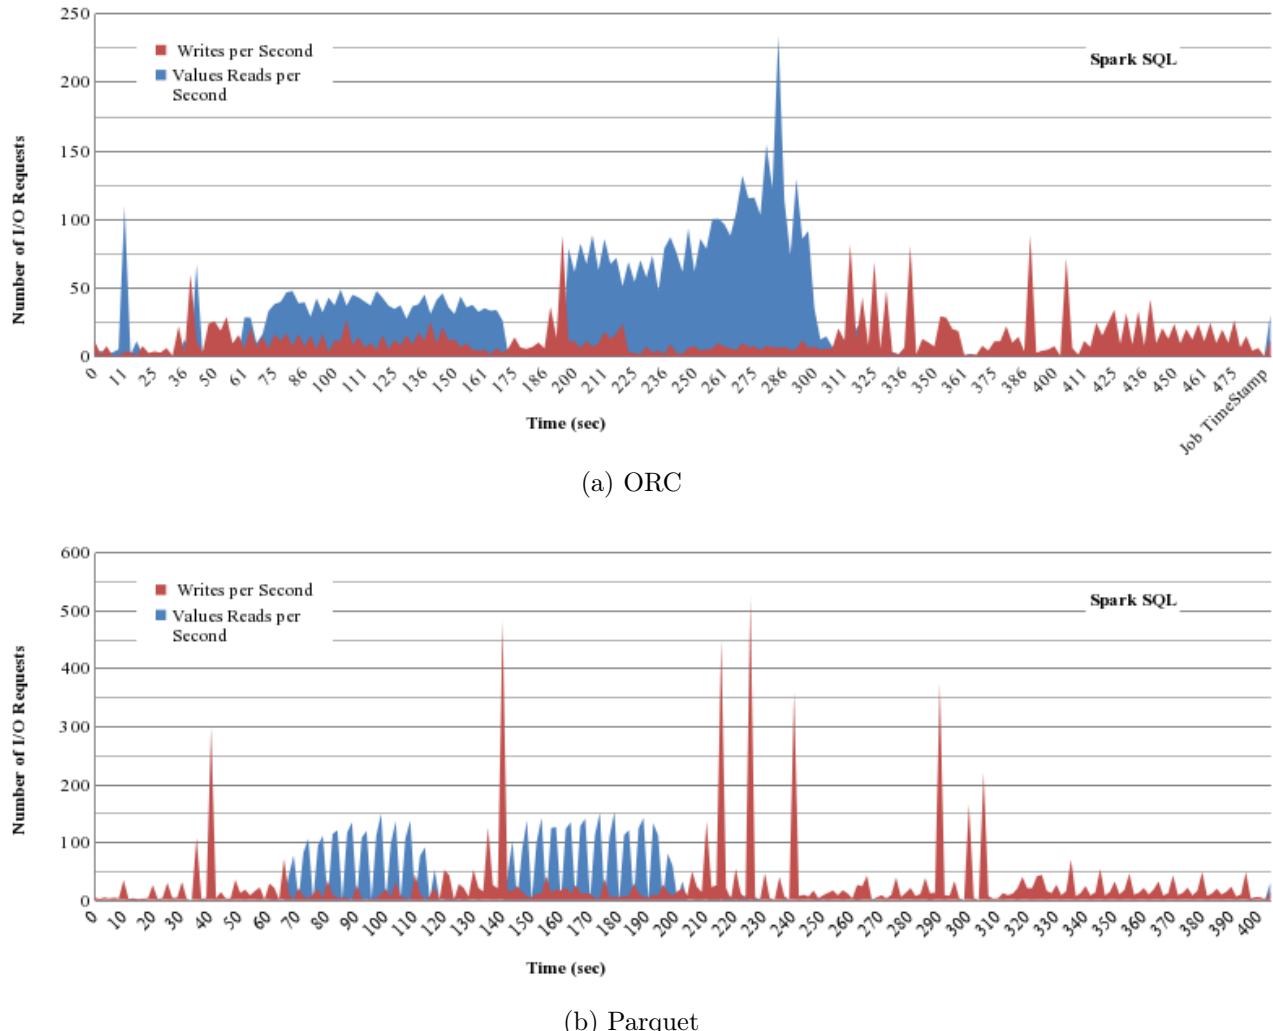


Figure 80: Disk Requests: Lower number of requests on Parquet.

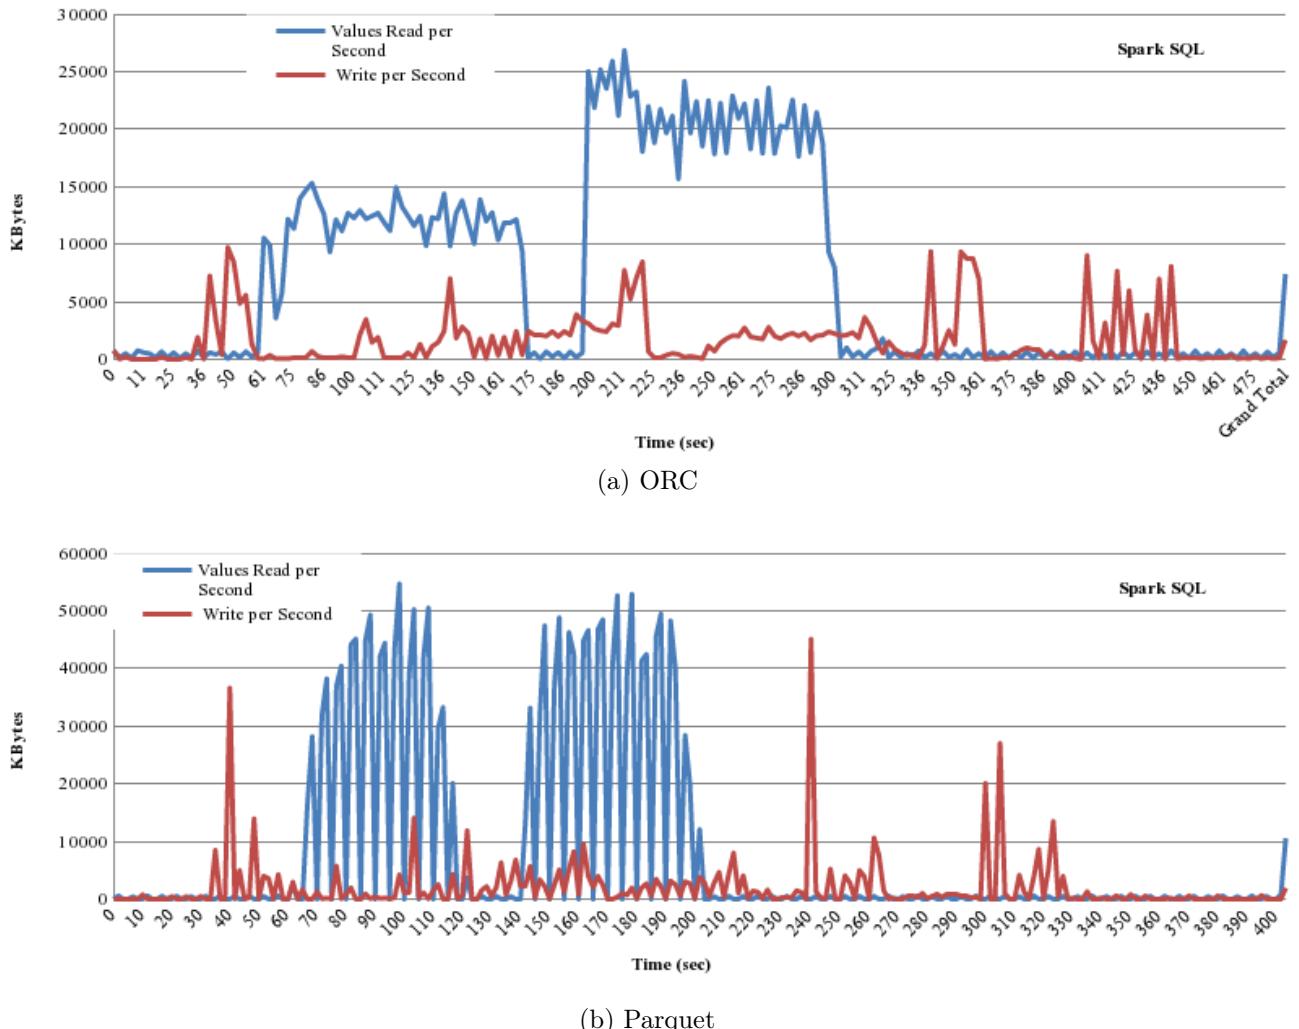


Figure 81: Disk Bandwidth: higher BW utilization on Parquet.

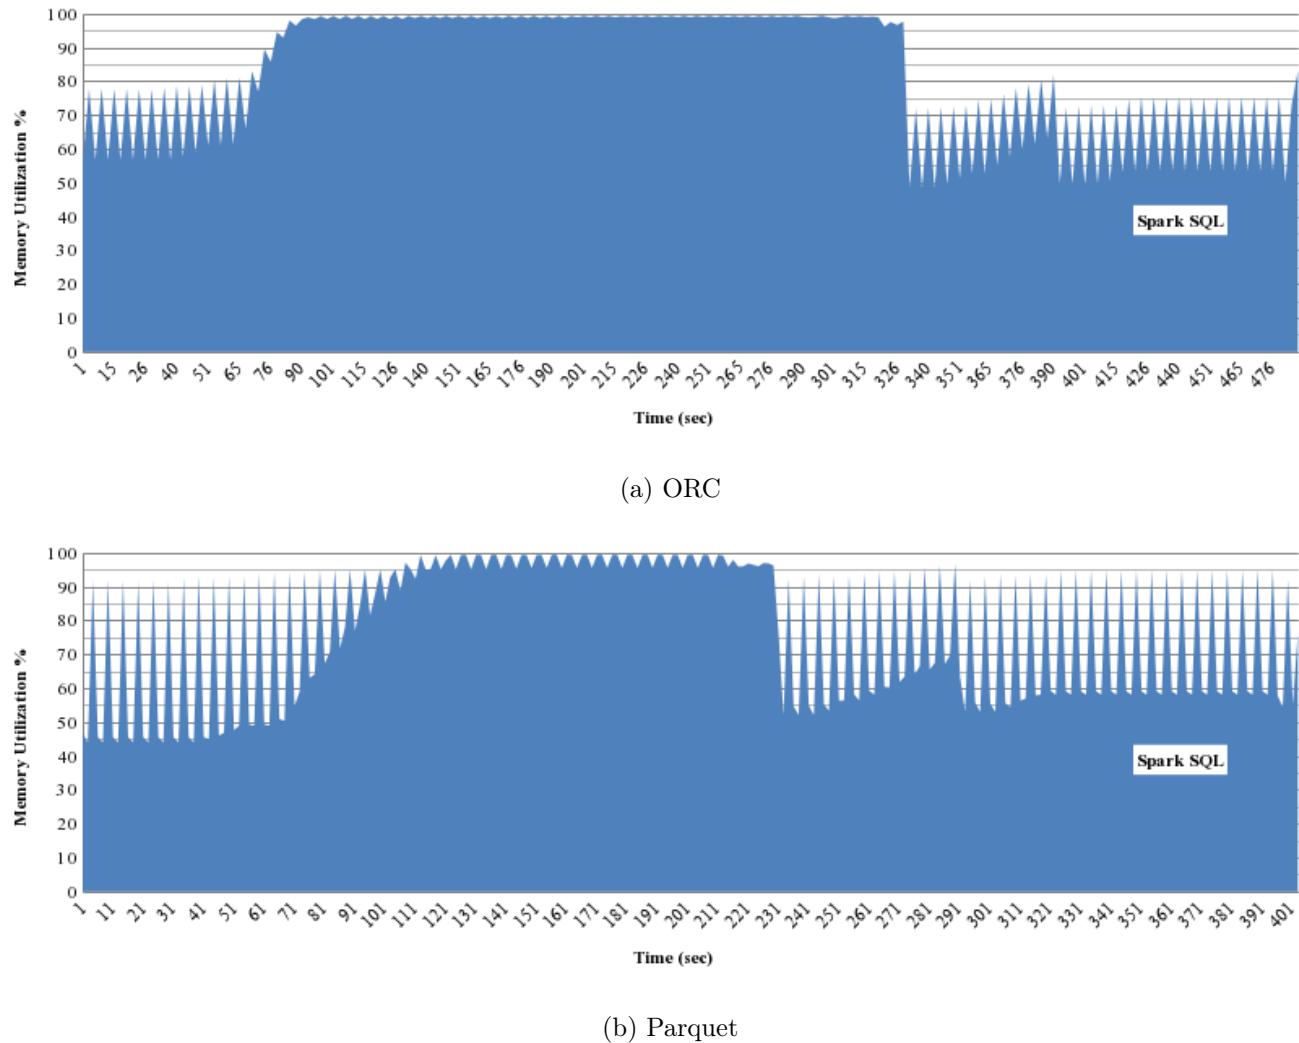


Figure 82: Free Memory.

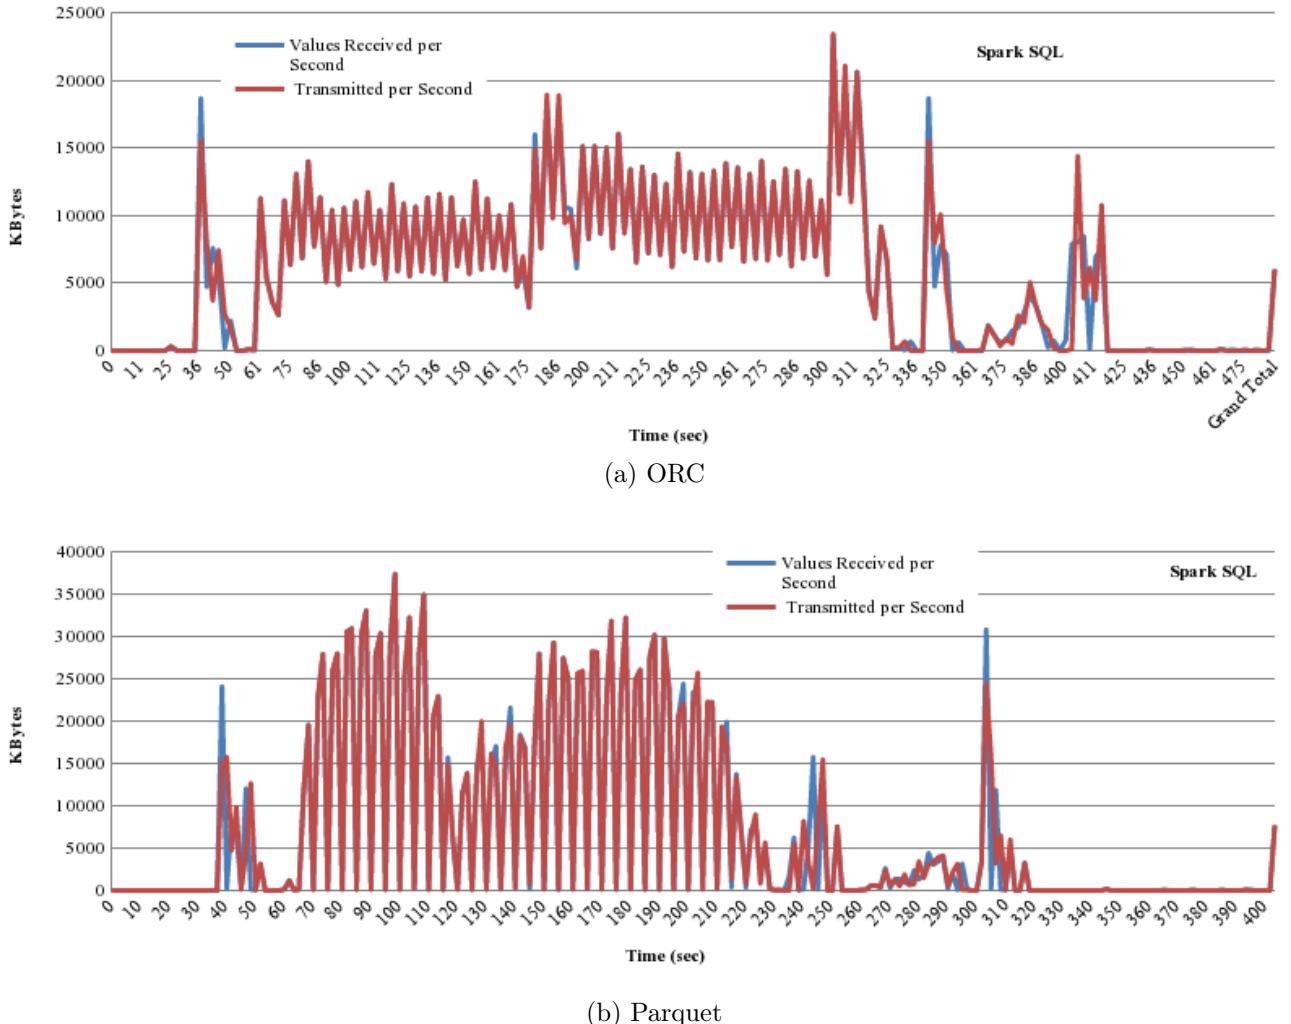


Figure 83: Network Utilization: more or less corresponding to DiskIO BW.

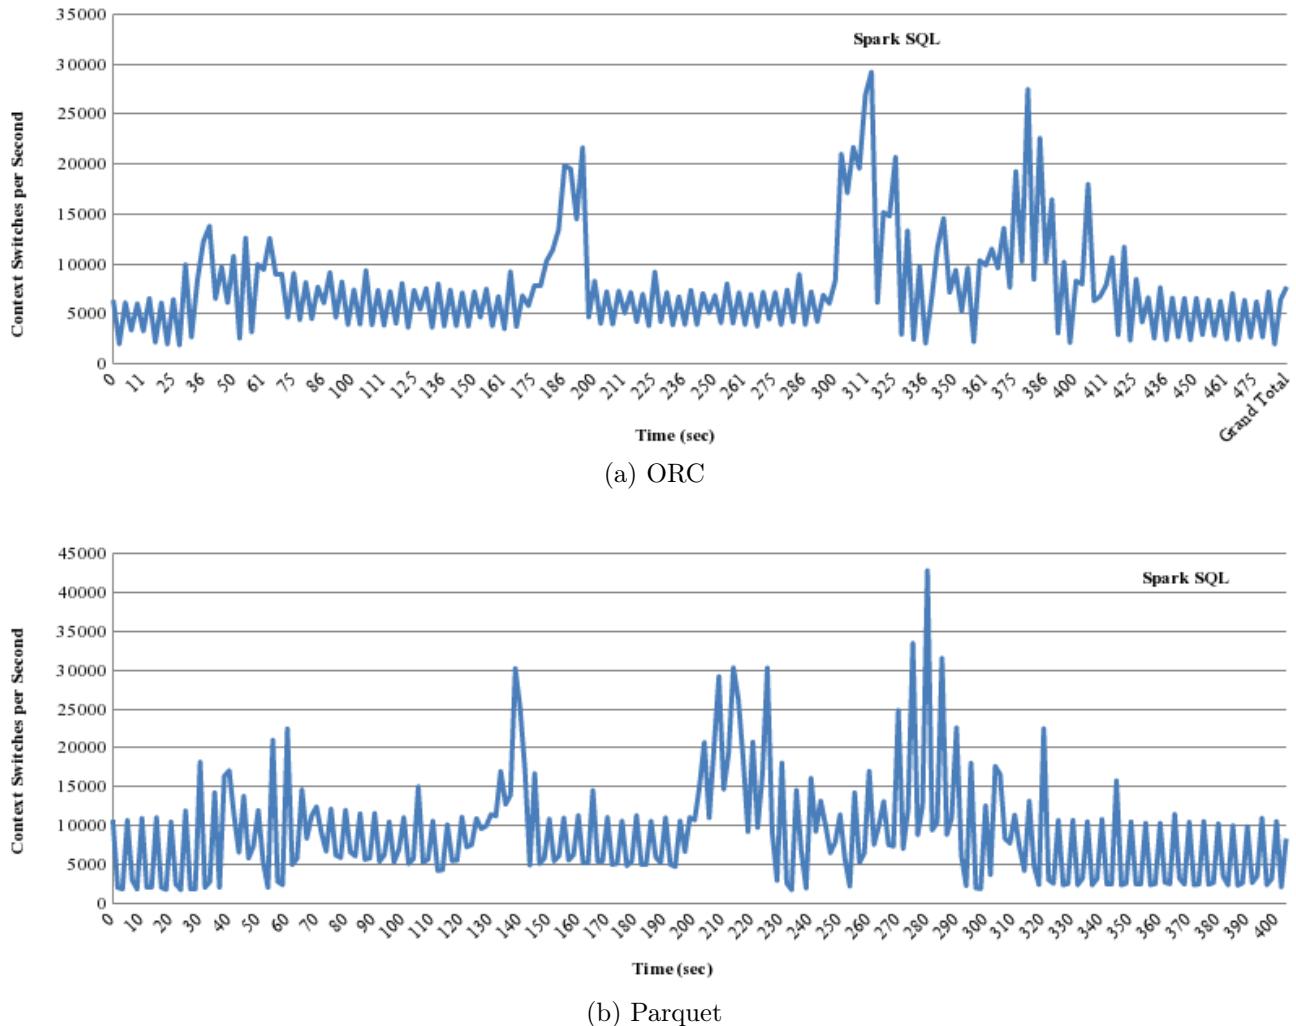


Figure 84: VMSTAT Context Switches

4.11 ORC - NoCompression vs. Snappy

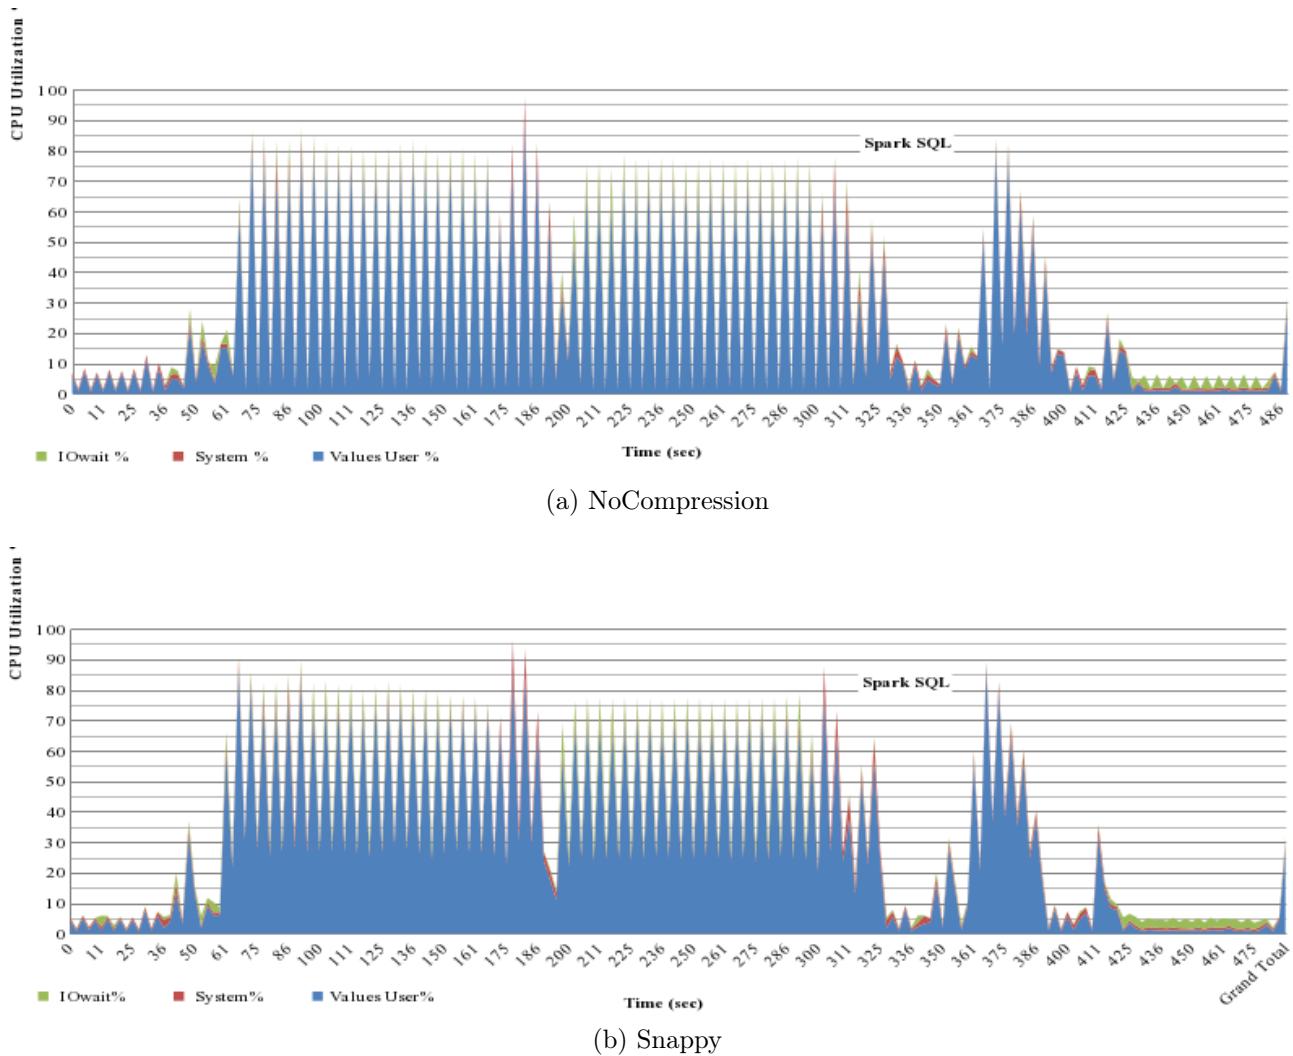


Figure 85: CPU.

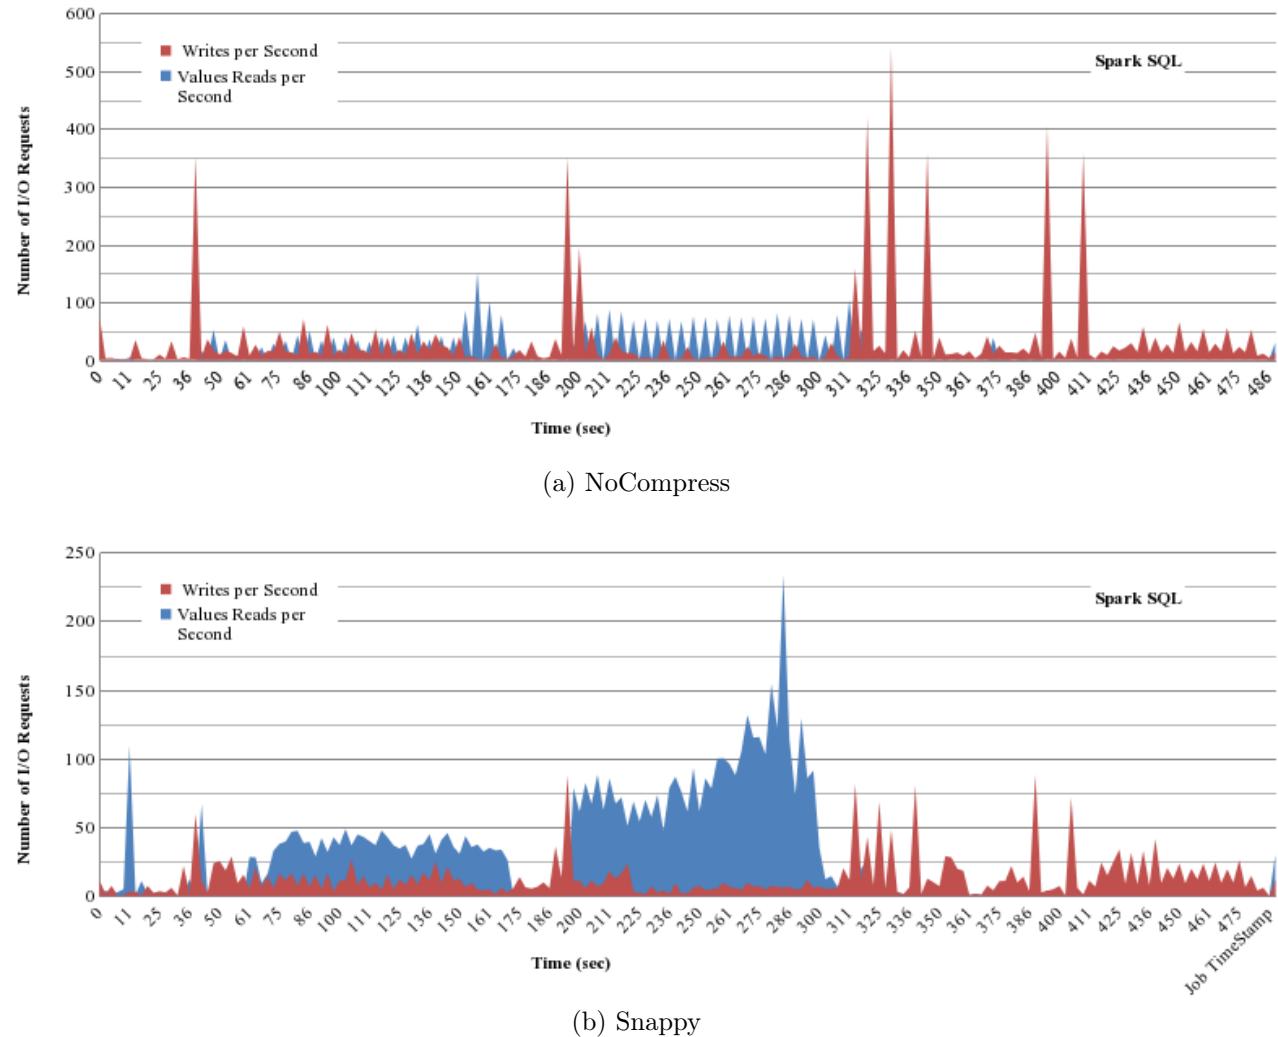


Figure 86: Disk Requests.

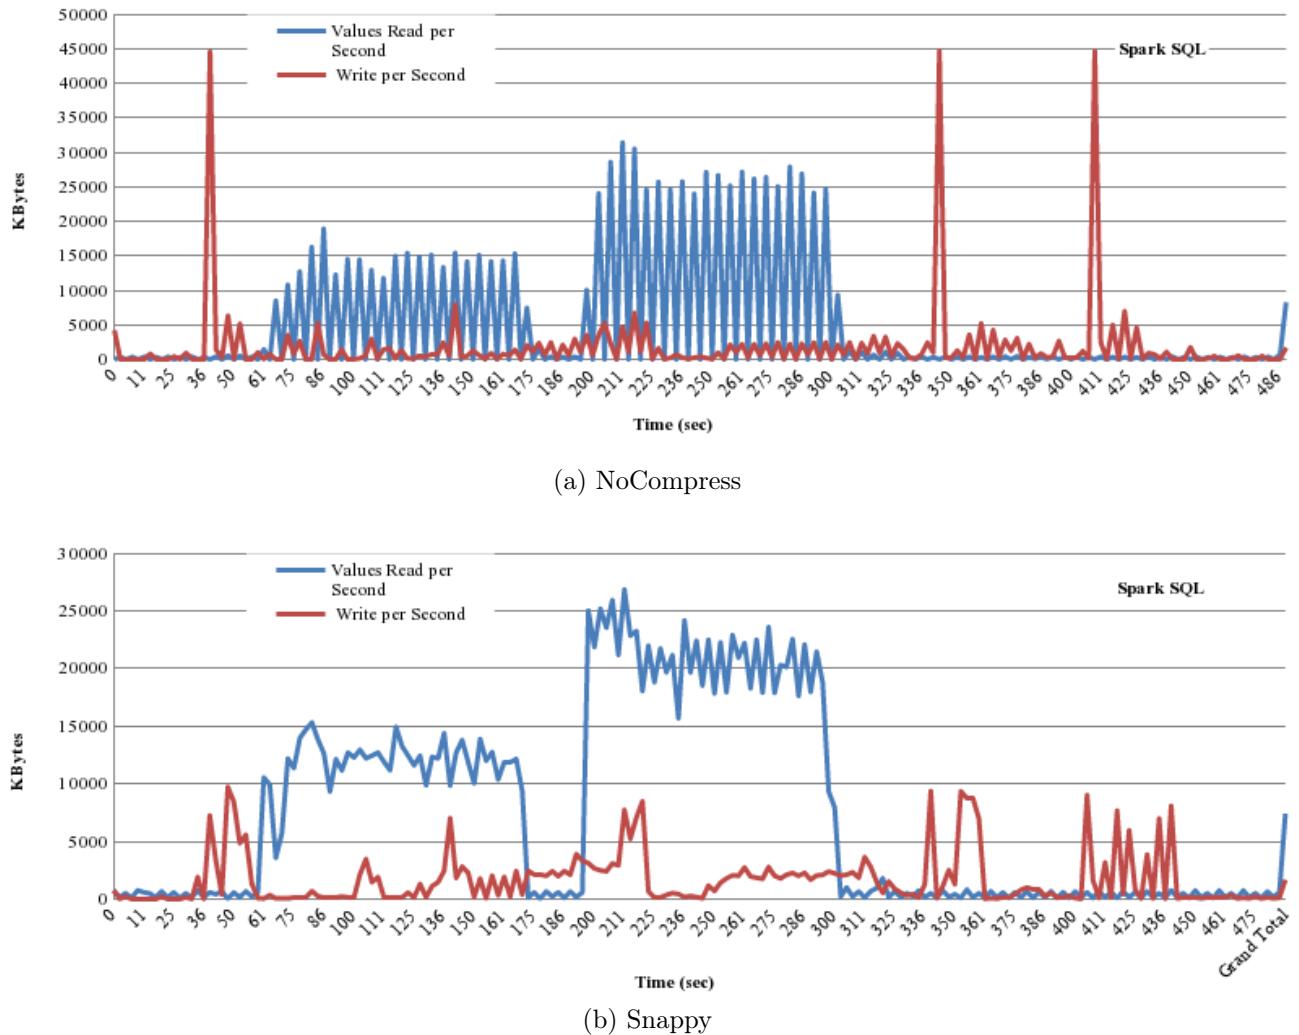


Figure 87: Disk Bandwidth.

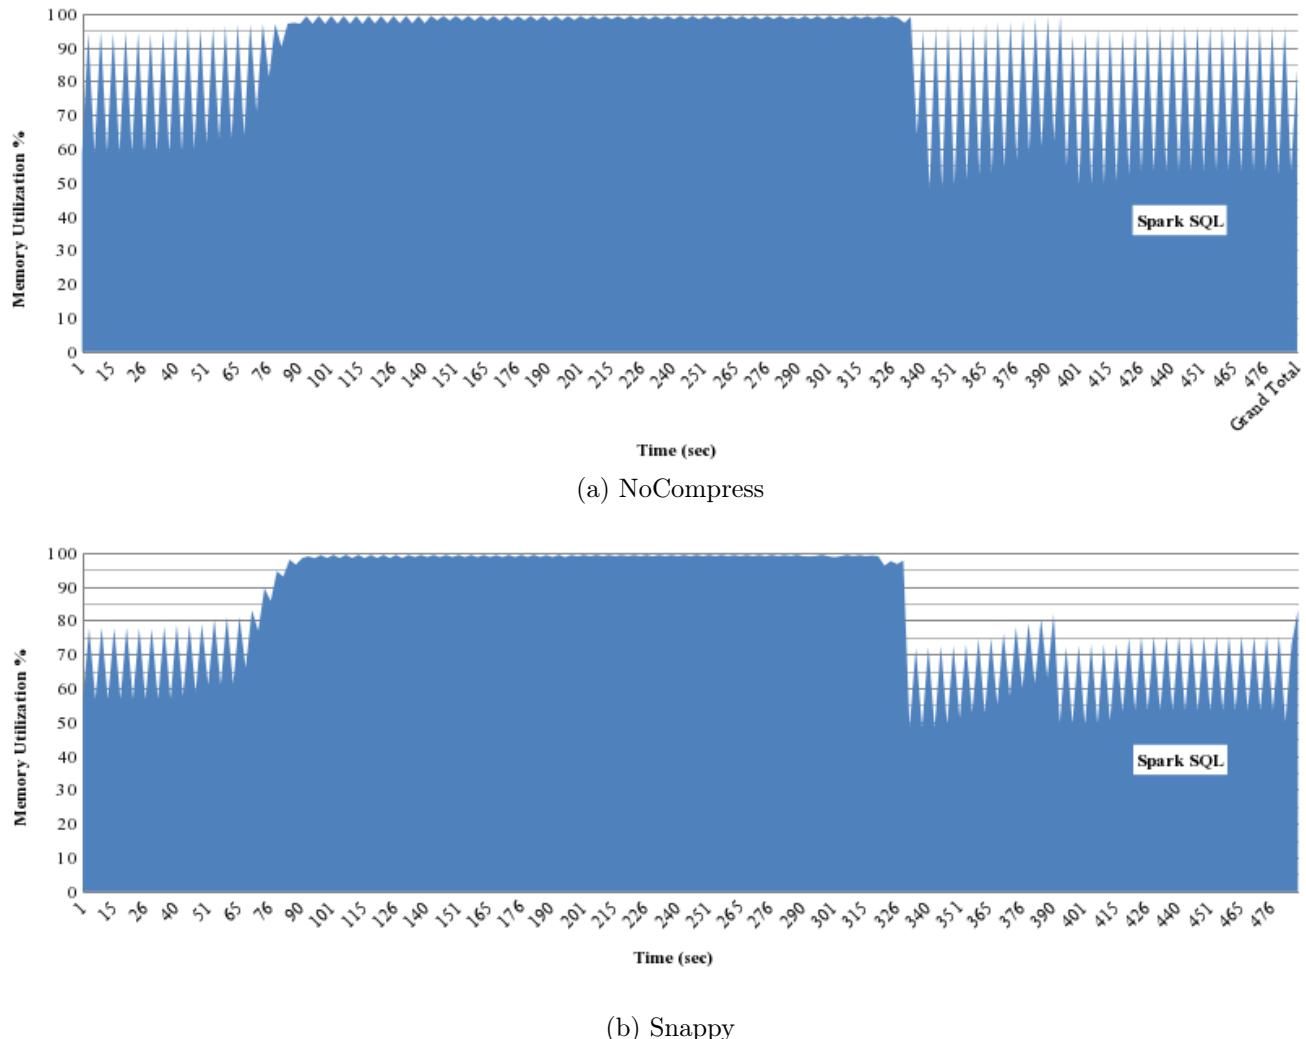


Figure 88: Free Memory.

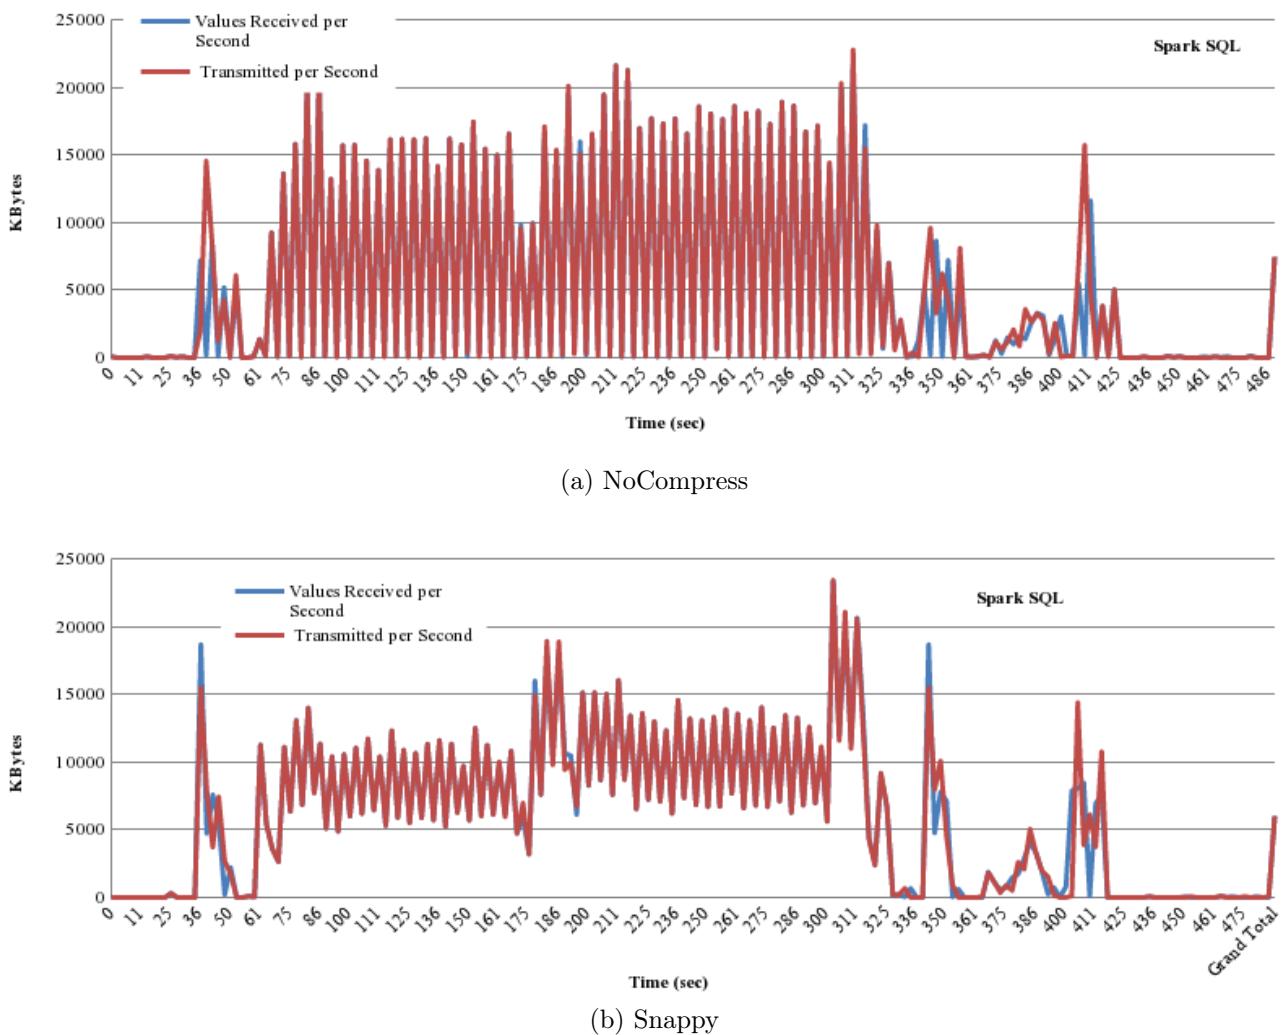


Figure 89: Network Utilization.

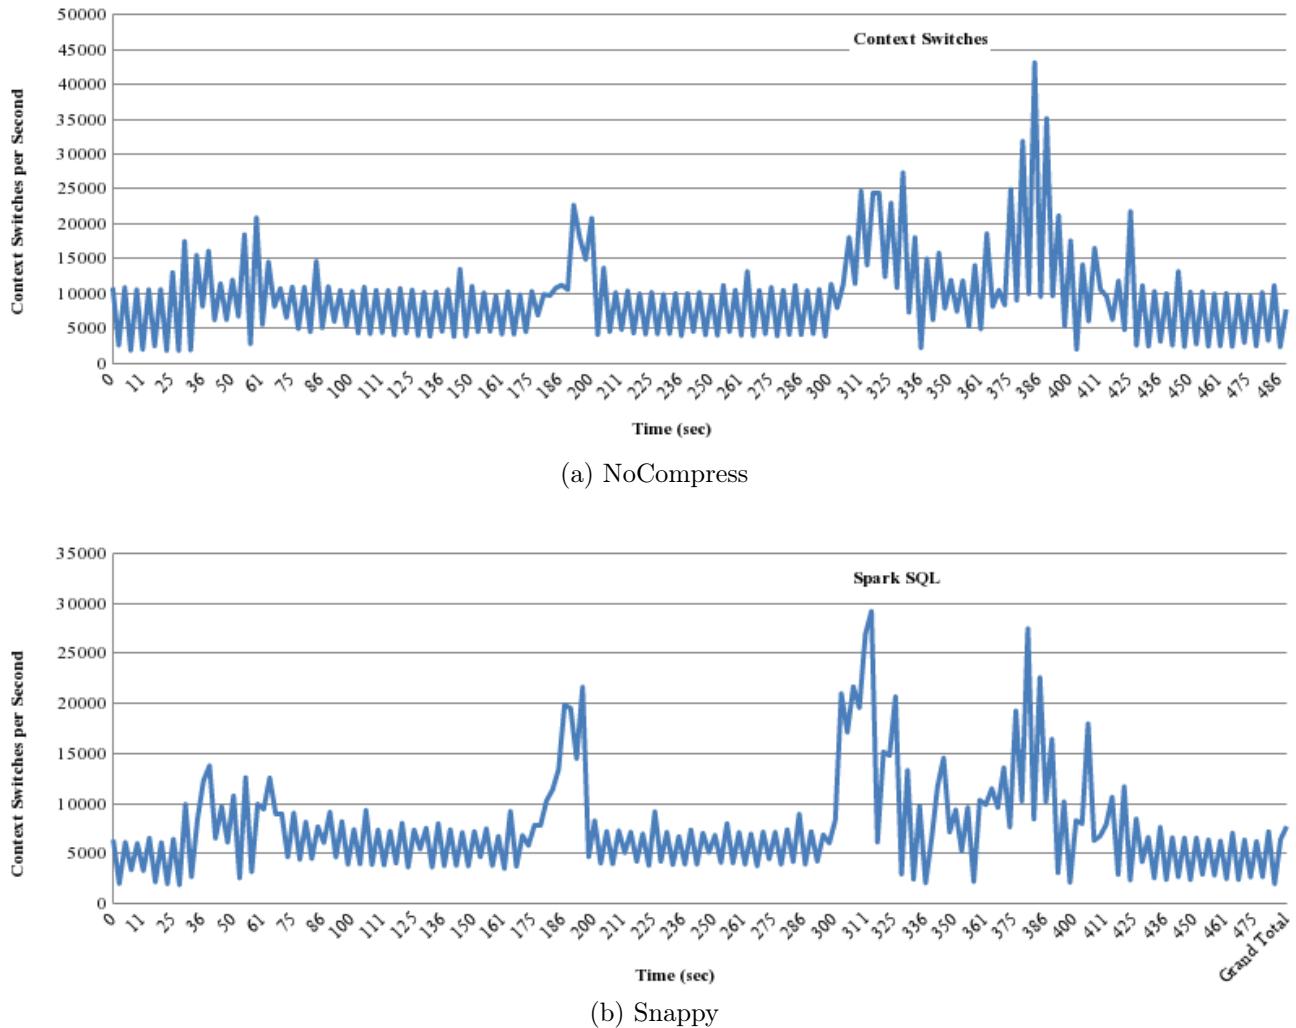


Figure 90: VMSTAT Context Switches

4.12 Parquet - NoCompression vs Snappy

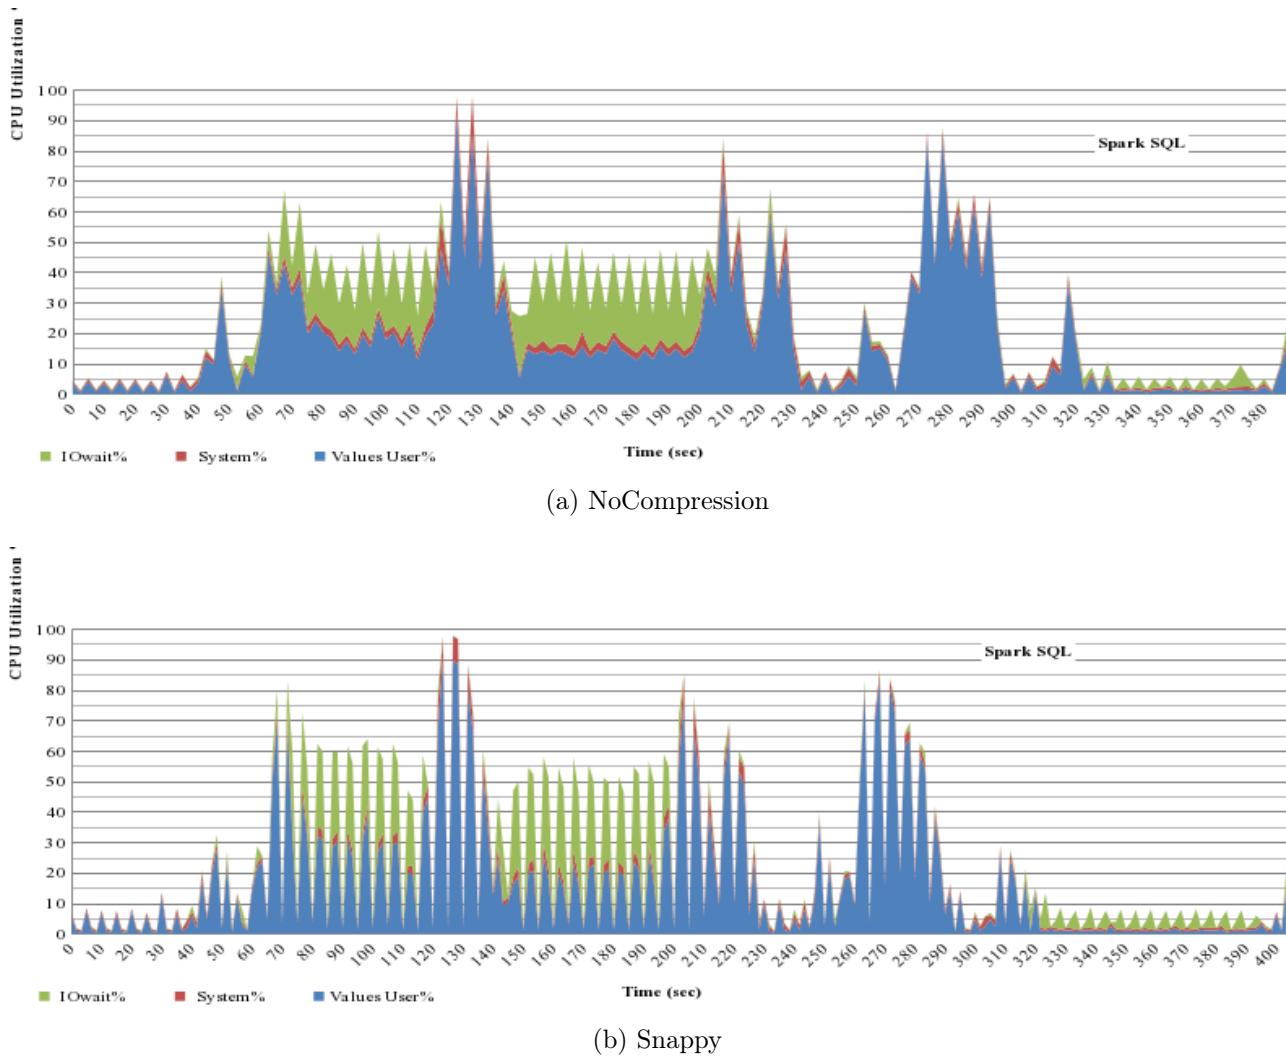


Figure 91: CPU.

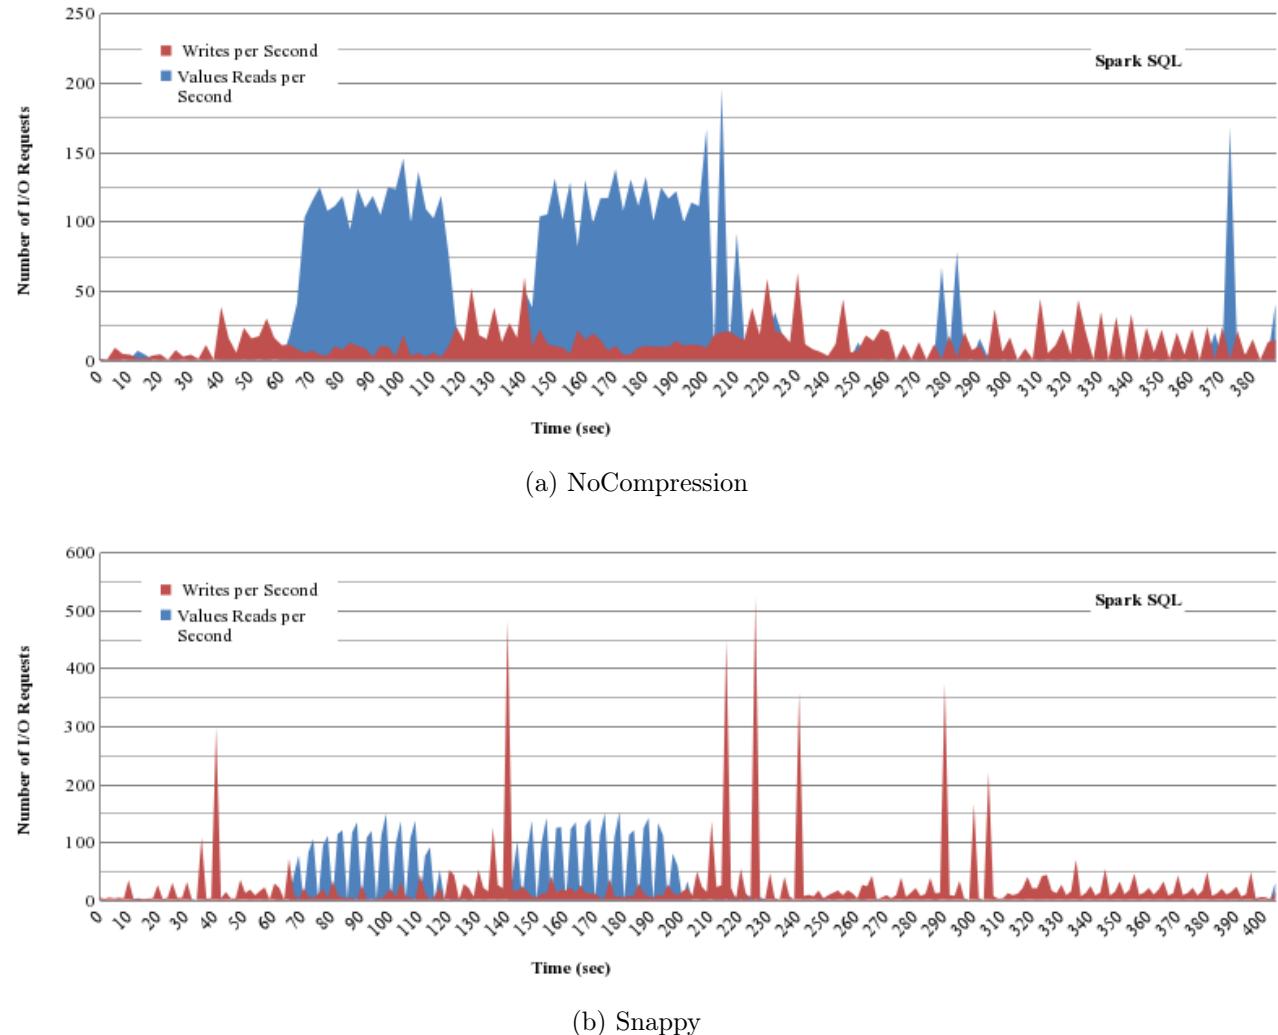


Figure 92: Disk Requests.

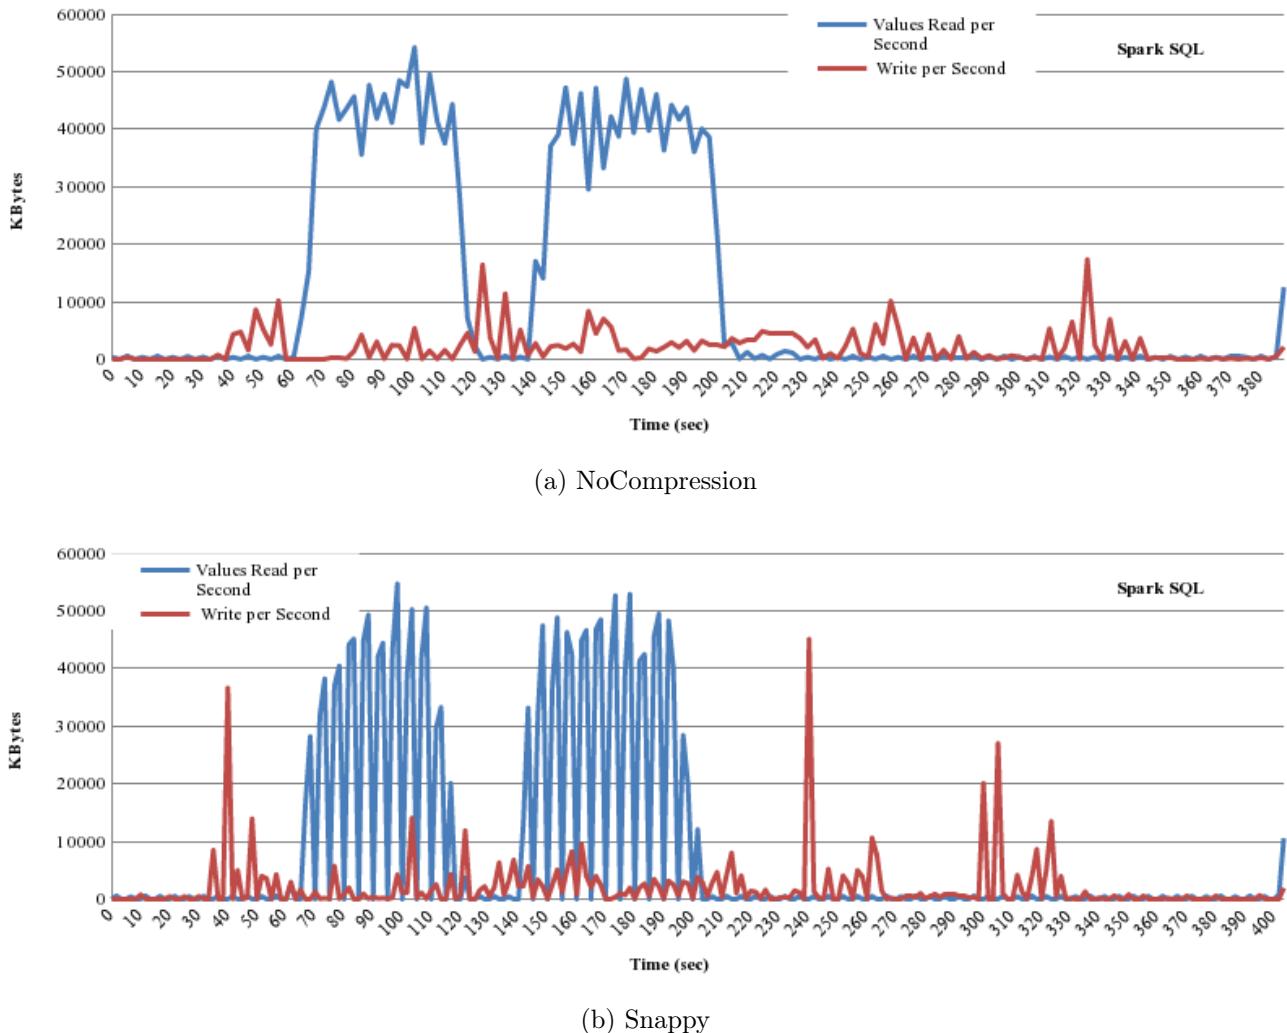


Figure 93: Disk Bandwidth.

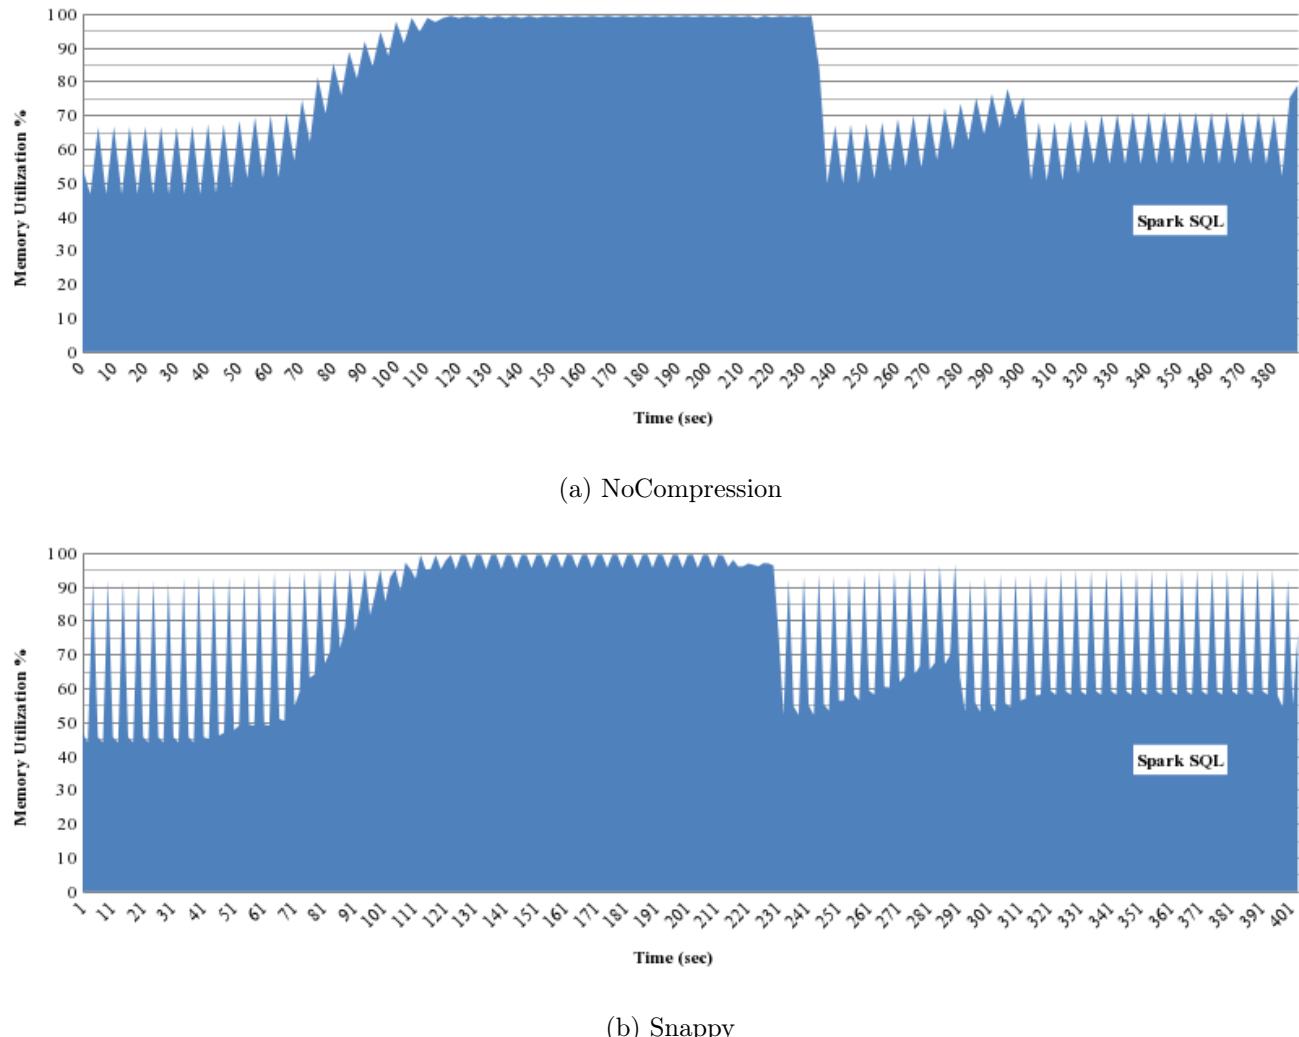


Figure 94: Free Memory.

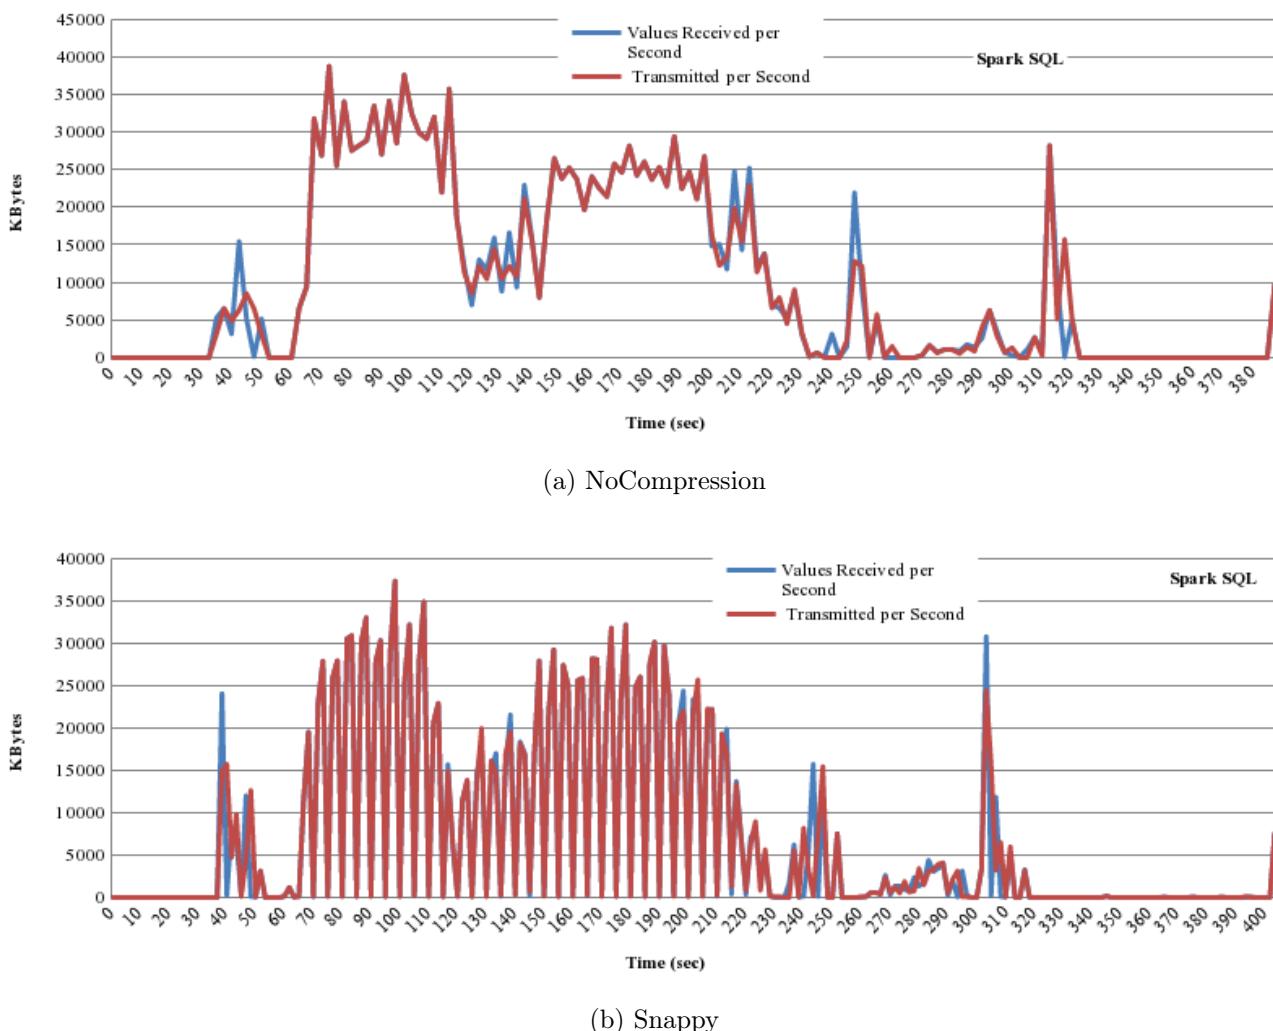


Figure 95: Network Utilization.

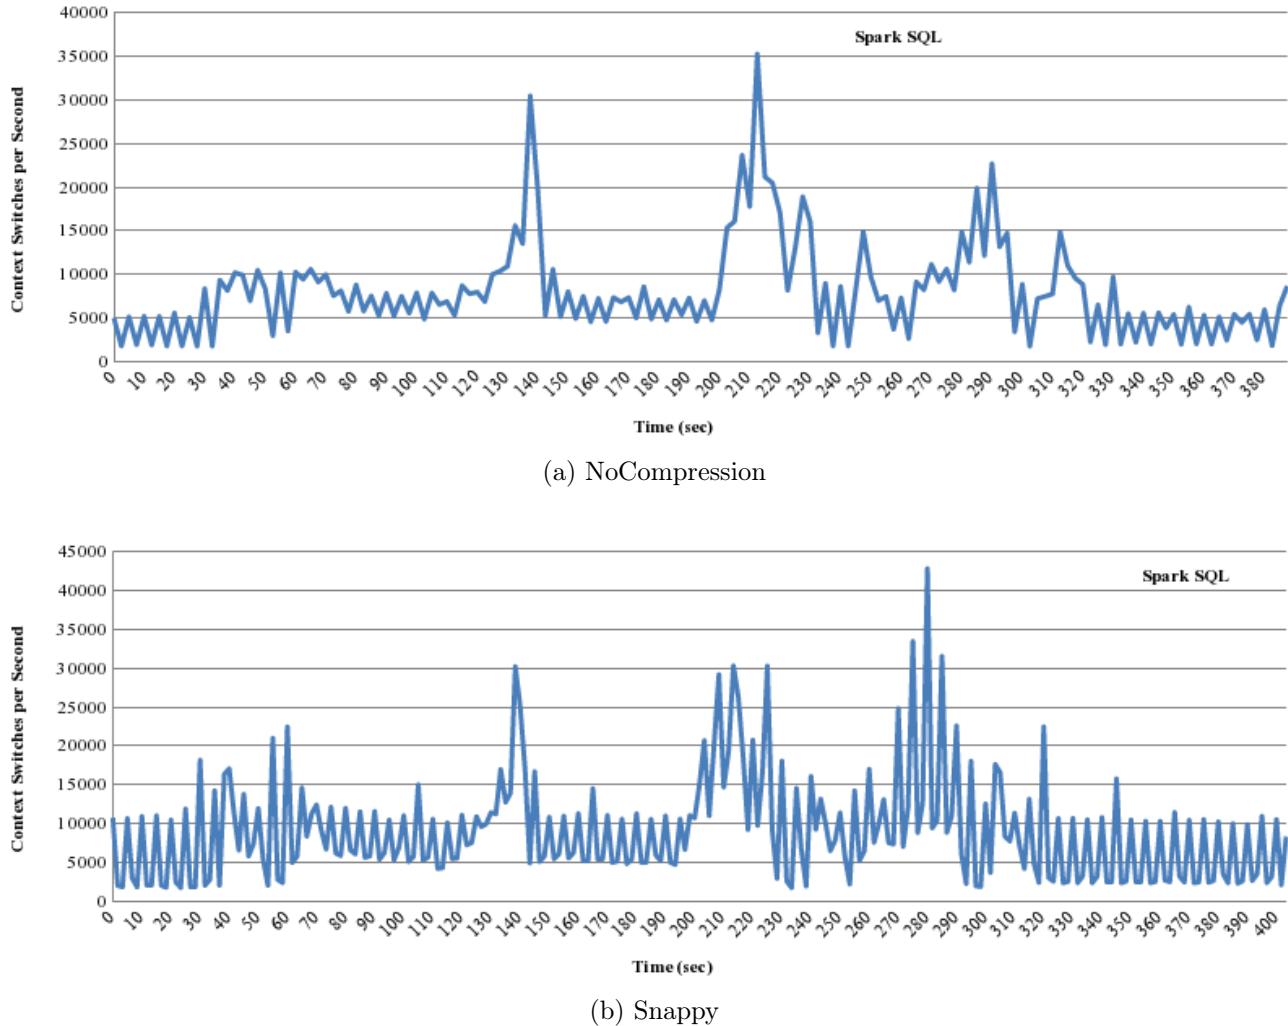


Figure 96: VMSTAT Context Switches